

**NANYANG  
TECHNOLOGICAL  
UNIVERSITY**  
**SINGAPORE**

## MH4510 Statistical Learning and Data Mining

Group: Blush of Dawn

### Classification of Blood Cancer Subtypes using Images

Name	Matriculation Number
Alvin Goh Jia Hao	U2140086C
Chong Wei Jun Kelvin	U2140300J
Shannon Lee Wei Kei	U2140284D
Sun Jia Yee Joyce	U2140792E
Sonya Annalise Thong Li Wen	U2140424K

#### Abstract

Blood cancer remains one of the most pressing concerns globally, and its means of diagnosis is still a challenge to this day. Current image-based research on blood cancer cells primarily focuses on single-cell images often neglecting the spatial features between multiple cells. Our aim is to train models for multi-classification task to classify leukemia subtypes using the Acute Lymphoblastic *Leukemia* (ALL) dataset, which comprises a diverse collection of 3,256 multi-cell blood smear images classed according to their maturity subtypes. In this project, we also proposed a new Convolutional Neural Network (CNN) architecture evaluated multiple machine learning models, including SVM, XGBoost, VGG16, DenseNet-201, and ConvNeXt. We focused on optimizing accuracy while considering the misclassification costs associated with different leukemia stages, and found that ConvNeXt and DenseNet-201 pretrained model combined with our proposed architecture performs the best, with DenseNet-201 topping in performance across accuracy and having minimal misclassification costs.

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	Literature Review . . . . .	4
1.2	Gaps in Current Knowledge . . . . .	4
1.3	Research Objectives . . . . .	4
<b>2</b>	<b>Dataset and Pre-processing</b>	<b>5</b>
2.1	Description of Dataset . . . . .	5
2.2	Exploratory Data Analysis . . . . .	6
<b>3</b>	<b>Feature Engineering</b>	<b>9</b>
3.1	Segmented and Original Images for Different Model Trainings . . . . .	10
3.2	Feature Extraction for Machine Learning Models . . . . .	10
3.2.6	Edge Detection Features . . . . .	11
3.2.7	Shape Features . . . . .	12
3.3	Segmentation . . . . .	12
3.4	Data Augmentation . . . . .	13
3.5	Stratified Sampling . . . . .	13
3.6	Class Weighting . . . . .	14
<b>4</b>	<b>Evaluating Misclassification Costs</b>	<b>15</b>
<b>5</b>	<b>Machine Learning Models</b>	<b>16</b>
5.1	Support Vector Machines . . . . .	16
5.1.1	Training process and tuning . . . . .	17
5.1.2	Test and train results . . . . .	18
5.1.3	Misclassification cost results . . . . .	19
5.2	XGBoost . . . . .	20
5.2.1	Training process and tuning . . . . .	21
5.2.2	Test and train results . . . . .	22
5.2.3	Misclassification cost results . . . . .	23
<b>6</b>	<b>Convolutional Neural Network</b>	<b>24</b>
6.1	Convolutional Layer . . . . .	25
6.2	Pooling Layer . . . . .	26
6.3	Proposed Architecture . . . . .	26
6.3.1	Differences between Research Paper's Architecture and Our Proposed Architecture . . . . .	26
6.3.2	Segmentation Block . . . . .	27
6.3.3	Architecture Overview . . . . .	28
6.4	VGG16 Architecture . . . . .	29
6.4.1	Model description . . . . .	29
6.4.2	Training process and tuning . . . . .	30
6.4.3	Test and train results . . . . .	31
6.4.4	Misclassification cost . . . . .	32
6.5	DenseNet-201 . . . . .	33
6.5.1	Model description . . . . .	33
6.5.2	Training process and tuning . . . . .	33
6.5.3	Test and train results . . . . .	34
6.5.4	Misclassification Cost . . . . .	36
6.6	ConvNeXt . . . . .	37
6.6.1	Model description . . . . .	37
6.6.2	Training process and tuning . . . . .	37
6.6.3	Test and train results . . . . .	38
6.6.4	Misclassification Costs . . . . .	39

<b>7</b>	<b>Results and Findings</b>	<b>40</b>
<b>8</b>	<b>Conclusion</b>	<b>40</b>
<b>9</b>	<b>Appendices</b>	<b>42</b>
9.1	SVM . . . . .	42
9.1.1	XGBoost . . . . .	54
9.1.2	VGG16 . . . . .	65
9.1.3	DenseNet-201 . . . . .	75
9.1.4	ConvNeXt . . . . .	85

# 1 Introduction

Blood cancer remains one of the most critical forms of cancer worldwide, and has been increasing annually since 1990, reaching a total of 1.3 million in 2019 [1]. Accurate risk assessments are vital for improving treatment outcomes, as early intervention can significantly increase survival rates [2]. Current diagnostic practices, however, including microscopic evaluation and histopathological methods, often rely on subjective visual analysis, which may lead to inconsistencies or delays in diagnosis [3]. As a result, deep learning models, particularly convolutional neural networks (CNNs), for image-based classification offer a promising direction for enhancing accuracy and efficiency in blood cancer diagnostics [4]. Such advancements could revolutionize how early-stage cancers like leukemia are identified and managed, especially by introducing automated tools capable of handling large datasets of blood cell images.

## 1.1 Literature Review

Many studies in artificial intelligence using images have successfully used convolutional neural networks (CNNs) for classifying types of blood cancer, demonstrating high accuracy levels. For instance, in leukemia classification, researchers have successfully applied CNNs to detect and classify various stages of leukemia cells from peripheral blood smear images with promising accuracy. One study used a CNN model to classify Acute Lymphoblastic Leukemia (ALL) cells from microscopic images, achieving an accuracy of over 90 percent on their test set, demonstrating the potential of CNNs in blood cancer diagnostics [5].

## 1.2 Gaps in Current Knowledge

While some studies have employed machine learning techniques and artificial neural networks for classification of leukemia using CNN, most have particularly focused on classification associated with single-cell images of blood cancer progression based on images of cell morphology. One particular study comes from research in Iran where a neural network model is used to classify developmental classes of B-Cell leukemia based on multi-cell images [6], but their objectives were limited to only using and developing CNN models. Another study was done recently that adapted ConvNeXt to blood cell classification, but they did not focus on specifically blood cancer subtypes [7]. Despite advancements in blood cancer classification through deep learning, there remains considerable space to explore in classifying the prediction of blood cancer maturity, using multi-cell images that are based on their appearances.

## 1.3 Research Objectives

The primary research objective of this study is to adapt machine learning methods and compare various pretrained CNN models with modern ConvNeXt architecture for classification tasks, using peripheral smeared blood images. We will first build a feature extraction function to extract features from images and use machine learning models such as Support Vector Machine (SVM) and XGBoost to train on those features for multi-class classification. Subsequently, we move on to propose a new CNN architecture that does both segmentation and classification, training directly on our image dataset for classification tasks, and using pre-trained network architectures such as VGG16, DenseNet-201 and a recent model ConvNeXt, to evaluate which pretrained models are the most optimal for our task and architecture.

## 2 Dataset and Pre-processing

### 2.1 Description of Dataset

The models in this study will be using the archived dataset of Acute Lymphoblastic Leukemia (ALL) blood cancer, which provides a large collection of 3,242 peripheral blood smear (PBS) images, organized into two subfolders where one contains the original images captured, and the other one pre-segmented. Within each of these two subfolders contains folders which hold images based on developmental subtypes: Benign, Early Pre-B, Pre-B, and Pro-B ALL. The dataset was prepared in the bone marrow laboratory of Taleqani Hospital in Tehran, Iran.

The dataset consists of image files in JPEG format. Figure 1 shows the folder structures.

```
archive/
|
└── segmented/
    ├── Benign/
    ├── Early/
    ├── Pre/
    └── Pro/
|
└── original/
    ├── Benign/
    ├── Early/
    ├── Pre/
    └── Pro/
```

Figure 1: Directory map for the chosen dataset

Figure 2 is a table showing the class names of the developmental subtypes, with their description explaining what each of them means. Benign is considered to be the least aggressive stage amongst all other developments, while Pro is considered to be the most severe and advanced stage of Leukemia B-Cells.

Subclass	Description
Benign	Non-cancerous, healthy cells
Early	Early stages of leukemia
Pre	Pre-stage abnormal cells
Pro	Advanced leukemia cells

Figure 2: Directory map for the chosen dataset

Figure 3 shows some sample images from the original folder of image datasets, each of which were labeled according to their maturity stages.

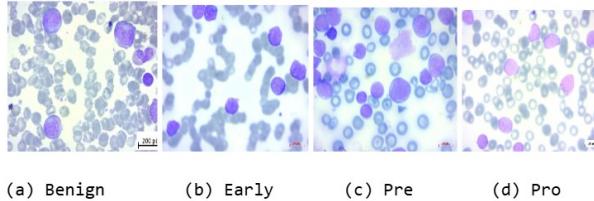


Figure 3: Directory map for the chosen dataset

## 2.2 Exploratory Data Analysis

We will be cleaning the data we have, and since our raw data consist mostly of pictures instead of numbers, we will be correcting the resolution of the images and taking appropriate measures during the model training process later on.

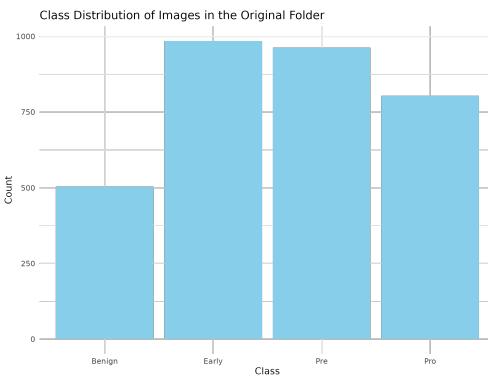


Figure 4: Class Distribution

The class distribution plot in Figure 4 indicates the number of images present in each category. The dataset comprises 3256 images, distributed across the classes as follows: Benign (504 images), Early (985 images), Pre (963 images), and Pro (804 images). The majority of images belong to the Early and Pre class, suggesting a potential imbalance that could affect certain models performance.

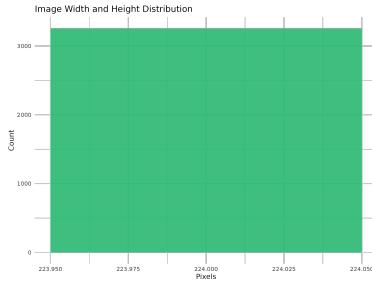


Figure 5: Image dimensions

The analysis of image dimensions as shown in the plot of Figure 5 revealed that all images were consistently sized at  $224 \times 224$  pixels. This uniformity is advantageous as it simplifies preprocessing steps and aligns with standard input requirements for convolutional neural networks (CNNs). The consistent aspect ratio of 1:1 further ensures that the images maintain their intended visual properties during training, thus reducing potential biases in model predictions.

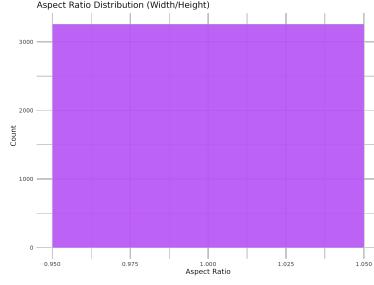


Figure 6: Aspect ratios

The aspect ratio of all images according to Figure 6 was consistently measured at 1.0, affirming that the dataset contains square images. This characteristic is essential for maintaining visual integrity when inputting the images into machine learning models. A uniform aspect ratio eliminates potential distortions that may arise from varying image shapes, thereby facilitating better feature extraction by the models.

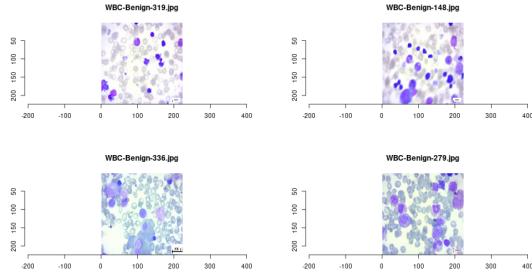


Figure 7: Benign

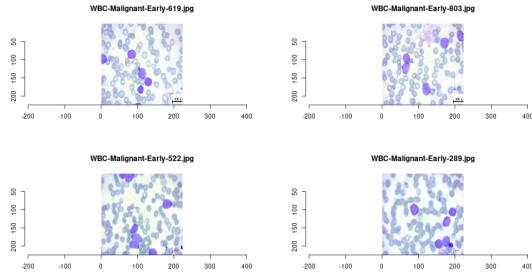


Figure 8: Early

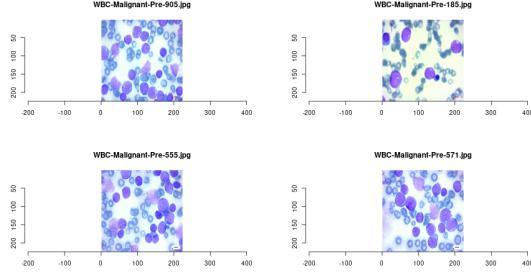


Figure 9: Pre

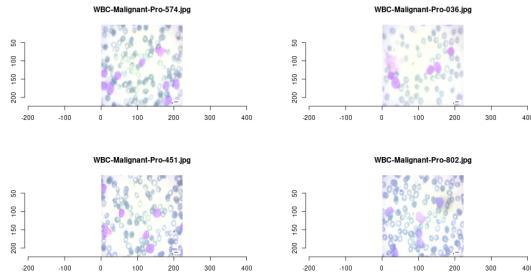


Figure 10: Pro

To further understand the dataset, a selection of sample images from each class was visually inspected. The images displayed clear differentiation between classes especially between Early and Pre, as well as Pro which shows a sudden lightness in colour tone of the cells, reflecting variations in cellular morphology that could be critical for accurate classification. This visual assessment is crucial for confirming the relevance of the dataset to the research objectives and the potential for successful application in automated diagnostic systems. From the exploratory data analysis of the image data and resolution, the only aspect we need to focus on is rebalancing the data's scarcity and unequal sizes in each class via augmentation techniques.

### 3 Feature Engineering

The loading of dataset process begins by defining the full dataset of image paths and their corresponding class labels. It first gathers all image file paths from the directory /kaggle/input/leukemia-images/Original for the four defined classes ("Benign", "Early", "Pre", "Pro") and stores them in the file \_paths vector. Each images class is recorded in the class\_labels vector. This initial step ensures that the dataset is categorized and ready for splitting. Figure 14 below shows the first 6 rows of the table output for the full dataframe before splitting.

A data.frame: 6 × 2		
	image_filename	class
	<chr>	<chr>
1	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-001.jpg	Benign
2	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-002.jpg	Benign
3	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-003.jpg	Benign
4	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-004.jpg	Benign
5	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-005.jpg	Benign
6	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-006.jpg	Benign

Figure 11: Full dataset

We call this DataFrame full\_df, which associates each image path with its corresponding class. A numeric representation of each class in the column class\_numeric is also generated by converting the class labels into numeric indices (0 for Benign, 1 for Early, 2 for Pre, and 3 for Pro). This numeric conversion is important because machine learning models typically require numerical data for training. The class labels are mapped to integers using the factor function, which is then adjusted to start from 0 (using as.numeric(factor(...)) - 1).

A data.frame: 6 × 3			
	image_filename	class	class_numeric
	<chr>	<chr>	<dbl>
1618	/kaggle/input/leukemia-images/Original/Pre/WBC-Malignant-Pre-129.jpg	Pre	2
1098	/kaggle/input/leukemia-images/Original/Early/WBC-Malignant-Early-594.jpg	Early	1
3176	/kaggle/input/leukemia-images/Original/Pro/WBC-Malignant-Pro-724.jpg	Pro	3
104	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-104.jpg	Benign	0
514	/kaggle/input/leukemia-images/Original/Early/WBC-Malignant-Early-010.jpg	Early	1
3161	/kaggle/input/leukemia-images/Original/Pro/WBC-Malignant-Pro-709.jpg	Pro	3

Figure 12: Shuffled train dataset

The dataset is then split into training, validation, and test sets using a stratified approach. This ensures that the class distribution is maintained in each subset. First, a random 64% of the data is allocated to the training set using createDataPartition from the caret library. The remaining 36% is set aside for the validation and test sets. After this, the remaining data (36%) is further divided, with 20% designated for validation and 16% for testing. Only the training data is specifically shuffled before creating the validation set.

A data.frame: 6 × 3			
	image_filename	class	class_numeric
	<chr>	<chr>	<dbl>
6	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-006.jpg	Benign	0
8	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-008.jpg	Benign	0
15	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-015.jpg	Benign	0
19	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-019.jpg	Benign	0
21	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-021.jpg	Benign	0
28	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-028.jpg	Benign	0

Figure 13: Validation dataset

A data.frame: 6 × 3			
	image_filename	class	class_numeric
	<chr>	<chr>	<dbl>
1	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-001.jpg	Benign	0
3	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-003.jpg	Benign	0
9	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-009.jpg	Benign	0
12	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-012.jpg	Benign	0
17	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-017.jpg	Benign	0
18	/kaggle/input/leukemia-images/Original/Benign/WBC-Benign-018.jpg	Benign	0

Figure 14: Test dataset

By keeping the validation and test sets in their original order, the model is tested in a consistent, replicable manner, ensuring that the evaluation reflects the model’s generalization capabilities rather than its ability to adapt to random data orderings. In terms of how the trained neural network model will be used in diagnosis, users will only need to input an image of blood cancer cells into the model and they will receive a classification result of one of the 4 subtypes as an output.

### 3.1 Segmented and Original Images for Different Model Trainings

The machine learning methods we will be using for this project are SVM and XGBoost, while the neural networks models we will be using are CNN pretrained models and a new architecture model ConvNeXt. Since these models thrive on rich feature extraction such as colour and intensity, images from the Original folder in the dataset will be chosen to train them. This is because the original images likely retain natural colour variations and contextual information that enhance feature diversity, allowing the model to use informative colour and texture features.

By extracting features like colour histograms and texture descriptors from the original images, SVM and XGBoost can differentiate between classes based on these informative characteristics more effectively compared to the pre-segmented images in the Segmented folder. For CNN models, we will not rely on existing segmented image dataset to train our neural network models. Instead, we propose a new architecture that uses U-Net to do the segmentation for us.

### 3.2 Feature Extraction for Machine Learning Models

Feature extraction is a fundamental step in image processing for machine learning pipelines. It transforms raw data into structured information, enhancing models’ abilities to recognize patterns and make accurate predictions [8]. This process is essential for models like Support Vector Machines (SVMs) and XGboost. Below, we provide a mathematical overview of key feature extraction techniques.

## Color Histogram Features

Color histograms represent the distribution of color intensities within an image. For an RGB image, each pixel is defined by red, green, and blue intensities. A histogram is created for each channel, where intensity values are split into bins, and the frequency of pixels in each bin is recorded.

For a color channel  $C$  with intensity values  $I_j$  divided into  $n$  bins, the histogram for bin  $i$  is defined as:

$$\text{hist}(C, i) = \frac{\sum_j \mathbb{1}(I_j \in \text{bin}_i)}{\text{Total number of pixels}}$$

where  $\mathbb{1}(I_j \in \text{bin}_i)$  is an indicator function that counts if pixel  $I_j$  falls into bin  $i$  [9]. By normalizing, this feature becomes size-invariant, allowing direct comparison across images. Color histograms capture color distribution compactly, simplifying classification.

## Texture Features

Texture features capture patterns in pixel intensities, describing an image's surface characteristics (Haralick et. al, 2007). Key descriptors include:

### Contrast

Contrast measures pixel intensity variation and is calculated as the variance  $\sigma^2$  around the mean intensity  $\mu$ :

$$\text{contrast} = \frac{1}{N} \sum_{j=1}^N (I_j - \mu)^2$$

where  $I_j$  is the intensity of the  $j$ -th pixel,  $N$  is the total number of pixels, and  $\mu$  is the mean intensity.

### Skewness

Skewness indicates asymmetry in intensity distribution. Positive skew suggests more light pixels, while negative skew indicates more dark pixels:

$$\text{skewness} = \frac{1}{N} \sum_{j=1}^N \left( \frac{I_j - \mu}{\sigma} \right)^3$$

### Kurtosis

Kurtosis captures the tailedness of intensity values, indicating how often extreme values occur:

$$\text{kurtosis} = \frac{1}{N} \sum_{j=1}^N \left( \frac{I_j - \mu}{\sigma} \right)^4 - 3$$

These texture features summarize the statistical distribution of intensities, aiding models by highlighting the complexity or smoothness of surfaces. In the kurtosis formula, the subtraction of 3 serves to "normalize" the kurtosis relative to the normal distribution, which has a kurtosis of 3. By subtracting 3, the resulting measure is known as excess kurtosis.

## Edge Detection Features

Edge detection identifies boundaries between regions in an image. Using the Sobel operator, intensity gradients are approximated across  $x$ - and  $y$ -axes [10]:

$$\text{edge magnitude} = \sqrt{\left( \frac{\partial I}{\partial x} \right)^2 + \left( \frac{\partial I}{\partial y} \right)^2}$$

where  $\frac{\partial I}{\partial x}$  and  $\frac{\partial I}{\partial y}$  represent intensity changes along the horizontal and vertical directions, respectively. The mean and standard deviation of gradients provide models with a compact measure of image edges, which help in recognizing structural boundaries.

### Shape Features

Shape features describe an object's geometric properties in an image, often derived after binary segmentation [11]. These descriptors include area, perimeter and circularity.

Area ( $A$ ) is the total pixel count in the object, representing size:

$$A = \sum_j \mathbb{1}(I_j \in \text{object pixels})$$

Perimeter ( $P$ ) measures the length of the boundary around the object, capturing contour details. Circularity indicates roundness, comparing the shape to a circle:

$$\text{circularity} = \frac{4\pi A}{P^2}$$

These features allow models to distinguish shapes by providing roundness or elongation dimensions, useful for object detection tasks.

Thus, feature extraction reduces high-dimensional pixel data into compact representations that capture essential image characteristics.

### 3.3 Segmentation

Segmentation in image classification refers to the process of dividing an image into meaningful regions, or segments, based on shared attributes like color, texture, or intensity. This technique is important because it allows for more accurate and efficient analysis by isolating relevant features that aid in object detection and classification. For example, in tasks such as medical imaging or autonomous driving, segmentation helps identify specific areas of interest, like tumors or pedestrians, reducing noise and improving the performance of classification models. By focusing on distinct regions rather than raw images, segmentation allows machine learning models to process information more effectively, leading to better accuracy in predictions [12]. We will explain more in our proposed architecture where we use U-Net for the segmentation process.



Figure 15: Before and after segmentation of Benign and Early images

### 3.4 Data Augmentation

Data augmentation is a crucial technique employed in this study to enhance the performance of the convolutional neural network (CNN) model for blood cancer cell classification tasks. Given the imbalanced distribution of the dataset, which consists of 3,256 images across four classes: Benign (504 images), Early (985 images), Pre (963 images), and Pro (804 images), data augmentation serves to artificially increase the diversity of the training set [13]. We will be using the Original images version of the dataset to illustrate the augmentation.

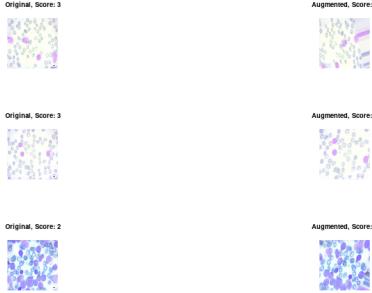


Figure 16: Before and after augmentation

Figure 16 involves a series of transformations applied to the original images, including random rotations (up to 30 degrees), width and height shifts (up to 20 percent of the total dimensions), shear transformations, and zooming (up to 20 percent). Additionally, horizontal flipping is implemented to introduce variations in image orientation. Each of these transformations helps to create multiple variations of the same image, thereby improving the model's ability to generalise and reducing the risk of overfitting.

By rescaling pixel values to a range of [0, 1] and utilising a validation split of 20 percent, the augmentation pipeline not only enhances the model against class imbalance but also aids in improving the overall accuracy of predictions based on their shapes and characteristics. This is particularly useful for the classes with fewer images (e.g., Benign), as it allows the model to learn from a more diverse set of examples without needing additional data.

Data generation in data augmentation involves creating modified versions of existing data samples to increase the diversity and size of the dataset, enhancing model training process and reducing overfitting. This process includes applying a range of transformations to the original data while preserving its underlying patterns and labels. For image data, common augmentations include rotations, flips, cropping, scaling, brightness adjustments, and introducing noise. For text, it might involve paraphrasing, synonym replacement, or shuffling words. By synthetically expanding the dataset, data augmentation helps machine learning models generalize better to unseen scenarios without requiring additional labeled data.

### 3.5 Stratified Sampling

Stratified sampling is a technique used to ensure that each subgroup (or "stratum") within a dataset is proportionally represented in the sample [14]. In this method, the population is divided into distinct, non-overlapping subgroups based on a specific characteristic, such as class labels. A sample is then drawn from each subgroup, with the sample size from each stratum determined proportionally to its size in the overall dataset. This approach is valuable in datasets where certain classes or groups may be underrepresented, especially such as the Benign class in our dataset. We will apply this to all the model trainings in this project.

### 3.6 Class Weighting

The class weighting technique serves as an effective strategy to address the unequal representation of classes in the context of classification. This methodology is especially pertinent in scenarios where certain classes are underrepresented, leading to potential bias in the model's learning process.

Let  $C$  represent the total number of unique classes in the dataset, and let  $n_i$  denote the number of instances in class  $i$ , where  $i \in \{0, 1, 2, \dots, C\}$ . The goal of class weighting is to assign a weight  $w_i$  to each class  $i$  such that the learning algorithm compensates for the underrepresentation of minority classes. The class weight  $w_i$  can be computed using the following formula:

$$w_i = \frac{N}{C \cdot n_i}$$

where:

- $w_i$  is the weight for a sample belonging to class  $i$ ,
- $N$  is the total number of instances in the dataset,
- $C$  is the total number of unique classes,
- $n_i$  is the number of instances in class  $i$ .

This equation serves a dual purpose: it ensures that classes with fewer instances are assigned higher weights, thereby increasing their importance during the training process. As a result, the model is penalised more heavily for misclassifying samples from minority classes, directing the learning algorithm to allocate more attention to these classes. By quantifying the significance of each class through the assigned weights, the model can learn more effectively from all available instances, thereby enhancing its predictive capabilities.

To incorporate the calculated class weights, the loss function of a model is modified to give more importance to certain samples based on their calculated weights. The weighted version of the loss function is:

$$\mathcal{L}_{\text{weighted}} = \frac{1}{N} \sum_{i=1}^N w_i \cdot L_{\text{original}}$$

where:

- $w_i$  is the weight assigned to the  $i$ -th sample,
- $L_{\text{original}}$  is the original loss function of the model.

Classes with larger weights contribute more to the loss. This is typically the case for underrepresented classes, which are given more importance during model training. Conversely, samples with smaller weights contribute less to the loss. This is often the case for overrepresented classes, where their contributions are minimized to avoid dominating the model's learning process.

Thus, the weighted loss function ensures that the model is trained to focus more on samples that are critical for improving the performance on minority classes (or other underrepresented samples), while not overfitting to the majority classes. During training, the model aims to minimize this weighted loss function. The gradient descent optimization will adjust the model's parameters to reduce the weighted error across all samples in the dataset, making the model more balanced in its predictions across different classes.

## 4 Evaluating Misclassification Costs

The cost of misclassification in the context of classifying blood cancer subtypes is significant, particularly due to the medical nature of the task. The impact of incorrectly classifying one class as another can have serious consequences, such as giving wrong treatment to the patients or disregarding serious existing conditions. Below is how we perceived the severity for each misclassification.

On the other hand, misclassifying early conditions as benign is a critical error. Such a mistake could lead to a delay in necessary interventions during the early stages of a condition, at a time when treatment is most effective and outcomes are more favorable. We give this a score of 10.

Misclassifying an "Early" stage as a "Pre" stage might still lead to relatively less severe consequences, such as slight misinterpretations of risk, which could justify a moderate cost of 15. However, misclassifying "Early" as "Pro" could have more serious ramifications, such as unnecessarily aggressive treatments or a failure to intervene early enough, which would justify a higher cost of 25.

The most severe misclassification occurs when pre or pro conditions are misclassified as benign. This error could result in overlooking advanced or aggressive conditions that require immediate attention. As a result, patient outcomes could worsen due to a lack of timely treatment. We give this a score of 20 for misclassifying pre, 30 for pro.

Another possible error is when pre is misclassified as pro, or vice versa. Although this is less severe than misclassifying a condition as benign, it still carries risks. It could lead to inappropriate treatment protocols, such as overtreatment or undertreatment, depending on the incorrect progression stage. We give this score a 10.

With these in mind, we formulate a cost matrix that assigns a unit cost to each misclassification according to their severity:

$$\text{Cost Matrix} = \begin{bmatrix} 0 & 10 & 20 & 30 \\ 10 & 0 & 15 & 25 \\ 20 & 15 & 0 & 10 \\ 30 & 25 & 10 & 0 \end{bmatrix}$$

Here, diagonal values represent correct classifications (cost = 0), and off-diagonal values represent misclassification costs. For the simplicity of this project, we assume that the cost of misclassifying one class with the other is equal as vice versa.

Using the confusion matrix, the total misclassification cost can be calculated by multiplying the frequency of each type of misclassification by its corresponding cost in the cost matrix and summing all costs:

$$\text{Total Cost} = \sum_{i \neq j} \text{ConfusionMatrix}[i, j] \times \text{CostMatrix}[i, j]$$

This metric aggregates the costs of all misclassifications, providing an overall measure of how costly the errors made by the model are.

To normalize the total misclassification cost, the average cost per sample is calculated by dividing the total cost by the total number of samples. This metric provides a more practical evaluation by considering the cost of misclassification on an individual basis, making it easier to compare the model's performance across different datasets or models of varying sizes:

$$\text{Average Cost Per Sample} = \frac{\text{Total Cost}}{\text{Total Number of Samples}}$$

Rather than relying on traditional accuracy metrics, we incorporating a cost-sensitive accuracy which takes into the account of each misclassification cost:

$$\text{Cost-Sensitive Accuracy} = 1 - \frac{\text{Total Misclassification Cost}}{\text{Maximum Possible Cost}}$$

Where the maximum possible cost is the sum of all misclassifications if every sample were misclassified.

The cost-benefit ratio is another important evaluation tool, comparing the benefits of correct classifications with the penalties of misclassifications:

$$\text{Cost-Benefit Ratio} = \frac{\text{Correct Classification Benefit}}{\text{Misclassification Cost}}$$

We can assign a monetary or impact value to correct classifications based on domain knowledge.

Adjust the F1-score for each class by incorporating the misclassification costs:

$$\text{Weighted F1-Score} = \sum_i (F1_i \times W_i)$$

Where  $W_i = \frac{1}{\text{CostMatrix}[i,j]}$  is the weight based on misclassification cost.

These formulas will be used to evaluate the cost of misclassification of every chosen model in this project, and will also be used as metrics for comparison between the models.

## 5 Machine Learning Models

In this section, we will introduce and train chosen models of machine learning for our classification task. Descriptions for each model and their respective techniques of model tuning and adjustments, together with their training and test error results are also included.

### 5.1 Support Vector Machines

Support Vector Machines (SVM) are a class of supervised learning models commonly used for classification and regression tasks. The fundamental idea behind SVM is to find a hyperplane (or decision boundary) that best separates data points of different classes in a high-dimensional feature space. In the case of classification tasks, such as in our project, SVM can predict continuous values by fitting a hyperplane that minimizes the error margin while maintaining a balance between model complexity and accuracy.

Given a dataset  $(\mathbf{x}_i, y_i)$ , where  $\mathbf{x}_i \in \mathbb{R}^n$  are feature vectors and  $y_i \in \{-1, 1\}$  are class labels, SVM seeks to maximize the margin  $M$  between the decision boundary and the nearest data points, known as the support vectors. The decision function is given by:

$$f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b,$$

where  $\mathbf{w}$  is the weight vector and  $b$  is the bias term. The optimization problem for a soft-margin SVM is:

$$\min_{\mathbf{w}, b, \xi} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \xi_i,$$

subject to the constraints:

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0,$$

where  $\xi_i$  are slack variables that allow for misclassification, and  $C > 0$  is a regularization parameter that controls the trade-off between margin maximization and misclassification penalties.

When data is not linearly separable, the SVM uses a kernel trick to map the input features to a higher-dimensional space, where linear separation becomes possible. The Radial Basis Function (RBF) kernel is given by:

$$K(x, y) = \exp \left( -\frac{\|x - y\|^2}{2\sigma^2} \right)$$

where:

- $K(x, y)$ : The value of the kernel function between the two data points  $x$  and  $y$ .
- $x, y$ : The input data points (vectors) in the feature space.
- $\|x - y\|^2$ : The squared Euclidean distance between  $x$  and  $y$ , which measures how far apart the points are in the feature space.
- $\sigma$ : The bandwidth parameter (also known as the kernel parameter) that controls the spread or smoothness of the kernel. It defines the width of the Gaussian function. Smaller values of  $\sigma$  make the kernel more sensitive to local variations, while larger values make it smoother.

The dual form of the optimization problem becomes [15]:

$$\max_{\alpha} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j),$$

subject to the constraints:

$$\sum_{i=1}^N \alpha_i y_i = 0, \quad 0 \leq \alpha_i \leq C.$$

Here,  $\alpha_i$  are the Lagrange multipliers, and  $K(\mathbf{x}_i, \mathbf{x}_j)$  is the kernel function.

To find the optimal values of the regularization parameter  $C$  and the kernel parameter  $\sigma$ , a grid search with cross-validation is used. The parameter  $C$  determines the trade-off between achieving low error on the training data and maximizing the margin. A high  $C$  prioritizes minimizing classification errors, which may lead to overfitting, while a low  $C$  favors larger margins, potentially leading to underfitting. The parameter  $\sigma$  controls the influence of individual data points. A smaller  $\sigma$  produces a smoother decision boundary, while a larger  $\sigma$  allows the model to capture more detailed variations in the data.

In the case of class imbalance, we assign weights to the classes based on their frequencies. The weight for each class is calculated with the same formula as given in section 3.6.

The performance of the model is evaluated using metrics such as accuracy, precision, recall, and the confusion matrix. The decision function  $f(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x}) + b$ , where  $\Phi(\mathbf{x})$  is the nonlinear mapping induced by the RBF kernel, determines the predicted class. Cross-validation is used to assess the accuracy of the model by averaging performance over multiple data splits.

### 5.1.1 Training process and tuning

The training process first begins by setting up the environment and loading libraries such as keras, tensorflow, and EBImage for image processing and machine learning tasks. Image data is loaded from a directory structure, where each subdirectory represents a class. The code organizes the data into a structured dataframe, assigns numeric labels to classes, and splits the dataset into training, validation, and test sets using a stratified approach to maintain class distribution. Class weights are computed to address potential class imbalance issues.

A data.frame: 6 × 64									
	feature_1	feature_2	feature_3	feature_4	feature_5	feature_6	feature_7	feature_8	feature_9
1	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
2	0.03163656	2.6236003	3.1068061	1.5554515	0.59113280	0.1519493	-0.1634713	-0.3242494	-0.50210272
3	-0.06400671	-0.2602499	-0.3348147	-0.4120166	-0.46574645	-0.5371078	-0.6452848	-0.6382741	-0.09371064
4	-0.25529325	-0.2883850	-0.3505538	-0.4146364	-0.46873620	-0.5939824	-0.7390973	-0.8789369	-0.94534665
5	0.22292310	-0.1195743	-0.1721771	-0.1395576	-0.02176747	0.8902246	2.9614860	3.2212775	1.79365058
6	-0.25529325	-0.2602499	-0.3453074	-0.3989176	-0.29234052	0.2481986	2.2455968	2.0868520	0.49340883

Figure 17: Extracted and filtered train features

Images are resized and processed to generate features, including color histograms, texture descriptors, edge intensity, and shape features derived from segmentation. The feature extraction function supports validation to ensure consistent output dimensions across samples. The script also analyzes and resolves inconsistencies in feature lengths by padding or truncating as necessary. Features are standardized to facilitate effective model training, and near-zero variance (NZV) features are identified and excluded to prevent model overfitting or inefficiency. The remaining features are used to build training, validation, and test datasets, which are ready for SVM training. Figure 16 shows the train features after extraction, standardization and filtered of NZV.

The SVM model undergoes hyperparameter tuning using a radial basis function (RBF) kernel. A grid search is conducted over a range of values from  $10^{-3}$  to  $10^3$  for C (cost) and  $\sigma$  (RBF kernel parameter) via 5-fold cross-validation. Class weights are applied to handle class imbalances during training. The script selects the best model based on validation accuracy for each fold and evaluates its performance on the test set. Subsequently, we report the summary statistics for training and validation accuracy across folds. Selected optimal hyperparameters were then used to retrain the final model. The trained final model is then evaluated on all datasets (training, validation, test), with confusion matrix for its performance on test set generated for detailed insights.

### 5.1.2 Test and train results

```
Cross-validation results:
```

	fold	train_accuracy	valid_accuracy	sigma	C
Accuracy	1	0.9807692	0.8865031	0.01	10
Accuracy1	2	0.9976134	0.8742331	0.01	100
Accuracy2	3	0.9952038	0.8895766	0.01	100
Accuracy3	4	1.0000000	0.8757669	0.01	1000
Accuracy4	5	0.9976077	0.9064417	0.01	100

Figure 18: Cross validation results for SVM

From the cross-validation results in figure 18, the most optimal hyperparameters are  $\sigma = 0.01$  and  $C = 100$  is the 5-th fold having validation accuracy of 0.9064.

Reference					
Prediction	Benign	Early	Pre	Pro	
Benign	70	5	1	3	
Early	10	150	3	0	
Pre	0	2	148	0	
Pro	0	0	1	125	

Overall Statistics					
Accuracy :	0.9517				
95% CI :	(0.9296,	0.9685)			
No Information Rate :	0.3031				
P-Value [Acc > NIR] :	< 2.2e-16				
Kappa :	0.9344				

Figure 19: Confusion matrix for SVM

The confusion matrix reveals the models predictions across four classes Benign, Early, Pre, and Pro. It shows strong performance, with the majority of predictions falling along the diagonal (correct classifications). For instance, 70 out of 79 Benign cases and 150 out of 163 Early cases were correctly classified, although there is occasional confusion, such as Benign misclassified as Early (5 cases) and Early as Benign (10 cases). Similarly, "Pre" and "Pro" classes have minimal misclassifications (3 combined). With an overall accuracy of 95.17% and a Kappa score of 0.9344, the model demonstrates excellent performance and balanced classification, making it highly reliable for practical use.

Class-Specific Statistics:							
	Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision		
Class: Benign	0.8750000	0.9794521	0.8860759	0.9772210	0.8860759		
Class: Early	0.9554140	0.9639889	0.9282454	0.9802817	0.9202454		
Class: Pre	0.9673203	0.9945205	0.9866667	0.9864130	0.9866667		
Class: Pro	0.9765625	0.9974359	0.9920635	0.9923469	0.9920635		
	Recall	F1	Prevalence	Detection Rate			
Class: Benign	0.8750000	0.8805031	0.1544482	0.1351351			
Class: Early	0.9554140	0.9375000	0.3838888	0.2895753			
Class: Pre	0.9673203	0.9768977	0.2953668	0.2857143			
Class: Pro	0.9765625	0.9842520	0.2471042	0.2413127			
	Detection	Prevalence	Balanced Accuracy				
Class: Benign	0.1525097	0.9272260					
Class: Early	0.3146718	0.9597015					
Class: Pre	0.2895753	0.9809204					
Class: Pro	0.2432432	0.9869992					

Figure 20: Class statistics by SVM

The model demonstrates strong performance, with a cross-validated mean training accuracy of  $99.42\% \pm 0.77\%$  and validation accuracy of  $88.65\% \pm 1.30\%$ , indicating good generalization. On the test set, it achieved an impressive overall accuracy of 95.17% (95% CI: 92.96%, 96.85%) and a Kappa score of 0.9344, reflecting high agreement with actual labels. Class-specific metrics show balanced performance, with sensitivities ranging from 87.50% (Benign) to 97.66% (Pro) and balanced accuracies exceeding 92% for all classes. Precision and recall are particularly high for "Pre" and "Pro," demonstrating strong classification for these categories. The confusion matrix reveals occasional misclassifications, primarily between "Benign" and "Early," but these are infrequent. The model generalizes well, with a small gap between training accuracy (99.71%) and test accuracy (95.17%), suggesting minimal overfitting.

### 5.1.3 Misclassification cost results

Cost-Weighted Misclassification Matrix:

	Benign	Early	Pre	Pro
Benign	0	50	20	90
Early	100	0	45	0
Pre	0	30	0	0
Pro	0	0	10	0

Total Misclassification Cost: 345

Average Cost Per Sample: 0.6660232

Cost-Sensitive Accuracy: 0.9878905

Weighted F1-score 0.9447882

Figure 21: SVM misclassification cost results

The model demonstrates strong cost-sensitive performance, with a low Average Cost Per Sample (0.666) and high Cost-Sensitive Accuracy (0.988), indicating effective minimization of penalties from misclassifications. The Total Misclassification Cost (345) is reasonable given the dataset size and complexity, while the Weighted F1-Score (0.945) highlights a balanced ability to accurately classify samples across all classes while considering the cost of errors. These metrics collectively suggest the model is well-suited for scenarios where minimizing misclassification costs is critical.

## 5.2 XGBoost

XGBoost is a gradient boosting algorithm that can deliver higher speeds and performance when handling structured and tabular data. Unlike other Random Forest, XGBoost forms the tree one at a time, allowing the new tree to correct the error produced by the previous tree, thereby improving the prediction of the model [16]. XGBoost has built-in regularization techniques such as L-1 / L-2 regularization to prevent overfitting which produces a model that generalizes well on unseen data [17]. XGBoost is also able to train models with large datasets efficiently and quickly due to the utilization of parallel and distributed computing [18].

XGBoost incorporates regularization in its objective function and takes the form

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_j \Omega(f_k), \quad \text{where } \Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$

The first term  $l$  is a differentiable convex loss function that measures the difference between the prediction  $\hat{y}_i$  and the target  $y_i$ . As multiple labels are used to classify the images, we will select  $l$  to be the multi-class cross-entropy loss function. The second term  $\Omega$  penalizes the complexity of the model. Inside it,  $T$  is the number of leaf nodes and  $\|w\|^2$  is the squared sum of the leaf prediction values [19]. There are 2 hyperparameters to prevent overfitting of the model:  $\gamma$  penalizes the number of leaf nodes and  $\lambda$  is the L-2 regularization parameters for the leaf predicted values. When these hyperparameters are set to zero, the objective falls back to the traditional gradient tree boosting.

The model is trained in an additive manner. We let  $\hat{y}_i^{(t)}$  be the prediction of the  $i$ -th instance at the  $t$ -th iteration, and add  $f_t(x_i)$  to minimize the objective function, thereby giving us

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t)$$

It will be difficult for us to solve the objective directly so we apply 2nd order Taylor approximation to it and get

$$\mathcal{L}^{(t)} = \sum_{i=1}^n \left[ l(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i) \right] + \Omega(f_t)$$

where  $g_i = \partial_{\hat{y}_i^{(t-1)}} l(y_i, \hat{y}_i^{(t-1)})$  and  $h_i = \partial_{\hat{y}_i^{(t-1)}}^2 l(y_i, \hat{y}_i^{(t-1)})$ .

Since  $l(y_i, \hat{y}_i^{(t-1)})$  is a constant, we can remove it and simplifying the modified objective further, we arrive at the following equation that will give us our optimal value.

$$\tilde{\mathcal{L}}^{(t)}(q) = -\frac{1}{2} \sum_{j=1}^T \frac{\left( \sum_{i \in I_j} g_i \right)^2}{\sum_{i \in I_j} h_i} + \gamma T$$

$g_i$  and  $h_i$  are known as the first and second order gradient statistics on the loss function. Alternatively, they are called gradient and hessian respectively [20].

### 5.2.1 Training process and tuning

In our project for classifying images of leukemia using XGBoost, the pipeline begins by loading necessary libraries such as xgboost, caret, dplyr, and imager, among others, for tasks like image processing, feature extraction, and model evaluation. The dataset, which consists of images from four classes (Benign, Early, Pre, Pro), is loaded and structured by first gathering the file paths for images of each class and associating labels to them. It ensures that the dataset is organized into a training, validation, and test set (64%, 20%, and 16%, respectively). Class weights are calculated based on the sample sizes of each class, and the dataset is shuffled to ensure randomness.

The feature extraction process is crucial in converting image data into a format suitable for machine learning models. We defined a function, "extract\_features", which takes an image path, processes the image (resizes it, converts to grayscale, extracts color RGB histograms, calculates texture features like contrast, skewness, and kurtosis, and performs edge detection), and computes shape features like area, perimeter and circularity based on image segmentation. Additionally, the function get\_feature\_params is used to analyze the feature lengths across the training data, ensuring consistent extraction. The feature extraction pipeline is validated and adjusted to handle inconsistent feature lengths by padding or truncating features as needed, ensuring that all features have a uniform length.

Once the features are extracted, they are processed and combined with labels to create a final dataset that is ready for model training. Features are extracted for the training, validation, and test datasets, and any "near-zero variance" (NZV) features are removed to improve the model's performance.

The xgboost model is then trained using the filtered dataset, with parameters optimized for multi-class classification. Early stopping is used during training to avoid overfitting.

We will be using xgboost in 'xgboost' package to train our model. This pacakge consists of the following parameters.

1. **nrounds:** the maximum number of iterations
2. **eta:** the learning rate; The higher the eta, the faster the computation
3. **max.depth:** the depth of the tree; The higher the depth of the tree, the higher the chance of overfitting
4. **subsample:** the number of samples (observations) supplied to a tree
5. **colsample\_bytree:** the number of features (variables) supplied to a tree
6. **min\_child\_weight:** minimum weight needed to create a new node in the tree

After training, predictions are made on the test dataset, and a confusion matrix is generated to evaluate the model's performance. The trained model is saved for future use.

### 5.2.2 Test and train results

```

Confusion Matrix and Statistics

Reference
Prediction Benign Early Pre Pro
    Benign    66     2     0     0
    Early     10    153     6     0
    Pre       1      2   144     0
    Pro       3      0     3   128

Overall Statistics

    Accuracy : 0.9479
    95% CI  : (0.9251, 0.9654)
    No Information Rate : 0.3031
    P-Value [Acc > NIR] : < 2.2e-16

    Kappa : 0.9289

    Mcnemar's Test P-Value : NA

    Final Test Accuracy: 0.9478764

    Final Train Accuracy: 0.9947267

```

Figure 22: Confusion Matrix of XGBoost

In our model, we set nrounds = 1000, eta = 0.1, max.depth = 6, subsample = 0.8, colsample\_bytree = 0.8 and min\_child\_weight = 1. This gives us a training accuracy of 0.9947267 and test accuracy of 0.9478764. As both accuracies are close to each other, it suggests that the model is not overfitted. From the confusion matrix in Figure 43, we can deduce the following observations. XGBoost is able to classify all 'Pro' cases correctly. It is able to do sufficiently well in 'Early' and 'Pre' cases. However, it struggles to classify 'Benign' cases correctly, which is evident in the class having the highest number of false positives among the 4. This has serious consequences for both medical practitioners and patients. Patient who have no illness are given medication by a practitioner, who has come out with the wrong diagnosis. This could harm the patient's body in the long run. Legal action could be taken against the practitioner for the wrong diagnosis, resulting in a lose-lose situation for both parties.

```

Class-Specific Statistics:

    Sensitivity Specificity Pos Pred Value Neg Pred Value Precision
Class: Benign  0.8250000  0.9954338  0.9705882  0.9688889  0.9705882
Class: Early   0.9745223  0.9556787  0.9053254  0.9885387  0.9853254
Class: Pre     0.9411765  0.9917808  0.9795918  0.9757412  0.9795918
Class: Pro     1.0000000  0.9846154  0.9552239  1.0000000  0.9552239

    Recall      F1 Prevalence Detection Rate
Class: Benign  0.8250000  0.8918919  0.1544402  0.1274131
Class: Early   0.9745223  0.9386503  0.3030888  0.2953668
Class: Pre     0.9411765  0.9600000  0.2953668  0.2779923
Class: Pro     1.0000000  0.9770992  0.2471042  0.2471042

    Detection Prevalence Balanced Accuracy
Class: Benign   0.1312741  0.9102169
Class: Early    0.3262548  0.9651005
Class: Pre      0.2837838  0.9664786
Class: Pro      0.2586873  0.9923077

```

Figure 23: Overall statistics

Nevertheless, XGBoost appears to handle class imbalance well, with high sensitivity, specificity, and balanced accuracy. From Figure 18, the balanced accuracy, which accounts for class imbalances, is also high, particularly for the "Pro" class (99.23%), and other classes show strong results as well (Early: 96.51%, Pre:

96.65%, Benign: 91.02%). Sensitivity and specificity are generally excellent, with the "Pro" class achieving perfect sensitivity (100%) and high specificity (98.46%), while "Benign" class also performs well with 82.5% sensitivity and 99.54% specificity. Precision is good, particularly for the "Pre" (97.96%) and "Benign" (97.06%) classes, though the "Early" class has slightly lower precision (90.53%) and sensitivity (97.45%), suggesting that some "Early" instances may be misclassified as "Benign". The model also excels in Negative Predictive Value, especially for the "Pro" class (100%). The F1 scores, which balance precision and recall, are also high, with "Pro" achieving the highest score (0.9771).

### 5.2.3 Misclassification cost results

**Cost-Weighted Misclassification Matrix:**

	Benign	Early	Pre	Pro
Benign	0	20	0	0
Early	100	0	90	0
Pre	20	30	0	0
Pro	90	0	30	0

**Total Misclassification Cost: 380**

**Average Cost Per Sample: 0.7335907**

**Cost-Sensitive Accuracy: 0.986662**

**Weighted F1-score 0.9419104**

Figure 24: Misclassification cost results, including the cost weighted confusion matrix

The Total misclassification Cost of this model is 380 which gives the overall financial impact of incorrect predictions, calculated by multiplying the frequency of misclassifications by their respective penalties. The average cost per sample of 0.7335907 normalizes this total cost by the number of samples, providing a clearer picture of the cost per misclassified instance. Cost-sensitive accuracy of 0.986662 adjusts the traditional accuracy by factoring in the severity of errors, highlighting the model's ability to minimize costly misclassifications.

Therefore, even though the Benign class is less prevalent, the model maintains strong performance across all classes. This suggests that XGBoost is well-suited for imbalanced datasets and is handling the class distribution in a balanced way.

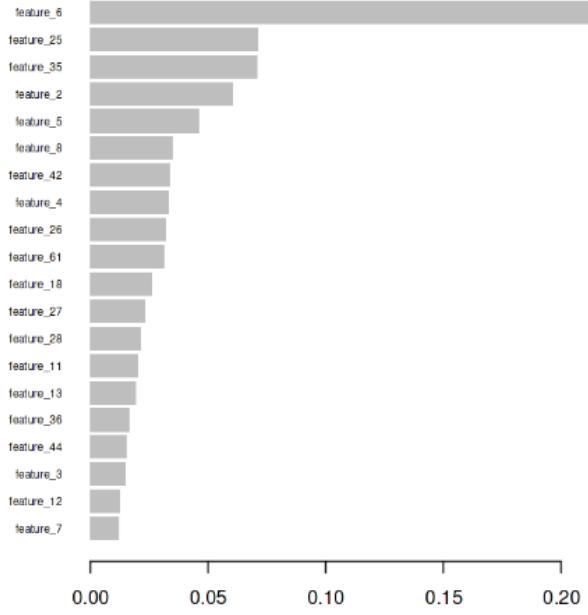


Figure 25: Feature importance of XGBoost

Finally, the top 20 important features are visualized in Figure 45. The top 5 features that are of high importance are feature\\_6, feature\\_25, feature\\_35, feature\\_2 and feature\\_5.

## 6 Convolutional Neural Network

A Convolutional Neural Network (CNN) is a type of deep learning model designed primarily for processing structured grid-like data, such as images. It is inspired by the human visual system, where different layers in the network learn hierarchical patterns and features from raw pixel data [21]. CNNs are particularly effective in image recognition, object detection, and segmentation tasks because they automatically learn and refine the features needed for these tasks during training, making them highly efficient for visual pattern recognition.

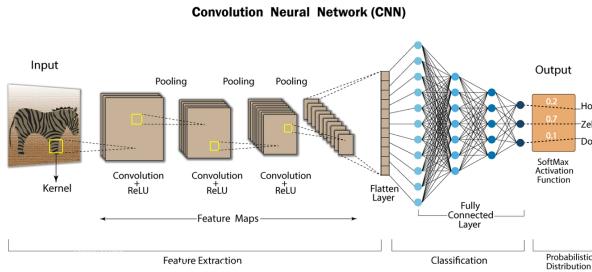


Figure 26: Convolutional Neural Network diagram

A CNN typically consists of multiple layers: convolutional layers, which apply filters to detect various features like edges, textures, and shapes; activation functions, usually ReLU (Rectified Linear Unit), that introduce non-linearity to the network; pooling layers, which downsample feature maps to reduce spatial dimensions and computational load while retaining important information; and fully connected layers as shown in Figure X.X, which interpret the extracted features for classification or other tasks.

## 6.1 Convolutional Layer

A convolutional layer in a Convolutional Neural Network (CNN) is designed to detect specific features in input data by applying a mathematical operation called convolution. In this layer, a small matrix of weights, known as a filter or kernel, slides over the input image (or previous layers output) to produce feature maps that highlight important patterns like edges, textures, or shapes. The convolution operation itself involves element-wise multiplication between the filter and a corresponding section of the input, followed by summing these products. Mathematically, the convolution operation for a given position can be described as:

$$Z(i, j) = \sum_{n=1}^N \sum_{m=1}^M X(i+m, j+n) \times K(m, n)$$

where:

- $Z(i, j)$  is the output at position  $(i, j)$ ,
- $X(i+m, j+n)$  is the pixel value from the input image region,
- $K(m, n)$  is the values of the filter.

The values of  $m$  and  $n$  are determined by the filters dimensions, and the filter is shifted across the input image by a fixed step size known as the stride. This operation captures spatial hierarchies in the input, enabling the network to learn patterns that represent different parts of an object or texture in the image. Multiple filters are usually applied within a convolutional layer, producing different feature maps that capture diverse aspects of the input data, which are then stacked together and passed to subsequent layers. Through training, these filters learn to identify increasingly complex features, making convolutional layers a crucial component for feature extraction in CNNs.

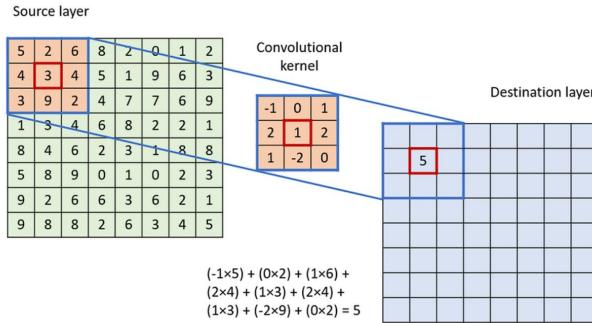


Figure 27: Convolutional operation example

Figure 27 illustrates a two-dimensional convolution operation, where a  $3 \times 3$  kernel (filter) is applied to a  $7 \times 7$  input matrix (source layer) to produce an output matrix (destination layer). At each step, a  $3 \times 3$  sub-matrix is selected from the input matrix, and an element-wise multiplication is performed between the sub-matrix and the kernel. The resulting values are summed to produce a single value, which is placed in the corresponding position of the output matrix. For example, given the sub-matrix

$$\begin{bmatrix} 5 & 2 & 6 \\ 4 & 3 & 4 \\ 3 & 9 & 2 \end{bmatrix}$$

and the kernel

$$\begin{bmatrix} -1 & 0 & 1 \\ 2 & 1 & 2 \\ 1 & -2 & 0 \end{bmatrix},$$

the computation is

$$(-1 \cdot 5) + (0 \cdot 2) + (1 \cdot 6) + (2 \cdot 4) + (1 \cdot 3) + (2 \cdot 4) + (1 \cdot 3) + (-2 \cdot 9) + (0 \cdot 2) = 5.$$

This value is placed in the output matrix. The kernel slides across the input matrix, repeating this process for all valid positions, forming the complete output matrix.

## 6.2 Pooling Layer

A pooling layer in a Convolutional Neural Network (CNN) is a downsampling operation that reduces the spatial dimensions (height and width) of the input feature map while retaining the most essential information. The primary goal of pooling is to decrease the computational load, reduce the number of parameters, and help the network generalize better by making it more invariant to small translations, rotations, or distortions in the input. The most commonly used pooling methods are max pooling and average pooling. In max pooling, a window or filter slides over the feature map and selects the maximum value within each region (typically 2x2 or 3x3).

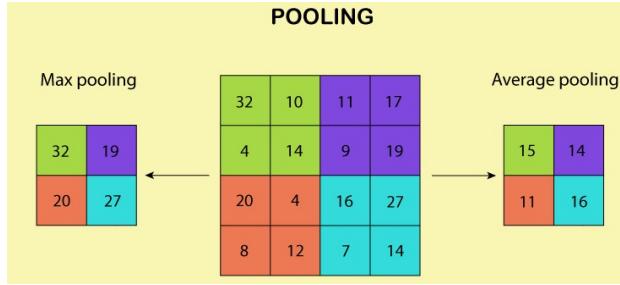


Figure 28: Average and max pooling example

In average pooling, it computes the average value within the same window. By reducing the spatial size, pooling helps make the model more efficient and less prone to overfitting, while also allowing it to focus on the most prominent features. Pooling layers are usually applied after convolutional layers to summarise the extracted features and progressively reduce the dimensionality of the feature maps, thus capturing more abstract patterns as the network deepens.

## 6.3 Proposed Architecture

### 6.3.1 Differences between Research Paper's Architecture and Our Proposed Architecture

The proposed architecture that we are introducing is based on the research paper that was done on the Acute Lymphoblastic Leukemia (ALL) image classification dataset that we are currently using [6]. Our architecture will be slightly different, in which it will incorporate a U-Net-like architecture into the segmentation block and a chosen pretrained model for the feature extraction block. The use of U-Net functionalities is a feature absent in the research paper.

For the classification block, the research paper employs two dense sequential blocks with LeakyReLU and ReLU activations, batch normalization, and dropout, culminating in a dense layer with four neurons and Softmax activation for multi-class classification. In contrast, our architecture includes more regularization with L2 penalties, higher dropout rates (0.7), and directly incorporates segmentation-based features into its dense layers, making the classification block more specialized towards leukemia cell images. Finally, while both approaches uses transfer learning, the research paper focuses on fine-tuning the CNNs, while we combined this with segmentation-informed learning.

Below is a table on the summary of key differences between the research paper's and our proposed architecture:

Aspect	Research paper	Our project
<b>Feature Extraction</b>	10 pretrained CNNs (e.g., DenseNet, ResNet, VGG).	Custom U-Net + 3 pretrained CNNs
<b>Segmentation</b>	Preprocessing step, features combined later.	Integrated into the pipeline with U-Net.
<b>Classification Block</b>	Two dense blocks with specific activations.	Dense layers with segmentation + pretrained CNN features and high dropout.
<b>Regularization</b>	Dropout, LeakyReLU, batch normalization.	L2 regularization, higher dropout rates (0.7).
<b>Final Output</b>	Softmax activation for 4 classes.	Softmax activation for the number of classes in <code>full_df</code> .

Table 1: Summary of key differences

### 6.3.2 Segmentation Block

Our architecture starts off with a purple segmentation layer for cancer cell shape identification based on their distinct purple colouration.

The U-Net architecture is a convolutional neural network originally developed for biomedical image segmentation tasks. It has a unique "U" shape formed by a contracting path and an expansive path. In the contracting path, high-level features are captured by successive convolutional and pooling layers, reducing the spatial dimensions while increasing feature depth.

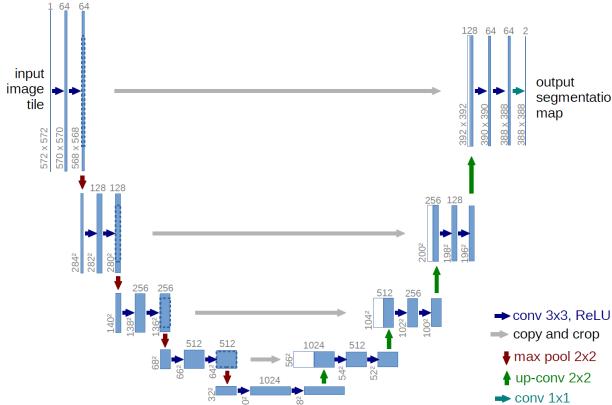


Figure 29: Average and max pooling example

Conversely, the expansive path uses transposed convolutions to upsample these features, reconstructing spatial information while simultaneously combining it with corresponding features from the contracting path through skip connections. These skip connections help retain important spatial context lost during downsampling, which improves the network's ability to accurately segment complex images. This design makes U-Net especially effective in tasks that require precise localization and boundary delineation [12].

### 6.3.3 Architecture Overview

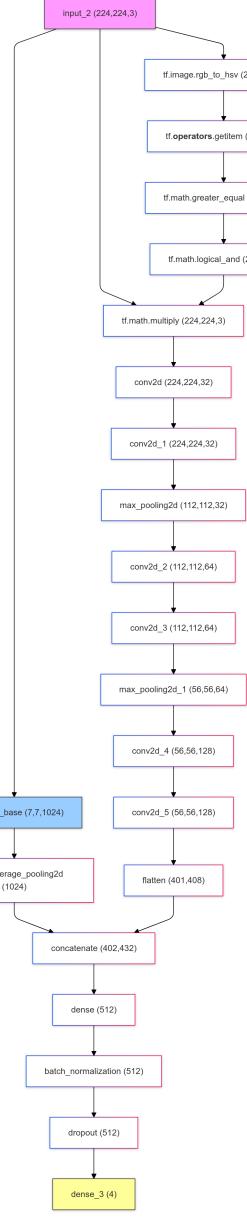


Figure 30: Proposed architecture flowchart, using ConvNeXt as an example for pretrained model

This model architecture integrates a pretrained base model with a U-Net-inspired segmentation block to perform segmentation and classification tasks. The pretrained model, initialized with ImageNet weights and excluding the top classification layers (`include_top = FALSE`), acts as a feature extractor, capturing high-level image features.

The architecture starts with a functional layer that processes an input image to segment purple regions corresponding to lymphoblasts, using a color threshold in the HSV (Hue, Saturation, Value) color space with specified lower and upper bounds for the hue, saturation, and value channels. The input image is then converted from RGB to HSV format using TensorFlow's `rgb_to_hsv` function. A binary mask is created by checking if each pixel in the image falls within the defined purple color range for each of the HSV channels.

This mask is then cast to a float32 type and expanded to match the original image's dimensions. The mask is finally applied to the original image by multiplying the RGB values by the mask, segregating the purple regions from the rest of the image, producing a segmented image where only the purple areas (lymphoblasts) are retained.

The segmentation block is based on the U-Net architecture, widely used in image segmentation. The encoder (contracting path) progressively extracts abstract features through convolutional layers followed by max-pooling. The encoder blocks have 32, 64, and 128 filters in conv1, conv2, and conv3, respectively, with max-pooling reducing spatial resolution. At the bottleneck, two convolutional layers with 256 filters learn abstract features before the decoder begins. The decoder (expansive path) upsamples features using transposed convolutional layers and incorporates skip connections from corresponding encoder layers to retain spatial information. Filters in the decoder decrease from 128 to 64 and 32, refining features at each step and reconstructing the spatial resolution for pixel-wise segmentation.

The pretrained base model also extracts features passed through a global average pooling layer, reducing spatial dimensions to a single feature vector. This vector is flattened and concatenated with the segmentation block's output, combining segmentation and pretrained features. The combined features are processed by the classification block, which consists of fully connected layers with 512, 256, and 128 units, each followed by batch normalization, dropout (0.7), and L2 regularization. The final fully connected layer is a softmax layer with four units (Benign, Early, Pre, Pro), producing class probabilities.

This architecture will be used consistently across all pretrained models in this project.

## 6.4 VGG16 Architecture

### 6.4.1 Model description

VGG16 is a popular CNN pre-trained model created by Visual Geometry Group (VGG) primarily used for image classification. It has a total of 16 layers, 13 of which are convolutional and the remaining 3 are fully connected. It utilises 3x3 convolutional filters throughout the network and for down-sampling, has 2x2 max-pooling layers. Our CNN model is built on this and we use a transfer learning approach using VGG16 with the pretrained weights from the Imagenet dataset [22].

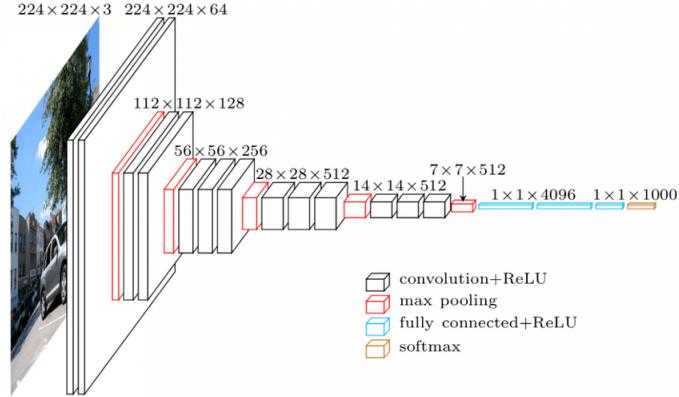


Figure 31: VGG16 architecture diagram

#### 6.4.2 Training process and tuning

Layer (Type)	Output Shape	Param #	Connected to	Trainable
input_2 (InputLayer)	(None, 224, 224, 3)	0	-	Y
tf.image.rgb_to_hsv (TFOpLambda)	(None, 224, 224, 3)	0	input_2[0][0]	Y
tf.__operators__.getitem (SlicingOpLambda)	(None, 224, 224)	0	tf.image.rgb_to_hsv[0][0]	Y
conv2d (Conv2D)	(None, 224, 224, 32)	896	tf.math.multiply[0][0]	Y
conv2d_1 (Conv2D)	(None, 224, 224, 32)	9248	conv2d[0][0]	Y
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0	conv2d_1[0][0]	Y
conv2d_2 (Conv2D)	(None, 112, 112, 64)	18496	max_pooling2d[0][0]	Y
conv2d_3 (Conv2D)	(None, 112, 112, 64)	36928	conv2d_2[0][0]	Y
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0	conv2d_3[0][0]	Y
conv2d_4 (Conv2D)	(None, 56, 56, 128)	73856	max_pooling2d_1[0][0]	Y
conv2d_5 (Conv2D)	(None, 56, 56, 128)	147584	conv2d_4[0][0]	Y
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 128)	0	conv2d_5[0][0]	Y
conv2d_6 (Conv2D)	(None, 28, 28, 256)	295168	max_pooling2d_2[0][0]	Y
conv2d_7 (Conv2D)	(None, 28, 28, 256)	590080	conv2d_6[0][0]	Y
conv2d_transpose (Conv2DTranspose)	(None, 56, 56, 128)	295040	conv2d_7[0][0]	Y
concatenate (Concatenate)	(None, 56, 56, 256)	0	conv2d_transpose[0][0], conv2d_5[0][0]	Y
dense (Dense)	(None, 512)	205783552	concatenate_3[0][0]	Y
batch_normalization (BatchNormalization)	(None, 512)	2048	dense[0][0]	Y
dropout (Dropout)	(None, 512)	0	batch_normalization[0][0]	Y
dense_3 (Dense)	(None, 4)	516	dropout_2[0][0]	Y

Table 2: Model architecture using VGG16 as pre-trained model

The training process begins by organizing the dataset for model training, starting with the preparation of file paths and corresponding class labels. The dataset is split into training (64%), validation (20%), and test (16%) sets using stratified sampling to ensure each class is proportionally represented. After the initial split, the training set (`train_df`) is further shuffled randomly. Hyperparameter tuning focuses on optimizing key parameters for model performance. The initial learning rate is set to  $1 \times 10^{-5}$  for stable convergence. Dense layers are regularized with  $L_2$  regularization (0.01) and dropout (0.7) to prevent overfitting. Early stopping mechanism is also implemented with a patience of 8, and using a learning rate plateau scheduler of patience 6, which reduces the learning rate by a factor of 0.1 if the validation loss stalls for 6 epochs. These callbacks help optimize training and prevent overfitting.

Data augmentation, including rotations, shifts, and zooms, is applied only to the training set. Class weights are adjusted based on class distribution to address imbalance, using the same formula from section 3.6. The model is trained for 220 epochs and evaluated using categorical cross-entropy loss and accuracy.

#### 6.4.3 Test and train results

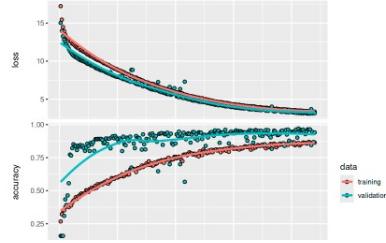


Figure 32: VGG16 loss and accuracy plot

The plots show training (red) and validation (blue) performance over epochs. Training loss decreases steadily, while validation loss stabilizes, indicating limited improvement in generalization. Interestingly, validation accuracy surpasses training accuracy at times, likely due to regularization techniques like dropout or data augmentation. While the model generalizes relatively well, the plateau in validation accuracy suggests potential for further fine-tuning.

Confusion Matrix and Statistics						
		Reference				
		Prediction	Benign	Early	Pre	Pro
Benign		68	0	0	0	
Early		8	156	10	0	
Pre		0	0	142	0	
Pro		4	1	1	128	
Overall Statistics						
Accuracy : 0.9537						
95% CI : (0.9318, 0.9701)						
No Information Rate : 0.3031						
P-Value [Acc > NIR] : < 2.2e-16						
Kappa : 0.9368						
McNemar's Test P-Value : NA						
Final Training Loss: 3.34835						
Final Validation Loss: 3.34835						
Final Train Accuracy: 0.8636806						
Final Validation Accuracy: 0.940625						
Test Accuracy: 0.953125						

Figure 33: Confusion Matrix of VGG16

The VGG16 model has demonstrated strong performance on this dataset, with a test accuracy of 95.31% and a test loss of 3.088493, indicating good generalization to unseen data. The confusion matrix shows that the model has correctly classified a majority of the images, particularly in the "Benign" and "Pro" classes, achieving perfect specificity in both. However, some misclassifications are observed in the "Early" and "Pre" classes, with "Early" images being misclassified as "Benign" and "Pre" images being misclassified as "Early."

The test accuracy (95.37%) is also slightly higher than the train accuracy (86.36%), which shows that there is no signs of overfitting.

Class-Specific Statistics:							
	Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision		
Class: Benign	0.8500000	1.0000000	1.0000000	0.9733333	1.0000000		
Class: Early	0.9936306	0.9501385	0.8965517	0.9970930	0.8965517		
Class: Pre	0.9281846	1.0000000	1.0000000	0.9707447	1.0000000		
Class: Pro	1.0000000	0.9846154	0.9552239	1.0000000	0.9552239		
	Recall	F1	Prevalence	Detection Rate			
Class: Benign	0.8500000	0.9189189	0.1544402	0.1312741			
Class: Early	0.9936306	0.9425982	0.3830888	0.3011583			
Class: Pre	0.9281846	0.9627119	0.2953668	0.2741313			
Class: Pro	1.0000000	0.9778992	0.2471042	0.2471042			
	Detection	Prevalence	Balanced Accuracy				
Class: Benign		0.1312741	0.9250000				
Class: Early		0.3359973	0.9718845				
Class: Pre		0.2741313	0.9640523				
Class: Pro		0.2586673	0.9923077				

Figure 34: VGG16 class specific statistics

The "Benign" class has slightly lower sensitivity (85%) but perfect precision, meaning the model is accurate when predicting Benign but may miss some true Benign cases. The "Early" class has the highest sensitivity (99.36%) but lower precision (89.66%), indicating that the model may falsely predict "Early" in some cases. For the "Pre" class, both sensitivity and specificity are good, and the "Pro" class shows excellent performance with 100% sensitivity and strong precision (95.52%). The model's F1 scores reflect these findings, with the "Pro" class achieving the highest F1 score (97.71%). Thus, with the balanced accuracy, the model is performing exceptionally well in terms of handling class imbalances, with only minor misclassifications occurring between certain classes.

#### 6.4.4 Misclassification cost

Cost-Weighted Misclassification Matrix:

	Benign	Early	Pre	Pro
Benign	0	0	0	0
Early	80	0	150	0
Pre	0	0	0	0
Pro	120	25	10	0

Total Misclassification Cost: 385

Average Cost Per Sample: 0.7432432

Cost-Sensitive Accuracy: 0.9864865

Weighted F1-score 0.9503321

Figure 35: VGG16 misclassification cost

The cost-weighted misclassification results highlight the model's strong performance, with a cost-sensitive accuracy of 98.65% and a weighted F1-score of 0.9503. The total misclassification cost of 385 and an average cost per sample of 0.743 reflect the impact of high-cost errors, particularly in misclassifying "Early" and "Pro" classes. While the model effectively minimizes costly mistakes, further fine-tuning with cost-sensitive methods could further reduce the misclassification impact.

## 6.5 DenseNet-201

### 6.5.1 Model description

DenseNet-201 is a deep convolutional neural network architecture that builds on the concept of dense connections between layers. Unlike traditional networks, where each layer receives input from only the previous one, DenseNet connects each layer to every other layer in a feed-forward manner. This results in a highly efficient network where each layer has direct access to the gradients from the loss function, enhancing the flow of information and reducing the risk of vanishing gradients. DenseNet-201, specifically, is a version of DenseNet with 201 layers, designed to perform exceptionally well in image classification tasks by leveraging its dense connectivity to improve feature reuse and reduce the number of parameters compared to traditional deep networks [23].

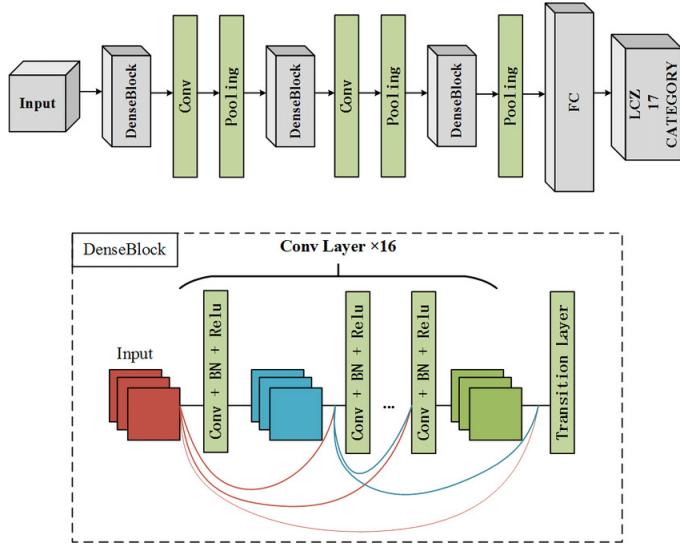


Figure 36: DenseNet-201 diagram

Each dense block is composed of multiple convolutional layers, where each layer receives input from all previous layers within the same block, as well as outputs feature maps that are passed to the subsequent layers. The key feature of DenseNet-201 is its "dense connectivity" pattern, where each layer is connected to every other layer in a feed-forward fashion. This ensures that each layer has access to all previously learned features, enhancing feature reuse and improving gradient flow during training. Between dense blocks, transition layers are used to reduce the feature map size and the number of channels, helping to control the computational complexity. DenseNet-201 is typically characterized by having fewer parameters than traditional deep networks because of the efficient feature reuse. It performs particularly well in tasks like image classification and object recognition by leveraging the dense connections to extract rich, multi-level features from input data.

### 6.5.2 Training process and tuning

The splitting and shuffling of dataset remains the same as all the other models in this project. For hyperparameter tuning, the learning rate is set to  $1 \times 10^{-4}$  for a quicker convergence. Dense layers are regularized with  $L_2$  regularization (0.01) and dropout (0.7) to prevent overfitting. Unlike VGG16 which requires a higher patience for stable convergence, DenseNet-201 is able to train smoothly with a early stopping patience of 5 and using a learning rate scheduler of patience 3, which reduces the learning rate by a factor of 0.1 if the validation loss stalls for 3 epochs. Data augmentation, including rotations, shifts, and zooms, are similarly only applied to the training set. The model is also trained for 220 epochs and evaluated using categorical cross-entropy loss and accuracy.

Layer (type)	Output Shape	Param #	Connected to	Trainable
input_2 (InputLayer)	(None, 224, 224, 3)	0	[]	Y
tf.image.rgb_to_hsv (TFOpLambda)	(None, 224, 224, 3)	0	[tf.image.rgb_to_hsv[0][0]]	Y
tf._operators._get_item (SlicingOpLambda)	(None, 224, 224)	0	[tf.image.rgb_to_hsv[0][0]]	Y
tf._operators._get_item_1 (SlicingOpLambda)	(None, 224, 224)	0	[tf.image.rgb_to_hsv[0][0]]	Y
tf.math.greater_equal (TFOpLambda)	(None, 224, 224)	0	[tf._operators._get_item[0][0]]	Y
tf.math.less_equal (TFOpLambda)	(None, 224, 224)	0	[tf._operators._get_item_1[0][0]]	Y
tf._operators._get_item_2 (SlicingOpLambda)	(None, 224, 224)	0	[tf.image.rgb_to_hsv[0][0]]	Y
tf.math.logical_and (TFOpLambda)	(None, 224, 224)	0	[tf.math.greater_equal[0][0], tf.math.less_equal[0][0]]	Y
tf.math.greater_equal_1 (TFOpLambda)	(None, 224, 224)	0	[tf._operators._get_item_2[0][0]]	Y
tf._operators._get_item_3 (SlicingOpLambda)	(None, 224, 224)	0	[tf.image.rgb_to_hsv[0][0]]	Y
tf.math.logical_and_1 (TFOpLambda)	(None, 224, 224)	0	[tf.math.logical_and[0][0], tf.math.greater_equal_1[0][0]]	Y
tf.math.greater_equal_2 (TFOpLambda)	(None, 224, 224)	0	[tf._operators._get_item_3[0][0]]	Y
tf.math.logical_and_2 (TFOpLambda)	(None, 224, 224)	0	[tf.math.logical_and_1[0][0], tf.math.greater_equal_2[0][0]]	Y
tf.cast (TFOpLambda)	(None, 224, 224)	0	[tf.math.logical_and_2[0][0]]	Y
tf.expand_dims (TFOpLambda)	(None, 224, 224, 1)	0	[tf.cast[0][0]]	Y
tf.tile (TFOpLambda)	(None, 224, 224, 3)	0	[tf.expand_dims[0][0]]	Y
tf.math.multiply (TFOpLambda)	(None, 224, 224, 3)	0	[input_2[0][0], tf.tile[0][0]]	Y
conv2d (Conv2D)	(None, 224, 224, 32)	896	[tf.math.multiply[0][0]]	Y
conv2d_1 (Conv2D)	(None, 224, 224, 32)	9248	[conv2d[0][0]]	Y
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0	[conv2d_1[0][0]]	Y
conv2d_2 (Conv2D)	(None, 112, 112, 64)	18496	[max_pooling2d[0][0]]	Y
conv2d_3 (Conv2D)	(None, 112, 112, 64)	36928	[conv2d_2[0][0]]	Y
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0	[conv2d_3[0][0]]	Y
conv2d_4 (Conv2D)	(None, 56, 56, 128)	73856	[max_pooling2d_1[0][0]]	Y
conv2d_5 (Conv2D)	(None, 56, 56, 128)	147584	[conv2d_4[0][0]]	Y
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 128)	0	[conv2d_5[0][0]]	Y
conv2d_6 (Conv2D)	(None, 28, 28, 256)	295168	[max_pooling2d_2[0][0]]	Y
conv2d_7 (Conv2D)	(None, 28, 28, 256)	590080	[conv2d_6[0][0]]	Y
conv2d_transpose (Conv2DTranspose)	(None, 56, 56, 128)	295040	[conv2d_7[0][0]]	Y
concatenate (Concatenate)	(None, 56, 56, 256)	0	[conv2d_transpose[0][0], conv2d_5[0][0]]	Y
conv2d_8 (Conv2D)	(None, 56, 56, 128)	295040	[concatenate[0][0]]	Y
densenet201 (Functional)	(None, 7, 7, 1920)	1832198	[input_2[0][0]]	Y
conv2d_9 (Conv2D)	(None, 56, 56, 128)	147584	[conv2d_8[0][0]]	Y
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1920)	0	[densenet201[0][0]]	Y
flatten (Flatten)	(None, 401408)	0	[conv2d_9[0][0]]	Y
concatenate_3 (Concatenate)	(None, 403328)	0	[global_average_pooling2d[0][0], flatten[0][0]]	Y
dense (Dense)	(None, 512)	2065044	[concatenate_3[0][0]]	Y
batch_normalization (BatchNormalization)	(None, 512)	2048	[dense[0][0]]	Y
dropout (Dropout)	(None, 512)	0	[batch_normalization[0][0]]	Y
dense_1 (Dense)	(None, 256)	131328	[dropout[0][0]]	Y
batch_normalization_1 (BatchNormalization)	(None, 256)	1024	[dense_1[0][0]]	Y
dropout_1 (Dropout)	(None, 256)	0	[batch_normalization_1[0][0]]	Y
dense_2 (Dense)	(None, 1)	257	[dropout_1[0][0]]	Y

Table 3: Model architecture using DenseNet-201

### 6.5.3 Test and train results

In the loss plot below (Figure 37), both training and validation losses decrease sharply at the beginning and eventually plateau at low values, indicating that the model is learning effectively and generalizing well without significant overfitting. In the accuracy plot, training and validation accuracy both increase rapidly and converge near 1.0, demonstrating strong model performance on both the training and validation datasets. The close alignment of the two curves suggests that the model is not overfitting and generalizes well to unseen data. This is an ideal outcome, reflecting a well-trained model.

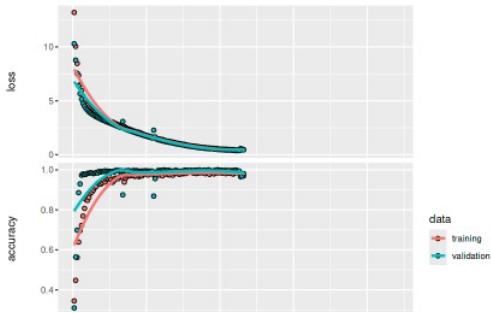


Figure 37: DenseNet201 loss and accuracy plot

Confusion Matrix and Statistics

		Reference				
		Prediction	Benign	Early	Pre	Pro
Prediction	Benign	79	0	0	0	
	Early	1	157	0	0	
	Pre	0	0	152	0	
	Pro	0	0	1	128	

Overall Statistics

Accuracy : 0.9961

95% CI : (0.9861, 0.9995)

No Information Rate : 0.3031

P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.9948

Mcnemar's Test P-Value : NA

Final Training Loss: 0.4692059

Final Validation Loss: 0.4692059

Final Train Accuracy: 0.969815

Final Validation Accuracy: 0.98125

Test Accuracy: 0.9960938

Figure 38: Confusion Matrix of DenseNet-201

The results in Figure 38 show exceptional performance in terms of classification accuracy and precision. The test accuracy of 99.61% suggests that the model is effectively distinguishing between classes, with minimal misclassification, as evidenced by the confusion matrix, which shows only a few false positives and false negatives. The overall test accuracy (0.9961) is significantly higher than the "No Information Rate" of 0.3031, indicating the model's strong predictive power. The kappa score of 0.9948 indicates excellent agreement between predicted and actual values.

Class-Specific Statistics:						
	Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision	
Class: Benign	0.9875000	1.0000000	1.0000000	0.9977221	1.0000000	
Class: Early	1.0000000	0.9972299	0.9936709	1.0000000	0.9936709	
Class: Pre	0.9934641	1.0000000	1.0000000	0.9972678	1.0000000	
Class: Pro	1.0000000	0.9974359	0.9922481	1.0000000	0.9922481	
	Recall	F1	Prevalence	Detection Rate		
Class: Benign	0.9875000	0.9937107	0.1544402	0.1525097		
Class: Early	1.0000000	0.9968254	0.3030888	0.3030888		
Class: Pre	0.9934641	0.9967213	0.2953668	0.2934363		
Class: Pro	1.0000000	0.9961089	0.2471042	0.2471042		
	Detection	Prevalence	Balanced	Accuracy		
Class: Benign		0.1525097		0.9937500		
Class: Early		0.3030193		0.9986150		
Class: Pre		0.2934363		0.9967320		
Class: Pro		0.2490347		0.9987179		

Figure 39: DenseNet201 class-specific statistics

Class-specific metrics further highlight the model's efficiency, with high sensitivity, specificity, and precision across all classes, particularly in detecting "Early" and "Pro" stages, where it achieves 100% sensitivity. The test loss of 0.3849 is relatively low, indicating good model generalization. However, while the model performs exceptionally well in validation, the training and validation losses and accuracies suggest that further tuning may be beneficial for improved generalization.

#### 6.5.4 Misclassification Cost

Cost-Weighted Misclassification Matrix:

	Benign	Early	Pre	Pro
Benign	0	0	0	0
Early	10	0	0	0
Pre	0	0	0	0
Pro	0	0	10	0

Total Misclassification Cost: 20

Average Cost Per Sample: 0.03861004

Cost-Sensitive Accuracy: 0.999298

Weighted F1-score 0.9958416

Figure 40: DenseNet201 misclassification costs

The misclassification costs of DenseNet-201, as indicated by the cost-weighted matrix, show a total misclassification cost of 20, which is relatively low considering the dataset's size. The misclassifications primarily involve "Early" and "Pro" classes, where "Early" samples were misclassified as "Benign" (10 times), and "Pro" samples were misclassified as "Pre" (10 times). These misclassifications contribute to a small cost in the overall performance. Despite these errors, the model maintains a high cost-sensitive accuracy of 99.93%, indicating that the misclassifications are rare and have minimal impact on the model's overall ability to make correct predictions. The weighted F1-score of 0.9958 reflects the model's ability to balance precision and recall well, even in the presence of some misclassifications. The average cost per sample of 0.0386 further emphasizes that the financial or operational impact of these misclassifications is minimal, though this is much lesser compared to VGG16.

## 6.6 ConvNeXT

### 6.6.1 Model description

ConvNeXT (Convolutional Neural Network Transformer) is a recent CNN architecture inspired by vision transformer-based models. It was designed to enhance the performance of CNNs by incorporating design principles by transformers, such as a larger receptive field and more global context. ConvNeXT optimizes CNNs by using techniques like normalized convolutional layers and transformer-inspired block structures while maintaining computational efficiency. This hybrid approach allows ConvNeXT to achieve state-of-the-art performance on image classification tasks, leveraging both the local feature extraction power of CNNs and the long-range dependency modeling of transformers [24].

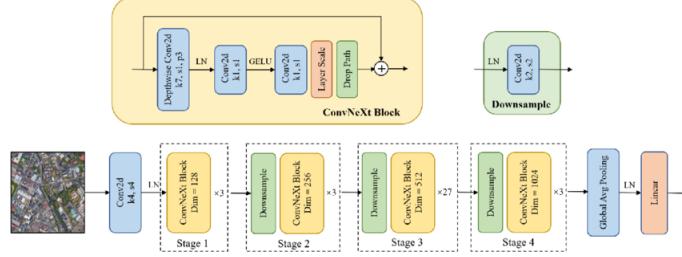


Figure 41: ConvNeXT architecture diagram

### 6.6.2 Training process and tuning

Layer (type)	Output Shape	Param #	Connected to	Trainable
input_2 (InputLayer)	(None, 224, 224, 3)	0	[]	Yes
tf.image.rgb_to_hsv (TFOpLambda)	(None, 224, 224, 3)	0	input_2[0][0]	Yes
tf._operators_.getitem	(None, 224, 224)	0	tf.image.rgb_to_hsv[0][0]	Yes
tf.math.greater_equal (TFOpLambda)	(None, 224, 224)	0	tf._operators_.getitem[0][0]	Yes
tf.math.logical_and (TFOpLambda)	(None, 224, 224)	0	tf.math.greater_equal[0][0]	Yes
tf.math.multiply (TFOpLambda)	(None, 224, 224, 3)	0	input_2[0][0], tf.tile[0][0]	Yes
conv2d (Conv2D)	(None, 224, 224, 32)	896	tf.math.multiply[0][0]	Yes
conv2d_1 (Conv2D)	(None, 224, 224, 32)	9,248	conv2d[0][0]	Yes
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0	conv2d_1[0][0]	Yes
conv2d_2 (Conv2D)	(None, 112, 112, 64)	18,496	max_pooling2d[0][0]	Yes
conv2d_3 (Conv2D)	(None, 112, 112, 64)	36,928	conv2d_2[0][0]	Yes
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0	conv2d_3[0][0]	Yes
conv2d_4 (Conv2D)	(None, 56, 56, 128)	73,856	max_pooling2d_1[0][0]	Yes
conv2d_5 (Conv2D)	(None, 56, 56, 128)	147,584	conv2d_4[0][0]	Yes
convnext_base (Functional)	(None, 7, 7, 1024)	87,566,464	input_2[0][0]	Yes
global_average_pooling2d	(None, 1024)	0	convnext_base[0][0]	Yes
flatten (Flatten)	(None, 401,408)	0	conv2d_9[0][0]	Yes
concatenate	(None, 402,432)	0	global_average_pooling2d, flatten	Yes
dense (Dense)	(None, 512)	206,045,696	concatenate[0][0]	Yes
batch_normalization	(None, 512)	2,048	dense[0][0]	Yes
dropout	(None, 512)	0	batch_normalization[0][0]	Yes
dense_3 (Dense)	(None, 4)	516	dropout[0][0]	Yes

Table 4: Model architecture using ConvNeXT

The splitting and shuffling of dataset remains the same as all the other models in this project. For hyperparameter tuning, initial learning rate is  $1 \times 10^{-4}$  for a quicker convergence. Layers are similarly regularized with  $L_2$  regularization (0.01) and dropout (0.7) to prevent overfitting. For ConvNeXT, an early stopping patience of 8 and using a learning rate scheduler of patience 5 is set. The values set for each augmentation property, including rotations, shifts, and zooms, remains the same. The model is also trained for 220 epochs and evaluated using categorical cross-entropy loss and accuracy.

### 6.6.3 Test and train results

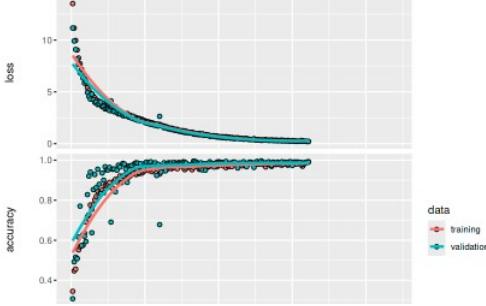


Figure 42: ConvNeXt loss and accuracy plot

The training and validation loss/accuracy plots for the ConvNeXt model reveal an effective and well-optimized training process. The training loss steadily decreases as epochs progress, while training accuracy rises smoothly, converging close to 1 (100%), indicating that the model learns the patterns in the training data effectively. Similarly, the validation loss shows a significant decline during the initial stages of training and stabilizes at a low value, while the validation accuracy increases and converges near 1. The close alignment between the training and validation curves for both loss and accuracy indicates that the model generalizes well to unseen data, with minimal signs of overfitting or underfitting.

```

Test Loss: 0.2232956
Test Accuracy: 0.9941406
Confusion Matrix and Statistics
Reference
Prediction Benign Early Pre Pro
Benign      80      0      0      0
Early        0    156      1      0
Pre          0      1   151      0
Pro          0      0      1  128
Overall Statistics
Accuracy : 0.9942
95% CI  : (0.9832, 0.9988)
No Information Rate : 0.3031
P-Value [Acc > NIR] : < 2.2e-16
Kappa : 0.9921
McNemar's Test P-Value : NA
Final Training Loss: 0.2005041
Final Validation Loss: 0.2005041
Final Train Accuracy: 0.9926972
Final Validation Accuracy: 0.990625

```

Figure 43: Confusion Matrix of ConvNeXt

The ConvNeXt model demonstrates outstanding performance, achieving a test accuracy of 99.41% with a test loss of 0.2233, indicating minimal errors on unseen data. The confusion matrix highlights that the model classifies most instances correctly across all four classes, with only a few misclassifications: one "Early" instance predicted as "Pre," one "Pre" instance predicted as "Early," and another as "Pro".

Class-Specific Statistics:						
	Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision	
Class: Benign	1.000000	1.000000	1.000000	1.000000	1.000000	
Class: Early	0.9936306	0.9972299	0.9936306	0.9972299	0.9936306	
Class: Pre	0.9869281	0.9972603	0.9934211	0.9945355	0.9934211	
Class: Pro	1.000000	0.9974359	0.9922481	1.000000	0.9922481	
	Recall	F1	Prevalence	Detection Rate		
Class: Benign	1.000000	1.000000	0.1544402	0.1544402		
Class: Early	0.9936306	0.9936306	0.3030888	0.3011583		
Class: Pre	0.9869281	0.9901639	0.2953668	0.2915058		
Class: Pro	1.000000	0.9961089	0.2471042	0.2471042		
	Detection	Prevalence	Balanced Accuracy			
Class: Benign	0.1544402	1.000000				
Class: Early	0.3030888	0.9954302				
Class: Pre	0.2934363	0.9920942				
Class: Pro	0.2490347	0.9987179				

Figure 44: ConvNeXt class-specific statistics

Class-specific metrics reveal excellent balance across all categories, indicating that the class weighting formula is able to work well with ConvNeXt model in rebalancing the dataset. Sensitivity and specificity values are close to 1 for each class, affirming the model's ability to identify true positives and avoid false positives with minimal bias. Precision values for all classes are similarly high, leading to F1-scores nearing 1, which confirms the models strong and consistent performance in handling the multi-class classification task. Despite the exceptional results, minor misclassifications between "Early" and "Pre" may indicate feature overlap, warranting further exploration.

#### 6.6.4 Misclassification Costs

Cost-Weighted Misclassification Matrix:

	Benign	Early	Pre	Pro
Benign	0	0	0	0
Early	0	0	15	0
Pre	0	15	0	0
Pro	0	0	10	0

Total Misclassification Cost: 40

Average Cost Per Sample: 0.07722008

Cost-Sensitive Accuracy: 0.998596

Weighted F1-score 0.9949759

Figure 45: ConvNeXt misclassification costs

The total misclassification cost, calculated by summing these values, is 40, which is slightly higher than that of the DenseNet-201 model (20) and much lower than VGG16 model (385). The average cost per sample is derived by dividing the total cost by the number of samples, yielding 0.077, indicating that, on average, each sample contributes a relatively low penalty for misclassification.

The cost-sensitive accuracy, which adjusts the traditional accuracy by factoring in the misclassification costs, is 0.9986. This indicates that the model performs exceptionally well in minimizing the cost impact of its errors, compared to VGG16 model, and slightly worse than DenseNet-201.

## 7 Results and Findings

Below is a summary of the train accuracy, test accuracy, and their respective misclassification cost metrics:

Model	Train Accuracy	Test Accuracy	Total Misclassification Cost	Average Cost Per Sample	Cost-Sensitive Accuracy	Weighted F1-Score
SVM	0.9971	0.9517	345	0.6660	0.9879	0.9448
XGBoost	0.9947	0.9479	380	0.7336	0.9867	0.9419
VGG16	0.8637	0.9531	385	0.7432	0.9865	0.9503
DenseNet-201	0.9698	0.9961	20	0.0386	0.9993	0.9958
ConvNeXt	0.9927	0.9941	40	0.0772	0.9986	0.9950

Table 5: Performance Metrics of Different Models

From the table, SVM model, with its training accuracy of 99.71% and test accuracy of 95.17%, provides a strong fit for the training data but suffers from a relatively high total misclassification cost (345). This suggests that while SVM can generalize well to some extent, it struggles with certain subtypes, particularly the "Early" and "Pre" classes, as reflected by its cost-sensitive accuracy of 98.79% and weighted F1-score of 94.48%. The higher misclassification cost in SVM indicates that, in a clinical setting, the cost of incorrect diagnosis could be significant, especially in cases where early-stage detection is critical.

XGBoost performs similarly well with a test accuracy of 94.79%, but its total misclassification cost (380) is higher than SVM, indicating that while it also provides a decent overall performance, it may be prone to misclassifying the "Pre" and "Early" classes. The cost-sensitive accuracy (98.67%) and weighted F1-score (94.19%) are lower than those of SVM, further suggesting that XGBoost struggles more with the delicate balance required in accurately diagnosing early and pre-leukemia stages. VGG16, despite its high test accuracy of 95.31%, shows lower training accuracy (86.37%) and a significantly higher misclassification cost (385). This higher cost, coupled with an average cost per sample of 0.74, suggests that VGG16's performance could be problematic in real-world scenarios where precise and cost-efficient classification is paramount.

DenseNet-201, on the other hand, outperforms the other models with a test accuracy of 99.61%, the lowest total misclassification cost (20), and the lowest average cost per sample (0.04). With a cost-sensitive accuracy of 99.93% and a weighted F1-score of 99.58%, DenseNet-201 stands out as the most efficient and reliable model for leukemia classification, particularly when minimizing the risks of misclassification is crucial in clinical decision-making. ConvNeXt, though not as efficient as DenseNet-201, offers a similarly high test accuracy of 99.41%, with a moderate misclassification cost (40) and a strong cost-sensitive accuracy of 99.86%. While ConvNeXt provides excellent results, its slightly higher misclassification cost indicates that DenseNet-201 is more optimal for real-world applications where accurate and cost-effective classification of leukemia subtypes is critical to patient care.

## 8 Conclusion

In conclusion, DenseNet-201 stands out as the most reliable model for leukemia subtype classification, offering the highest test accuracy, the lowest misclassification cost, and best cost-sensitive metrics. This makes it the most suitable for clinical applications where precision in identifying leukemia stages is critical. However, while DenseNet-201 outperforms other models, it is important to note that even the best models exhibit trade-offs in terms of misclassification, particularly for early-stage subtypes.

For the limitations of the methods used, deep learning models, such as DenseNet-201 and ConvNeXt, are often considered "black-box" models. Their complex architectures make it difficult for clinicians to understand the reasoning behind the classification decisions, which could hinder trust and adoption in medical practice. In healthcare, interpretability is crucial for ensuring clinical decisions are transparent and justifiable.

Furthermore, the misclassification cost metrics used in our project were based on our subjective judgment, which means they were designed according to our perceived severity of misclassifications rather than a rigorous, scientifically grounded framework. While this approach aligns with our understanding of the potential clinical impact of misclassifying leukemia subtypes, it lacks objective validation or a consensus from medical experts regarding the relative costs of these errors. Therefore, consultation with medical professionals to define the misclassification costs based on clinical significance, or a data-driven method would help give more accurate and corroborated misclassification cost metrics.

Also, there is a limitation of using such dataset for data mining and deep learning purposes. The dataset consists of 3,256 samples which is small relative to the number of skin cancer cases worldwide. This might mean that the results shown in this report might not be indicative of its actual performance in real life. Further research in this area of study is needed to ensure that the risks of misdiagnosis are mitigated.

## References

- [1] Nan Zhang, Jinxian Wu, Qian Wang, Yuxing Liang, Xinqi Li, Guopeng Chen, Linlu Ma, Xiaoyan Liu, and Fuling Zhou. Global burden of hematologic malignancies and evolution patterns over the past 30 years. *Blood Cancer Journal*, 13(1):82, 2023.
- [2] World Health Organization et al. Early cancer diagnosis saves lives, cuts treatment costs. *World Health Organization*, 2017.
- [3] Tibor Mezei, Melinda Kolcsár, András Joó, and Simona Gurzu. Image analysis in histopathology and cytopathology: From early days to current perspectives. *Journal of Imaging*, 10(10):252, 2024.
- [4] Ahmed J Abougarair, M Alshaibi, Amany K Alarbish, Mohammed Ali Qasem, Doaa Abdo, Othman Qasem, Fursan Thabit, and Ozgu Can. Blood cells cancer detection based on deep learning. *Journal of Advances in Artificial Intelligence*, 2(1):108–121, 2024.
- [5] Amjad Rehman, Naveed Abbas, Tanzila Saba, Syed Ijaz ur Rahman, Zahid Mehmood, and Hoshang Kolivand. Classification of acute lymphoblastic leukemia using deep learning. *Microscopy Research and Technique*, 81(11):1310–1317, 2018.
- [6] Mustafa Ghaderzadeh, Mehrad Aria, Azamossadat Hosseini, Farkhondeh Asadi, Davood Bashash, and Hassan Abolghasemi. A fast and efficient cnn model for b-cell diagnosis and its subtypes classification using peripheral blood smear images. *International Journal of Intelligent Systems*, 37(8):5113–5133, 2022.
- [7] Mehmet Erten, Prabal Datta Barua, Sengul Dogan, Turker Tuncer, Ru-San Tan, and UR Acharya. Concatnext: An automated blood cell classification with a new deep convolutional neural network. *Multimedia Tools and Applications*, pages 1–19, 2024.
- [8] Murray H Loew. Feature extraction. *Handbook of medical imaging*, 2:273–342, 2000.
- [9] Scott E Umbaugh. *Digital image processing and analysis: computer vision and image analysis*. CRC Press, 2023.
- [10] David A Forsyth and Jean Ponce. *Computer vision: a modern approach*. prentice hall professional technical reference, 2002.
- [11] Dengsheng Zhang and Guojun Lu. Review of shape representation and description techniques. *Pattern recognition*, 37(1):1–19, 2004.
- [12] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical image computing and computer-assisted intervention–MICCAI 2015: 18th international conference, Munich, Germany, October 5–9, 2015, proceedings, part III* 18, pages 234–241. Springer, 2015.

- [13] Mingle Xu, Sook Yoon, Alvaro Fuentes, and Dong Sun Park. A comprehensive survey of image augmentation techniques for deep learning. *Pattern Recognition*, 137:109347, 2023.
- [14] Ravindra Singh, Naurang Singh Mangat, Ravindra Singh, and Naurang Singh Mangat. Stratified sampling. *Elements of survey sampling*, pages 102–144, 1996.
- [15] Corinna Cortes. Support-vector networks. *Machine Learning*, 1995.
- [16] GeeksforGeeks. Difference between random forest vs xgboost, 2024. Accessed: 2024-11-18.
- [17] Saiwa.ai. Xgboost in machine learning: Everything you need to know, 2024. Accessed: 2024-11-18.
- [18] Guestrin Chen. Xgboost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016.
- [19] Random Realizations. Xgboost explained, 2024. Accessed: 2024-11-18.
- [20] The AI Edge. Understanding xgboost: From a to z, 2024. Accessed: 2024-11-19.
- [21] Ragav Venkatesan and Baoxin Li. *Convolutional neural networks in visual computing: a concise guide*. CRC Press, 2017.
- [22] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- [23] Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. Densely connected convolutional networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4700–4708, 2017.
- [24] Zhuang Liu, Hanzi Mao, Chao-Yuan Wu, Christoph Feichtenhofer, Trevor Darrell, and Saining Xie. A convnet for the 2020s. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 11976–11986, 2022.

## 9 Appendices

### 9.1 SVM

```

1 # Load necessary libraries
2 library(keras)
3 library(tensorflow)
4 library(dplyr) # For data manipulation
5 library(imager)
6 library(caret)
7 library(e1071)
8 library(ggplot2)
9 library(EBImage) # For bulabel and shape feature extraction
10
11 # Set the path to your dataset
12 data_dir <- "/kaggle/input/leukemia-images/Original"
13
14 # Define image size and parameters
15 img_height <- 150
16 img_width <- 150
17 batch_size <- 32
18
19 # Create a list to hold all file paths and corresponding labels
20 file_paths <- c()
21 class_labels <- c()
22

```

```

23 # Loop through each class to gather file paths and labels
24 for (class in c("Benign", "Early", "Pre", "Pro")) {
25   class_path <- file.path(data_dir, class)
26   files <- list.files(class_path, full.names = TRUE)
27
28   # Ensure files are unique
29   unique_files <- unique(files)
30
31   file_paths <- c(file_paths, unique_files)
32   class_labels <- c(class_labels, rep(class, length(unique_files))) # Store the
33   class name
34 }
35
36 # Create a DataFrame with only the required structure
37 full_df <- data.frame(
38   image_filename = file_paths,
39   class = class_labels,
40   stringsAsFactors = FALSE
41 )
42 head(full_df)
43
44 # Define class-to-index mapping
45 classes <- c("Benign", "Early", "Pre", "Pro")
46 class_to_index <- setNames(0:(length(classes) - 1), classes)
47
48 # Add numeric columns to the full and training DataFrames
49 full_df$class_numeric <- as.numeric(factor(full_df$class, levels = classes)) - 1
50
51 # Stratified split into training (64%), validation (20%), and test (16%) sets
52 set.seed(123)
53 trainIndex <- createDataPartition(full_df$class, p = 0.64, list = FALSE)
54 train_df <- full_df[trainIndex, ]
55 remaining_df <- full_df[-trainIndex, ]
56
57 # Shuffle only train_df
58 train_df <- train_df[sample(nrow(train_df)), ]
59
60 # Split the remaining 36% into validation (20%) and test (16%)
61 validationIndex <- createDataPartition(remaining_df$class, p = 20 / (20 + 16),
62                                         list = FALSE)
62 val_df <- remaining_df[validationIndex, ]
63 test_df <- remaining_df[-validationIndex, ]
64
65 head(train_df)
66 head(val_df)
67 head(test_df)
68
69 # Calculate class weights using only the training dataset
70 n_samples <- nrow(train_df)
71 n_classes <- length(unique(train_df$class))
72
73 # Get class frequencies from training dataset
74 class_counts <- table(train_df$class_numeric)
75
76 # Calculate balanced weights
77 class_weights <- n_samples / (n_classes * class_counts)
78
79 # Convert to a named list with numeric indices as names

```

```

80 class_weights <- as.list(class_weights)
81 names(class_weights) <- as.character(0:(n_classes-1))
82
83 # Print diagnostics
84 cat("\nClass distribution in training dataset:\n")
85 print(table(train_df$class))
86
87 cat("\nFinal Class Weights (based on training dataset):\n")
88 print(class_weights)
89
90 # Validation steps
91 cat("\nValidation:\n")
92 cat("Total samples in training dataset:", n_samples, "\n")
93 cat("Sum of class counts:", sum(class_counts), "\n")
94 cat("Number of classes:", n_classes, "\n")
95 cat("Maximum weight ratio:", max(unlist(class_weights)) /
96     min(unlist(class_weights)), "\n\n")
97
98 # Additional validation checks
99 cat("Verification of weight calculation:\n")
100 for(i in 0:(n_classes-1)) {
101   class_name <- classes[i+1]
102   weight <- class_weights[[as.character(i)]]
103   count <- class_counts[as.character(i)]
104   cat(sprintf("%s: count=%d, weight=%.4f, count*weight=%.4f\n",
105         class_name, count, weight, count*weight))
106 }
107
108 # Features extraction function
109 extract_features <- function(image_path, bins = 16, feature_params = NULL) {
110   # Load and resize the image using imager
111   image <- load.image(image_path)
112   image <- resize(image, img_width, img_height)
113
114   # Convert image to array and split into color channels
115   red_channel <- as.vector(image[, , 1, 1])
116   green_channel <- as.vector(image[, , 1, 2])
117   blue_channel <- as.vector(image[, , 1, 3])
118
119   # Color histogram features for each channel
120   red_hist <- hist(red_channel, breaks = bins, plot = FALSE)$counts /
121     length(red_channel)
122   green_hist <- hist(green_channel, breaks = bins, plot = FALSE)$counts /
123     length(green_channel)
124   blue_hist <- hist(blue_channel, breaks = bins, plot = FALSE)$counts /
125     length(blue_channel)
126
127   # Convert to grayscale for texture and edge features
128   gray_image <- 0.299 * red_channel + 0.587 * green_channel + 0.114 * blue_channel
129
130   # Texture features: contrast, skewness, and kurtosis
131   contrast <- var(gray_image)
132   skewness <- e1071::skewness(gray_image)
133   kurtosis <- e1071::kurtosis(gray_image)
134
135   # Statistical descriptors
136   mean_intensity <- mean(gray_image)
137   sd_intensity <- sd(gray_image)
138   max_intensity <- max(gray_image)

```

```

135 min_intensity <- min(gray_image)
136
137 # Edge detection using gradient magnitude approximation (Sobel-like)
138 gray_img_cimg <- as.cimg(gray_image, x = img_width, y = img_height)
139 grad_x <- imgradient(gray_img_cimg, "x")
140 grad_y <- imgradient(gray_img_cimg, "y")
141 edge_magnitude <- sqrt(grad_x^2 + grad_y^2) # Approximate edge strength
142 edge_intensity_mean <- mean(edge_magnitude)
143 edge_intensity_sd <- sd(edge_magnitude)
144
145 # Shape features using segmentation (using EBImage)
146 gray_img_matrix <- matrix(gray_image, nrow = img_height, ncol = img_width)
147 binary_image <- gray_img_matrix > 0.5 # Apply threshold to get binary mask
148 binary_image <- as.Image(binary_image) # Convert to EBImage object
149
150 # Label the binary image
151 labeled_image <- bwlablel(binary_image)
152
153 # Compute shape features
154 shape_features <- computeFeatures.shape(labeled_image)
155
156 # Shape descriptors - Mean values of area, perimeter, and circularity
157 area_mean <- mean(shape_features[, "s.area"], na.rm = TRUE)
158 perimeter_mean <- mean(shape_features[, "s.perimeter"], na.rm = TRUE)
159 circularity_mean <- mean((4 * pi * shape_features[, "s.area"]) /
160   (shape_features[, "s.perimeter"]^2), na.rm = TRUE)
161
162 # Combine all features into a single vector
163 feature_vector <- c(red_hist, green_hist, blue_hist, contrast, skewness,
164 kurtosis,
165 mean_intensity, sd_intensity, max_intensity, min_intensity,
166 edge_intensity_mean, edge_intensity_sd,
167 area_mean, perimeter_mean, circularity_mean)
168
169 return(feature_vector)
170 }
171
172 # Function to analyze feature lengths across training data
173 get_feature_params <- function(train_image_paths, sample_size = NULL) {
174   # If sample_size is provided, randomly sample images for efficiency
175   if (!is.null(sample_size) && sample_size < length(train_image_paths)) {
176     set.seed(123) # for reproducibility
177     image_paths <- sample(train_image_paths, sample_size)
178   } else {
179     image_paths <- train_image_paths
180   }
181
182   cat("Analyzing features from", length(image_paths), "training images...\n")
183
184   # Extract features from all training images and analyze
185   feature_lengths <- sapply(image_paths, function(path) {
186     tryCatch(
187       features <- extract_features(path)
188       length(features)
189     ), error = function(e) {
190       warning(sprintf("Error processing image %s: %s", path, e$message))
191       return(NA)
192     })
193   })

```

```

192
193 # Analyze feature length distribution
194 length_summary <- summary(feature_lengths)
195 length_table <- table(feature_lengths)
196
197 # Check for inconsistencies
198 if (length(unique(feature_lengths[!is.na(feature_lengths)])) > 1) {
199   warning("Inconsistent feature lengths detected in training data!")
200   print(length_table)
201   print(length_summary)
202 }
203
204 # Get most common feature length (mode)
205 modal_length <- as.numeric(names(length_table)[which.max(length_table)])
206
207 # Get feature names from a successful extraction
208 sample_features <- NULL
209 for (path in image_paths) {
210   tryCatch({
211     sample_features <- extract_features(path)
212     if (length(sample_features) == modal_length) break
213   }, error = function(e) {
214     next
215   })
216 }
217
218 params <- list(
219   feature_length = modal_length,
220   feature_names = paste0("feature_", seq_len(modal_length)),
221   length_distribution = length_table,
222   length_summary = length_summary,
223   sample_size = length(image_paths),
224   success_rate = sum(!is.na(feature_lengths)) / length(feature_lengths),
225   consistent_length = length(unique(feature_lengths[!is.na(feature_lengths)]))
226   == 1
227 )
228
229 # Print summary
230 cat("\nFeature extraction analysis summary:\n")
231 cat("Modal feature length:", params$feature_length, "\n")
232 cat("Success rate:", sprintf("%.2f%%", params$success_rate * 100), "\n")
233 cat("Length consistency:", ifelse(params$consistent_length, "Consistent",
234   "Inconsistent"), "\n")
235
236 if (!params$consistent_length) {
237   cat("\nFeature length distribution:\n")
238   print(params$length_distribution)
239 }
240
241 return(params)
242 }
243
244 # Modified feature extraction pipeline
245 set.seed(123) # for reproducibility
246
247 # Step 1: Get feature parameters from training data
248 # Option to sample if training set is very large
249 sample_size <- if(nrow(train_df) > 1000) 1000 else NULL
250 feature_params <- get_feature_params(train_df$image_filename, sample_size)

```

```

249
250 # Enhanced feature extraction function with validation
251 extract_and_validate_features <- function(image_paths, feature_params,
252   dataset_name) {
252   features_list <- lapply(seq_along(image_paths), function(i) {
253     path <- image_paths[i]
254
255     if (i %% 100 == 0) {
256       cat(sprintf("Processing %s image %d of %d\n",
257           dataset_name, i, length(image_paths)))
258     }
259
260     tryCatch({
261       features <- extract_features(path)
262
263       # Validate feature length
264       if (length(features) != feature_params$feature_length) {
265
266         # Handle inconsistent length
267         if (length(features) < feature_params$feature_length) {
268           # Pad with NA if shorter
269           features <- c(features, rep(NA, feature_params$feature_length -
270             length(features)))
271         } else {
272           # Truncate if longer
273           features <- features[1:feature_params$feature_length]
274         }
275
276       return(features)
277
278     }, error = function(e) {
279       warning(sprintf("Error processing %s image %s: %s",
280           dataset_name, path, e$message))
281       return(rep(NA, feature_params$feature_length))
282     })
283   })
284
285   # Convert to matrix/data frame
286   features_matrix <- do.call(rbind, features_list)
287   colnames(features_matrix) <- feature_params$feature_names
288
289   # Report statistics about NA values
290   na_counts <- colSums(is.na(features_matrix))
291   if (any(na_counts > 0)) {
292     cat("\nNA counts in", dataset_name, "dataset:\n")
293     print(na_counts[na_counts > 0])
294   }
295
296   return(features_matrix)
297 }
298
299 # Extract features for each dataset
300 cat("\nExtracting training features...\n")
301 train_features <- extract_and_validate_features(
302   train_df$image_filename,
303   feature_params,
304   "training"
305 )

```

```

306
307 cat("\nExtracting validation features...\n")
308 val_features <- extract_and_validate_features(
309   val_df$image_filename,
310   feature_params,
311   "validation"
312 )
313
314 cat("\nExtracting test features...\n")
315 test_features <- extract_and_validate_features(
316   test_df$image_filename,
317   feature_params,
318   "test"
319 )
320
321 # Ensure that class labels are factors with consistent levels
322 train_df$class <- factor(train_df$class, levels = c("Benign", "Early", "Pre",
323   "Pro"))
323 test_df$class <- factor(test_df$class, levels = c("Benign", "Early", "Pre",
324   "Pro"))
324 val_df$class <- factor(val_df$class, levels = c("Benign", "Early", "Pre", "Pro"))
325
326 train_data <- cbind(as.data.frame(train_features), class_numeric =
327   train_df$class_numeric)
327 validation_data <- cbind(as.data.frame(val_features), class_numeric =
328   val_df$class_numeric)
328 test_data <- cbind(as.data.frame(test_features), class_numeric =
329   test_df$class_numeric)
330
330 train_data[is.na(train_data)] <- 0
331 validation_data[is.na(validation_data)] <- 0
332 test_data[is.na(test_data)] <- 0
333
334 # Standardize the features
335 preprocess_params <- preProcess(train_data[, -ncol(train_data)], method =
336   c("center", "scale"))
336 train_data_scaled <- predict(preprocess_params, train_data)
337 validation_data_scaled <- predict(preprocess_params, validation_data)
338 test_data_scaled <- predict(preprocess_params, test_data)
339
340 # Convert class_numeric to factor with proper labels for training and testing
341 train_data_scaled$class <- factor(train_data_scaled$class_numeric,
342   levels = 0:3,
343   labels = classes)
344 validation_data_scaled$class <- factor(validation_data_scaled$class_numeric,
345   levels = 0:3,
346   labels = classes)
347 test_data_scaled$class <- factor(test_data_scaled$class_numeric,
348   levels = 0:3,
349   labels = classes)
350
351 library(caret)
352 library(e1071)
353
354
355 # Function to identify NZV features from training data
356 identify_nzv_features <- function(train_data) {
357   nzv <- nearZeroVar(train_data, saveMetrics = TRUE)
358   nzv_features <- rownames(nzv)[nzv$zeroVar | nzv$nzv]

```

```

359
360 # Log the features being removed and their properties
361 cat("Features removed due to near-zero variance:\n")
362 for(feature in nzv_features) {
363   cat(sprintf("Feature: %s, Variance: %.6f\n",
364         feature,
365         var(train_data[[feature]], na.rm = TRUE)))
366 }
367
368 return(nzv_features)
369 }
370
371 # Function to check for differences in feature variance between datasets
372 check_feature_variance <- function(train_data, val_data, test_data, nzv_features)
373 {
374   warning_features <- list()
375
376   for(feature in nzv_features) {
377     if(feature %in% names(val_data) && feature %in% names(test_data)) {
378       train_var <- var(train_data[[feature]], na.rm = TRUE)
379       val_var <- var(val_data[[feature]], na.rm = TRUE)
380       test_var <- var(test_data[[feature]], na.rm = TRUE)
381
382       # Check if feature has significant variance in validation or test
383       if(val_var > 0.1 || test_var > 0.1) { # Threshold can be adjusted
384         warning_features[[feature]] <- list(
385           train_variance = train_var,
386           val_variance = val_var,
387           test_variance = test_var
388         )
389       }
390     }
391   }
392
393   return(warning_features)
394 }
395
396 # Apply the fixed NZV filtering process
397 # First, identify NZV features from training data only
398 nzv_features <- identify_nzv_features(train_data_scaled[, -ncol(train_data_scaled)])
399
400 # Check for potential issues in validation and test sets
401 variance_warnings <- check_feature_variance(
402   train_data_scaled[, -ncol(train_data_scaled)],
403   validation_data_scaled[, -ncol(validation_data_scaled)],
404   test_data_scaled[, -ncol(test_data_scaled)],
405   nzv_features
406 )
407
408 # Print warnings if any features show different behavior in val/test
409 if(length(variance_warnings) > 0) {
410   cat("\nWARNING: Some features show different variance patterns in
411         validation/test sets:\n")
412   print(variance_warnings)
413   cat("\nConsider keeping these features in the model.\n")
414 }
415
416 # Remove NZV features from all datasets

```

```

415 train_data_filtered <- train_data_scaled[, !names(train_data_scaled) %in%
416   nzv_features]
416 validation_data_filtered <- validation_data_scaled[, !names(validation_data_scaled) %in% nzv_features]
417 test_data_filtered <- test_data_scaled[, !names(test_data_scaled) %in% nzv_features]
418
419 head(train_data_filtered)
420 head(validation_data_filtered)
421 head(test_data_filtered)
422
423 # Set the number of folds for cross-validation
424 num_folds <- 5
425
426 # Create a 5-fold cross-validation partition for the training set
427 set.seed(123) # For reproducibility
428 folds <- createFolds(train_data_filtered$class, k = num_folds, list = TRUE)
429
430 # Define the tune grid for SVM hyperparameters (sigma and C)
431 tune_grid <- expand.grid(
432   sigma = 10^(-3:3),
433   C = 10^(-3:3)
434 )
435
436 # Store performance results for each fold
437 cv_results <- data.frame(
438   fold = integer(),
439   train_accuracy = numeric(),
440   valid_accuracy = numeric(),
441   sigma = numeric(),
442   C = numeric(),
443   stringsAsFactors = FALSE
444 )
445
446 # Loop through each fold to perform training and validation
447 for (fold_idx in 1:num_folds) {
448   cat("Training fold", fold_idx, "of", num_folds, "\n")
449
450   # Get the training data for the current fold
451   train_indices <- unlist(folds[[fold_idx]])
452   train_fold <- train_data_filtered[train_indices, ]
453
454   # Create class weights for the training fold
455   class_weights_fold <- sapply(train_fold$class_numeric,
456                                 function(x) class_weights[[as.character(x)]])
457
458   # Initialize tracking for best model
459   best_valid_accuracy <- -Inf
460   best_train_accuracy <- -Inf
461   best_svm_model <- NULL
462   best_sigma <- NULL
463   best_C <- NULL
464
465   # Grid search over hyperparameters
466   for (sigma_val in tune_grid$sigma) {
467     for (C_val in tune_grid$C) {
468       # Train the SVM model
469       tryCatch({
470         model <- svm(

```

```

471     class ~ .,
472     data = train_fold[, !names(train_fold) %in% c("class_numeric")],
473     kernel = "radial",
474     cost = C_val,
475     gamma = sigma_val,
476     weights = class_weights_fold,
477     probability = TRUE,
478     scale = FALSE # Disable internal scaling since data is already scaled
479   )
480
481   # Make predictions on the external validation set
482   valid_predictions <- predict(model, validation_data_filtered)
483   valid_accuracy <- confusionMatrix(valid_predictions,
484
485   validation_data_filtered$class)$overall['Accuracy']
486
487   # Make predictions on training fold
488   train_predictions <- predict(model, train_fold)
489   train_accuracy <- confusionMatrix(train_predictions,
490           train_fold$class)$overall['Accuracy']
491
492   # Track the best model based on validation accuracy
493   if (valid_accuracy > best_valid_accuracy) {
494     best_valid_accuracy <- valid_accuracy
495     best_train_accuracy <- train_accuracy
496     best_svm_model <- model
497     best_sigma <- sigma_val
498     best_C <- C_val
499   }
500   }, error = function(e) {
501     cat("Error with sigma =", sigma_val, "and C =", C_val, ":", e$message,
502         "\n")
503   }
504 }
505
506 # Store the results for the current fold
507 cv_results <- rbind(cv_results, data.frame(
508   fold = fold_idx,
509   train_accuracy = best_train_accuracy,
510   valid_accuracy = best_valid_accuracy,
511   sigma = best_sigma,
512   C = best_C,
513   stringsAsFactors = FALSE
514 ))
515 cat("Best model for fold", fold_idx, ":\n",
516     "sigma =", best_sigma, "\n",
517     "C =", best_C, "\n",
518     "Training accuracy:", best_train_accuracy, "\n",
519     "Validation accuracy:", best_valid_accuracy, "\n\n")
520 }
521
522 # Print cross-validation results summary
523 cat("Cross-validation results:\n")
524 print(cv_results)
525
526 # Calculate summary statistics
527 mean_train_accuracy <- mean(cv_results$train_accuracy)

```

```

528 sd_train_accuracy <- sd(cv_results$train_accuracy)
529 mean_valid_accuracy <- mean(cv_results$valid_accuracy)
530 sd_valid_accuracy <- sd(cv_results$valid_accuracy)
531
532 cat("\nSummary Statistics:\n")
533 cat("Mean training accuracy:", mean_train_accuracy, "±", sd_train_accuracy, "\n")
534 cat("Mean validation accuracy:", mean_valid_accuracy, "±", sd_valid_accuracy,
      "\n")
535
536 # Select best hyperparameters based on validation performance
537 best_params <- cv_results[which.max(cv_results$valid_accuracy), ]
538
539 # Train final model on full training data
540 final_svm_model <- svm(
541   class ~ .,
542   data = train_data_filtered[, !names(train_data_filtered) %in%
543     c("class_numeric")],
544   kernel = "radial",
545   cost = best_params$C,
546   gamma = best_params$sigma,
547   weights = sapply(train_data_filtered$class_numeric,
548     function(x) class_weights[[as.character(x)]]),
549   probability = TRUE,
550   scale = FALSE
551 )
552
553 # Evaluate final model on all datasets
554 evaluate_model <- function(model, data, set_name) {
555   predictions <- predict(model, data)
556   cm <- confusionMatrix(predictions, data$class)
557   cat("\n", set_name, "set evaluation:\n")
558   print(cm)
559   return(cm)
560 }
561
562 # Final evaluations
563 train_cm <- evaluate_model(final_svm_model, train_data_filtered, "Training")
564 validation_cm <- evaluate_model(final_svm_model, validation_data_filtered,
565   "Validation")
566 test_cm <- evaluate_model(final_svm_model, test_data_filtered, "Test")
567
568 cat("\nClass-Specific Statistics:\n")
569 print(test_cm$byClass)
570
571 cat("\nConfusion Matrix Mode:\n")
572 print(test_cm$mode)
573
574 cat("\nAdditional Arguments Passed to confusionMatrix():\n")
575 print(test_cm$dots)
576
577 # Print overall accuracy results
578 accuracy_test <- test_cm$overall['Accuracy']
579 accuracy_train <- train_cm$overall['Accuracy']
580 cat("Final Test Accuracy:", accuracy_test, "\n")
581 cat("Final Train Accuracy:", accuracy_train, "\n")
582
583 # Evaluating misclassification cost
584 # Define the cost matrix (replace with your actual costs)
585 cost_matrix <- matrix(

```

```

584   c(0, 10, 20, 30,    # Benign -> Benign, Early, Pre, Pro
585     10, 0, 15, 25,    # Early -> Benign, Early, Pre, Pro
586     20, 15, 0, 10,    # Pre -> Benign, Early, Pre, Pro
587     30, 25, 10, 0),   # Pro -> Benign, Early, Pre, Pro
588   nrow = 4, byrow = TRUE
589 )
590 rownames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
591 colnames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
592
593 # Extract confusion matrix data
594 cm_table <- test_cm$table
595 classes <- rownames(cm_table)
596
597 visualize_cost_confusion_matrix <- function(test_cm, cost_matrix) {
598   # Extract the confusion matrix from caret's result
599   cm_table <- test_cm$table
600
601   # Create cost-weighted matrix
602   cost_weighted_cm <- matrix(0,
603                               nrow = nrow(cost_matrix),
604                               ncol = ncol(cost_matrix))
605
606   for (i in 1:nrow(cm_table)) {
607     for (j in 1:ncol(cm_table)) {
608       if (i != j) {
609         cost_weighted_cm[i, j] <- cm_table[i, j] * cost_matrix[i, j]
610       }
611     }
612   }
613
614   # Set row and column names
615   rownames(cost_weighted_cm) <- rownames(cost_matrix)
616   colnames(cost_weighted_cm) <- colnames(cost_matrix)
617
618   return(cost_weighted_cm)
619 }
620
621 # Usage
622 cost_weighted_matrix <- visualize_cost_confusion_matrix(test_cm, cost_matrix)
623 cat("\nCost-Weighted Misclassification Matrix:\n")
624 print(cost_weighted_matrix)
625
626 # Initialize total misclassification cost
627 total_cost <- 0
628
629 # Calculate total misclassification cost by iterating over confusion matrix
630   entries
631 for (i in 1:nrow(cm_table)) {
632   for (j in 1:ncol(cm_table)) {
633     if (i != j) { # Misclassifications (i != j)
634       total_cost <- total_cost + cm_table[i, j] * cost_matrix[i, j]
635     }
636   }
637
638 # Output the total misclassification cost
639 cat("Total Misclassification Cost:", total_cost, "\n")
640
641 # Calculate average cost per sample

```

```

642 total_samples <- sum(cm_table)
643 avg_cost_per_sample <- total_cost / total_samples
644 cat("Average Cost Per Sample:", avg_cost_per_sample, "\n")
645
646 # Calculate cost-sensitive accuracy
647 # Compute the maximum possible cost (if all instances were misclassified to the
# most costly class)
648 max_possible_cost <- sum(cost_matrix) * (total_samples / nrow(cost_matrix))
649 cost_sensitive_accuracy <- 1 - (total_cost / max_possible_cost)
650 cat("Cost-Sensitive Accuracy:", cost_sensitive_accuracy, "\n")
651
652 # Using cost matrix for cost-weighted F1 calculation
653 f1_scores <- test_cm$byClass[, "F1"]
654 weights <- 1 / (diag(cost_matrix) + 1) # Add 1 to avoid divide-by-zero
655 weights <- weights / sum(weights)
656 weighted_f1 <- sum(f1_scores * weights, na.rm = TRUE)
657 cat("Weighted F1-score", weighted_f1, "\n")
658
659 # Save final model and results
660 save(final_svm_model, cv_results, train_cm, validation_cm, test_cm,
661       file = "final_svm_model.RData")

```

### 9.1.1 XGBoost

```

1 # Load necessary libraries
2 library(xgboost)
3 library(caret)
4 library(dplyr)
5 library(imager)
6 library(EBImage) # For bwlabel and shape feature extraction
7 library(ggplot2)
8 library(e1071) # For skewness and kurtosis
9
10 # Set the path to your dataset
11 data_dir <- "/kaggle/input/leukemia-images/Original"
12
13 # Define image size and parameters
14 img_height <- 150
15 img_width <- 150
16 batch_size <- 32 # Note: batch_size is defined but not used in the code
17
18 # Create a list to hold all file paths and corresponding labels
19 file_paths <- character() # More explicit initialization
20 class_labels <- character()
21
22 # Loop through each class to gather file paths and labels
23 for (class in c("Benign", "Early", "Pre", "Pro")) {
24   class_path <- file.path(data_dir, class)
25
26   # Add error handling for directory access
27   if (!dir.exists(class_path)) {
28     stop(paste("Directory not found:", class_path))
29   }
30
31   files <- list.files(class_path, full.names = TRUE, pattern =
32                       "\\\.(?i)(jpg|jpeg|png|tiff)$")
33
34   # Ensure files are unique

```

```

34 unique_files <- unique(files)
35
36 file_paths <- c(file_paths, unique_files)
37 class_labels <- c(class_labels, rep(class, length(unique_files)))
38 }
39
40 # Create a DataFrame with only the required structure
41 full_df <- data.frame(
42   image_filename = file_paths,
43   class = factor(class_labels, levels = c("Benign", "Early", "Pre", "Pro")),
44   stringsAsFactors = FALSE
45 )
46
47 # Define class-to-index mapping
48 classes <- c("Benign", "Early", "Pre", "Pro")
49 class_to_index <- setNames(0:(length(classes) - 1), classes)
50
51 # Add numeric class column (0-based indexing)
52 full_df$class_numeric <- as.numeric(full_df$class) - 1
53
54 # Stratified split into training (64%), validation (20%), and test (16%) sets
55 set.seed(123)
56 trainIndex <- createDataPartition(full_df$class, p = 0.64, list = FALSE)
57 train_df <- full_df[trainIndex, ]
58 remaining_df <- full_df[-trainIndex, ]
59
60 # Shuffle only train_df
61 train_df <- train_df[sample(nrow(train_df)), ]
62
63 validationIndex <- createDataPartition(remaining_df$class, p = 20/36, list =
  FALSE)
64 val_df <- remaining_df[validationIndex, ]
65 test_df <- remaining_df[-validationIndex, ]
66
67 # Calculate class weights using only the training dataset
68 n_samples <- nrow(train_df)
69 n_classes <- length(unique(train_df$class))
70
71 # Get class frequencies from training dataset
72 class_counts <- table(train_df$class_numeric)
73
74 # Calculate balanced weights
75 class_weights <- n_samples / (n_classes * class_counts)
76
77 # Convert to a named list with numeric indices as names
78 class_weights <- as.list(class_weights)
79 names(class_weights) <- as.character(0:(n_classes-1))
80
81 # Print diagnostics
82 cat("\nClass distribution in training dataset:\n")
83 print(table(train_df$class))
84
85 cat("\nFinal Class Weights (based on training dataset):\n")
86 print(class_weights)
87
88 # Validation steps
89 cat("\nValidation:\n")
90 cat("Total samples in training dataset:", n_samples, "\n")
91 cat("Sum of class counts:", sum(class_counts), "\n")

```

```

92 cat("Number of classes:", n_classes, "\n")
93 cat("Maximum weight ratio:", max(unlist(class_weights)) /
94     min(unlist(class_weights)), "\n\n")
95 # Additional validation checks
96 cat("Verification of weight calculation:\n")
97 for(i in 0:(n_classes-1)) {
98   class_name <- classes[i+1]
99   weight <- class_weights[[as.character(i)]]
100  count <- class_counts[as.character(i)]
101  cat(sprintf("%s: count=%d, weight=%.4f, count*weight=%.4f\n",
102          class_name, count, weight, count*weight))
103 }
104
105 # Features extraction function
106 extract_features <- function(image_path, bins = 16, feature_params = NULL) {
107   # Load and resize the image using imager
108   image <- load.image(image_path)
109   image <- resize(image, img_width, img_height)
110
111   # Convert image to array and split into color channels
112   red_channel <- as.vector(image[, , 1])
113   green_channel <- as.vector(image[, , 2])
114   blue_channel <- as.vector(image[, , 3])
115
116   # Color histogram features for each channel
117   red_hist <- hist(red_channel, breaks = bins, plot = FALSE)$counts /
118     length(red_channel)
119   green_hist <- hist(green_channel, breaks = bins, plot = FALSE)$counts /
120     length(green_channel)
121   blue_hist <- hist(blue_channel, breaks = bins, plot = FALSE)$counts /
122     length(blue_channel)
123
124   # Convert to grayscale for texture and edge features
125   gray_image <- 0.299 * red_channel + 0.587 * green_channel + 0.114 * blue_channel
126
127   # Texture features: contrast, skewness, and kurtosis
128   contrast <- var(gray_image)
129   skewness <- e1071::skewness(gray_image)
130   kurtosis <- e1071::kurtosis(gray_image)
131
132   # Statistical descriptors
133   mean_intensity <- mean(gray_image)
134   sd_intensity <- sd(gray_image)
135   max_intensity <- max(gray_image)
136   min_intensity <- min(gray_image)
137
138   # Edge detection using gradient magnitude approximation (Sobel-like)
139   gray_img_cimg <- as.cimg(gray_image, x = img_width, y = img_height)
140   grad_x <- imgradient(gray_img_cimg, "x")
141   grad_y <- imgradient(gray_img_cimg, "y")
142   edge_magnitude <- sqrt(grad_x^2 + grad_y^2) # Approximate edge strength
143   edge_intensity_mean <- mean(edge_magnitude)
144   edge_intensity_sd <- sd(edge_magnitude)
145
146   # Shape features using segmentation (using EBImage)
147   gray_img_matrix <- matrix(gray_image, nrow = img_height, ncol = img_width)
148   binary_image <- gray_img_matrix > 0.5 # Apply threshold to get binary mask
149   binary_image <- as.Image(binary_image) # Convert to EBImage object

```

```

147
148 # Label the binary image
149 labeled_image <- bwlabel(binary_image)
150
151 # Compute shape features
152 shape_features <- computeFeatures.shape(labeled_image)
153
154 # Shape descriptors - Mean values of area, perimeter, and circularity
155 area_mean <- mean(shape_features[, "s.area"], na.rm = TRUE)
156 perimeter_mean <- mean(shape_features[, "s.perimeter"], na.rm = TRUE)
157 circularity_mean <- mean((4 * pi * shape_features[, "s.area"]) /
158   (shape_features[, "s.perimeter"]^2), na.rm = TRUE)
159
160 # Combine all features into a single vector
161 feature_vector <- c(red_hist, green_hist, blue_hist, contrast, skewness,
162   kurtosis,
163   mean_intensity, sd_intensity, max_intensity, min_intensity,
164   edge_intensity_mean, edge_intensity_sd,
165   area_mean, perimeter_mean, circularity_mean)
166
167 return(feature_vector)
168 }
169
170 # Function to analyze feature lengths across training data
171 get_feature_params <- function(train_image_paths, sample_size = NULL) {
172   # If sample_size is provided, randomly sample images for efficiency
173   if (!is.null(sample_size) && sample_size < length(train_image_paths)) {
174     set.seed(123) # for reproducibility
175     image_paths <- sample(train_image_paths, sample_size)
176   } else {
177     image_paths <- train_image_paths
178   }
179
180   cat("Analyzing features from", length(image_paths), "training images...\n")
181
182   # Extract features from all training images and analyze
183   feature_lengths <- sapply(image_paths, function(path) {
184     tryCatch({
185       features <- extract_features(path)
186       length(features)
187     }, error = function(e) {
188       warning(sprintf("Error processing image %s: %s", path, e$message))
189       return(NA)
190     })
191   })
192
193   # Analyze feature length distribution
194   length_summary <- summary(feature_lengths)
195   length_table <- table(feature_lengths)
196
197   # Check for inconsistencies
198   if (length(unique(feature_lengths[!is.na(feature_lengths)])) > 1) {
199     warning("Inconsistent feature lengths detected in training data!")
200     print(length_table)
201     print(length_summary)
202   }
203
204   # Get most common feature length (mode)
205   modal_length <- as.numeric(names(length_table)[which.max(length_table)])

```

```

204
205 # Get feature names from a successful extraction
206 sample_features <- NULL
207 for (path in image_paths) {
208   tryCatch({
209     sample_features <- extract_features(path)
210     if (length(sample_features) == modal_length) break
211   }, error = function(e) {
212     next
213   })
214 }
215
216 params <- list(
217   feature_length = modal_length,
218   feature_names = paste0("feature_", seq_len(modal_length)),
219   length_distribution = length_table,
220   length_summary = length_summary,
221   sample_size = length(image_paths),
222   success_rate = sum(!is.na(feature_lengths)) / length(feature_lengths),
223   consistent_length = length(unique(feature_lengths[!is.na(feature_lengths)])))
224   == 1
225 )
226
227 # Print summary
228 cat("\nFeature extraction analysis summary:\n")
229 cat("Modal feature length:", params$feature_length, "\n")
230 cat("Success rate:", sprintf("%.2f%%", params$success_rate * 100), "\n")
231 cat("Length consistency:", ifelse(params$consistent_length, "Consistent",
232   "Inconsistent"), "\n")
233
234 if (!params$consistent_length) {
235   cat("\nFeature length distribution:\n")
236   print(params$length_distribution)
237 }
238
239 return(params)
240}
241
242 # Modified feature extraction pipeline
243 set.seed(123) # for reproducibility
244
245 # Step 1: Get feature parameters from training data
246 # Option to sample if training set is very large
247 sample_size <- if(nrow(train_df) > 1000) 1000 else NULL
248 feature_params <- get_feature_params(train_df$image_filename, sample_size)
249
250 # Enhanced feature extraction function with validation
251 extract_and_validate_features <- function(image_paths, feature_params,
252   dataset_name) {
253   features_list <- lapply(seq_along(image_paths), function(i) {
254     path <- image_paths[i]
255
256     if (i %% 100 == 0) {
257       cat(sprintf("Processing %s image %d of %d\n",
258             dataset_name, i, length(image_paths)))
259     }
260
261     tryCatch({
262       features <- extract_features(path)

```

```

260
261     # Validate feature length
262     if (length(features) != feature_params$feature_length) {
263
264         # Handle inconsistent length
265         if (length(features) < feature_params$feature_length) {
266             # Pad with NA if shorter
267             features <- c(features, rep(NA, feature_params$feature_length -
268             length(features)))
269         } else {
270             # Truncate if longer
271             features <- features[1:feature_params$feature_length]
272         }
273     }
274
275     return(features)
276 }, error = function(e) {
277     warning(sprintf("Error processing %s image %s: %s",
278                 dataset_name, path, e$message))
279     return(rep(NA, feature_params$feature_length))
280 })
281 })
282
283 # Convert to matrix/data frame
284 features_matrix <- do.call(rbind, features_list)
285 colnames(features_matrix) <- feature_params$feature_names
286
287 # Report statistics about NA values
288 na_counts <- colSums(is.na(features_matrix))
289 if (any(na_counts > 0)) {
290     cat("\nNA counts in", dataset_name, "dataset:\n")
291     print(na_counts[na_counts > 0])
292 }
293
294 return(features_matrix)
295 }
296
297 # Extract features for each dataset
298 cat("\nExtracting training features...\n")
299 train_features <- extract_and_validate_features(
300     train_df$image_filename,
301     feature_params,
302     "training"
303 )
304
305 cat("\nExtracting validation features...\n")
306 val_features <- extract_and_validate_features(
307     val_df$image_filename,
308     feature_params,
309     "validation"
310 )
311
312 cat("\nExtracting test features...\n")
313 test_features <- extract_and_validate_features(
314     test_df$image_filename,
315     feature_params,
316     "test"
317 )

```

```

318
319 # Ensure that class labels are factors with consistent levels
320 train_df$class <- factor(train_df$class, levels = c("Benign", "Early", "Pre",
321   "Pro"))
321 test_df$class <- factor(test_df$class, levels = c("Benign", "Early", "Pre",
322   "Pro"))
322 val_df$class <- factor(val_df$class, levels = c("Benign", "Early", "Pre", "Pro"))
323
324 train_data <- cbind(as.data.frame(train_features), class_numeric =
325   train_df$class_numeric)
325 validation_data <- cbind(as.data.frame(val_features), class_numeric =
326   val_df$class_numeric)
326 test_data <- cbind(as.data.frame(test_features), class_numeric =
327   test_df$class_numeric)
328
328 train_data[is.na(train_data)] <- 0
329 validation_data[is.na(validation_data)] <- 0
330 test_data[is.na(test_data)] <- 0
331
332 # Convert class_numeric to factor with proper labels for training and testing
333 train_data$class <- factor(train_data$class_numeric,
334   levels = 0:3,
335   labels = classes)
336 validation_data$class <- factor(validation_data$class_numeric,
337   levels = 0:3,
338   labels = classes)
339 test_data$class <- factor(test_data$class_numeric,
340   levels = 0:3,
341   labels = classes)
342
343 library(caret)
344 library(e1071)
345
346
347 # Function to identify NZV features from training data
348 identify_nzv_features <- function(train_data) {
349   nzv <- nearZeroVar(train_data, saveMetrics = TRUE)
350   nzv_features <- rownames(nzv)[nzv$zeroVar | nzv$nzv]
351
352   # Log the features being removed and their properties
353   cat("Features removed due to near-zero variance:\n")
354   for(feature in nzv_features) {
355     cat(sprintf("Feature: %s, Variance: %.6f\n",
356           feature,
357           var(train_data[[feature]], na.rm = TRUE)))
358   }
359
360   return(nzv_features)
361 }
362
363 # Function to check for differences in feature variance between datasets
364 check_feature_variance <- function(train_data, val_data, test_data, nzv_features)
365   {
366     warning_features <- list()
367
368     for(feature in nzv_features) {
369       if(feature %in% names(val_data) && feature %in% names(test_data)) {
370         train_var <- var(train_data[[feature]], na.rm = TRUE)
371         val_var <- var(val_data[[feature]], na.rm = TRUE)

```

```

371     test_var <- var(test_data[[feature]], na.rm = TRUE)
372 
373     # Check if feature has significant variance in validation or test
374     if(val_var > 0.1 || test_var > 0.1) { # Threshold can be adjusted
375         warning_features[[feature]] <- list(
376             train_variance = train_var,
377             val_variance = val_var,
378             test_variance = test_var
379         )
380     }
381 }
382 }
383 
384 return(warning_features)
385 }
386 
387 # Apply the fixed NZV filtering process
388 # First, identify NZV features from training data only
389 nzv_features <- identify_nzv_features(train_data[, -ncol(train_data)])
390 
391 # Check for potential issues in validation and test sets
392 variance_warnings <- check_feature_variance(
393     train_data[, -ncol(train_data)],
394     validation_data[, -ncol(validation_data)],
395     test_data[, -ncol(test_data)],
396     nzv_features
397 )
398 
399 # Print warnings if any features show different behavior in val/test
400 if(length(variance_warnings) > 0) {
401     cat("\nWARNING: Some features show different variance patterns in
402         validation/test sets:\n")
403     print(variance_warnings)
404     cat("\nConsider keeping these features in the model.\n")
405 }
406 
407 # Remove NZV features and 'class_numeric' from all datasets
408 train_data_filtered <- train_data[, !names(train_data) %in% c(nzv_features,
409     "class_numeric", "class")]
410 validation_data_filtered <- validation_data[, !names(validation_data) %in%
411     c(nzv_features, "class_numeric", "class")]
412 test_data_filtered <- test_data[, !names(test_data) %in% c(nzv_features,
413     "class_numeric", "class")]
414 
415 # Ensure 'class_numeric' is preserved separately for labels
416 train_labels <- train_data$class_numeric
417 validation_labels <- validation_data$class_numeric
418 test_labels <- test_data$class_numeric
419 
420 # Ensure the label vector is preserved separately
421 train_labels <- train_data$class_numeric
422 validation_labels <- validation_data$class_numeric
423 test_labels <- test_data$class_numeric
424 
425 str(train_data_filtered)
426 
427 head(train_data_filtered)
428 head(validation_data_filtered)
429 head(test_data_filtered)

```

```

426
427 class_weights <- as.numeric(class_weights)
428 str(class_weights)
429
430 train_weights <- class_weights[train_df$class_numeric + 1]
431
432 # Check the train_labels values
433 cat("Train Labels:\n")
434 print(head(train_labels)) # Display the first few train labels
435 cat("\n")
436
437 # Check the class_weights values
438 cat("Class Weights:\n")
439 print(class_weights) # Display the class weights
440 cat("\n")
441
442 # Ensure that the class_weights vector is correctly aligned with the class labels
443 # If we have 4 classes, class_weights should correspond to the following:
444 # 0 -> 1.615 (Benign), 1 -> 0.826 (Early), 2 -> 0.845 (Pre), 3 -> 1.013 (Pro)
445
446 # Assign the weights based on train_labels
447 train_weights <- class_weights[train_labels + 1] # Adjusting for 1-based
        indexing in R
448
449 # Check the first few train_weights to confirm they match the expected values
450 cat("Assigned Train Weights (first few):\n")
451 print(head(train_weights)) # Display the first few train weights
452 cat("\n")
453
454 # Check for NA values in the train_weights
455 cat("Checking for NA values in train_weights:\n")
456 na_count <- sum(is.na(train_weights)) # Should return 0 if no NA values
457 cat("Number of NA values in train_weights: ", na_count, "\n")
458 cat("\n")
459
460 # If there are NA values, investigate further
461 if (na_count > 0) {
462   cat("Investigate the indices that caused NAs:\n")
463   na_indices <- which(is.na(train_weights))
464   print(train_labels[na_indices]) # Investigate which train_labels caused the NA
        values
465   cat("\n")
466 }
467
468 # Finally, print the first few rows of train_labels and train_weights to confirm
        alignment
469 cat("Train Labels and Corresponding Train Weights:\n")
470 df_check <- data.frame(train_labels = train_labels[1:10], train_weights =
        train_weights[1:10]) # First 10 rows for quick inspection
471 print(df_check)
472 cat("\n")
473
474 # Prepare DMatrix objects
475 train_matrix <- xgb.DMatrix(data = as.matrix(train_data_filtered), label =
        train_labels, weight = train_weights)
476 validation_matrix <- xgb.DMatrix(data = as.matrix(validation_data_filtered),
        label = validation_labels)
477 test_matrix <- xgb.DMatrix(data = as.matrix(test_data_filtered), label =
        test_labels)

```

```

478
479 # Set XGBoost parameters
480 params <- list(
481   objective = "multi:softmax",
482   num_class = 4,
483   eval_metric = c("merror", "mlogloss"),
484   max_depth = 6,
485   eta = 0.1,
486   subsample = 0.8,
487   colsample_bytree = 0.8,
488   min_child_weight = 1,
489   nthread = parallel::detectCores() - 1
490 )
491
492 # Create watchlist for early stopping
493 watchlist <- list(train = train_matrix, val = validation_matrix)
494
495 # Train model with early stopping
496 xgb_model <- xgb.train(
497   params = params,
498   data = train_matrix,
499   nrounds = 1000,
500   watchlist = watchlist,
501   early_stopping_rounds = 20,
502   verbose = 1
503 )
504
505 # Make predictions on test matrix
506 test_predictions <- predict(xgb_model, test_matrix)
507
508 # Make predictions on train matrix
509 train_predictions <- predict(xgb_model, train_matrix)
510
511 # Evaluate results
512 cm_test <- confusionMatrix(
513   factor(test_predictions, levels = 0:3, labels = levels(full_df$class)),
514   factor(test_df$class_numeric, levels = 0:3, labels = levels(full_df$class)))
515 )
516
517 # Evaluate results
518 cm_train <- confusionMatrix(
519   factor(train_predictions, levels = 0:3, labels = levels(full_df$class)),
520   factor(train_df$class_numeric, levels = 0:3, labels = levels(full_df$class)))
521 )
522
523 # Print results
524 print(cm_test)
525
526 cat("\nClass-Specific Statistics:\n")
527 print(cm_test$byClass)
528
529 cat("\nConfusion Matrix Mode:\n")
530 print(cm_test$mode)
531
532 cat("\nAdditional Arguments Passed to confusionMatrix():\n")
533 print(cm_test$dots)
534
535 # Print overall accuracy results
536 accuracy_test <- cm_test$overall['Accuracy']

```

```

537 accuracy_train <- cm_train$overall['Accuracy']
538 cat("Final Test Accuracy:", accuracy_test, "\n")
539 cat("Final Train Accuracy:", accuracy_train, "\n")
540
541 feature_names <- colnames(train_data_filtered)
542 xgb_importance <- xgb.importance(feature_names = feature_names, model = xgb_model)
543
544 # Select Top 20 Features
545 top_20_features <- head(xgb_importance, 20)
546
547 # Plot the top 20 features
548 xgb.plot.importance(top_20_features)
549
550 # Define the cost matrix (replace with your actual costs)
551 cost_matrix <- matrix(
552   c(0, 10, 20, 30,    # Benign -> Benign, Early, Pre, Pro
553     10, 0, 15, 25,    # Early -> Benign, Early, Pre, Pro
554     20, 15, 0, 10,    # Pre -> Benign, Early, Pre, Pro
555     30, 25, 10, 0),   # Pro -> Benign, Early, Pre, Pro
556   nrow = 4, byrow = TRUE
557 )
558 rownames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
559 colnames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
560
561 # Extract confusion matrix data
562 cm_table <- cm_test$table
563 classes <- rownames(cm_table)
564
565 visualize_cost_confusion_matrix <- function(cm_test, cost_matrix) {
566   # Extract the confusion matrix from caret's result
567   cm_table <- cm_test$table
568
569   # Create cost-weighted matrix
570   cost_weighted_cm <- matrix(0,
571                             nrow = nrow(cost_matrix),
572                             ncol = ncol(cost_matrix))
573
574   for (i in 1:nrow(cm_table)) {
575     for (j in 1:ncol(cm_table)) {
576       if (i != j) {
577         cost_weighted_cm[i, j] <- cm_table[i, j] * cost_matrix[i, j]
578       }
579     }
580   }
581
582   # Set row and column names
583   rownames(cost_weighted_cm) <- rownames(cost_matrix)
584   colnames(cost_weighted_cm) <- colnames(cost_matrix)
585
586   return(cost_weighted_cm)
587 }
588
589 # Usage
590 cost_weighted_matrix <- visualize_cost_confusion_matrix(cm_test, cost_matrix)
591 cat("\nCost-Weighted Misclassification Matrix:\n")
592 print(cost_weighted_matrix)
593
594 # Initialize total misclassification cost
595 total_cost <- 0

```

```

596
597 # Calculate total misclassification cost by iterating over confusion matrix
598 entries
599 for (i in 1:nrow(cm_table)) {
600   for (j in 1:ncol(cm_table)) {
601     if (i != j) { # Misclassifications (i != j)
602       total_cost <- total_cost + cm_table[i, j] * cost_matrix[i, j]
603     }
604   }
605 }
606
607 # Output the total misclassification cost
608 cat("Total Misclassification Cost:", total_cost, "\n")
609
610 # Calculate average cost per sample
611 total_samples <- sum(cm_table)
612 avg_cost_per_sample <- total_cost / total_samples
613 cat("Average Cost Per Sample:", avg_cost_per_sample, "\n")
614
615 # Calculate cost-sensitive accuracy
616 # Compute the maximum possible cost (if all instances were misclassified to the
617 # most costly class)
618 max_possible_cost <- sum(cost_matrix) * (total_samples / nrow(cost_matrix))
619 cost_sensitive_accuracy <- 1 - (total_cost / max_possible_cost)
620 cat("Cost-Sensitive Accuracy:", cost_sensitive_accuracy, "\n")
621
622 # Using cost matrix for cost-weighted F1 calculation
623 f1_scores <- cm_test$byClass[, "F1"]
624 weights <- 1 / (diag(cost_matrix) + 1) # Add 1 to avoid divide-by-zero
625 weights <- weights / sum(weights)
626 weighted_f1 <- sum(f1_scores * weights, na.rm = TRUE)
627 cat("Weighted F1-score", weighted_f1, "\n")
628
629 # Save the trained model
630 xgb.save(xgb_model, "xgb_model.model")

```

### 9.1.2 VGG16

```

1 # Load necessary libraries
2 library(keras)
3 library(tensorflow)
4 library(dplyr) # For data manipulation
5 library(tidyverse) # For tidyverse functions
6 library(caret) # For confusionMatrix and stratified sampling
7
8 # Set the path to your dataset
9 data_dir <- "/kaggle/input/leukemia-images/Original"
10
11 # Define image size and parameters
12 img_height <- 224
13 img_width <- 224
14 batch_size <- 32
15 epochs <- 220
16
17 # Gather file paths and corresponding labels
18 file_paths <- c()
19 class_labels <- c()
20

```

```

21 for (class in c("Benign", "Early", "Pre", "Pro")) {
22   class_path <- file.path(data_dir, class)
23   files <- list.files(class_path, full.names = TRUE)
24   unique_files <- unique(files)
25   file_paths <- c(file_paths, unique_files)
26   class_labels <- c(class_labels, rep(class, length(unique_files)))
27 }
28
29 # Create DataFrame with the necessary structure
30 full_df <- data.frame(
31   image_filename = file_paths,
32   class = class_labels,
33   stringsAsFactors = FALSE
34 )
35
36 # Stratified split into training (64%), validation (20%), and test (16%) sets
37 set.seed(123)
38 trainIndex <- createDataPartition(full_df$class, p = 0.64, list = FALSE)
39 train_df <- full_df[trainIndex, ]
40 remaining_df <- full_df[-trainIndex, ]
41
42 # Shuffle only train_df
43 train_df <- train_df[sample(nrow(train_df)), ]
44
45 # Split the remaining 36% into validation (20%) and test (16%)
46 validationIndex <- createDataPartition(remaining_df$class, p = 20 / (20 + 16),
47                                         list = FALSE)
47 val_df <- remaining_df[validationIndex, ]
48 test_df <- remaining_df[-validationIndex, ]
49
50 # Load VGG16 model from Keras
51 base_model <- application_vgg16(
52   weights = 'imagenet',
53   include_top = FALSE,
54   input_shape = c(img_height, img_width, 3)
55 )
56
57 # Freeze most of the layers of the base model, unfreezing the top few for
58 # fine-tuning
59 for (layer in base_model$layers[1:500]) {
60   layer$trainable <- FALSE
61 }
62 for (layer in base_model$layers[501:length(base_model$layers)]) {
63   layer$trainable <- TRUE
64 }
65
66 # Set the color threshold for purple in HSV space
67 lower_purple <- c(0.7, 0.5, 0.4) # Adjust as needed
68 upper_purple <- c(0.8, 1.0, 1.0) # Adjust as needed
69
70 # Define the model input layer
71 input_image <- layer_input(shape = c(img_height, img_width, 3))
72
73 # Convert the input image to HSV
74 input_image_hsv <- tf$image$rgb_to_hsv(input_image)
75
76 # Create a binary mask for purple regions (lymphoblasts) based on the HSV
77 # threshold
78 mask_purple <- tf$math$logical_and(

```

```

77  input_image_hsv[,,1] >= lower_purple[1], input_image_hsv[,,1] <=
    upper_purple[1]
78 )
79 mask_purple <- tf$math$logical_and(mask_purple, input_image_hsv[,,2] >=
    lower_purple[2])
80 mask_purple <- tf$math$logical_and(mask_purple, input_image_hsv[,,3] >=
    lower_purple[3])
81
82 # Convert logical mask to float and expand dims to match RGB channels
83 mask_purple <- tf$cast(mask_purple, dtype = tf$float32) # Ensure the mask is a
    float tensor
84 mask_purple <- tf$expand_dims(mask_purple, axis = as.integer(-1)) # Add a channel
    dimension
85 mask_purple <- tf$title(mask_purple, multiples = as.integer(c(1, 1, 1, 3))) # 
    Repeat for each color channel
86
87 # Apply the mask to the original RGB image
88 segmented_image <- tf$math$multiply(input_image, mask_purple)
89
90 # Contracting Path (Encoder) with fewer filters
91 conv1 <- layer_conv_2d(segmented_image, filters = 32, kernel_size = c(3, 3),
    activation = "relu", padding = "same")
92 conv1 <- layer_conv_2d(conv1, filters = 32, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
93 pool1 <- layer_max_pooling_2d(conv1, pool_size = c(2, 2))
94
95 conv2 <- layer_conv_2d(pool1, filters = 64, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
96 conv2 <- layer_conv_2d(conv2, filters = 64, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
97 pool2 <- layer_max_pooling_2d(conv2, pool_size = c(2, 2))
98
99 conv3 <- layer_conv_2d(pool2, filters = 128, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
100 conv3 <- layer_conv_2d(conv3, filters = 128, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
101 pool3 <- layer_max_pooling_2d(conv3, pool_size = c(2, 2))
102
103 # Bottleneck layer with fewer filters
104 conv4 <- layer_conv_2d(pool3, filters = 256, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
105 conv4 <- layer_conv_2d(conv4, filters = 256, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
106
107 # Expansive Path (Decoder) with skip connections and fewer filters
108
109 # Upsampling + skip connection from conv3
110 upconv3 <- layer_conv_2d_transpose(conv4, filters = 128, kernel_size = c(3, 3),
    strides = c(2, 2), padding = "same")
111 concat3 <- layer_concatenate(list(upconv3, conv3))
112 conv5 <- layer_conv_2d(concat3, filters = 128, kernel_size = c(3, 3), activation
    = "relu", padding = "same")
113 conv5 <- layer_conv_2d(conv5, filters = 128, kernel_size = c(3, 3), activation =
    "relu", padding = "same")
114
115 # Upsampling + skip connection from conv2
116 upconv2 <- layer_conv_2d_transpose(conv5, filters = 64, kernel_size = c(3, 3),
    strides = c(2, 2), padding = "same")
117 concat2 <- layer_concatenate(list(upconv2, conv2))

```

```

118 conv6 <- layer_conv_2d(concat2, filters = 64, kernel_size = c(3, 3), activation =
119   "relu", padding = "same")
120 conv6 <- layer_conv_2d(conv6, filters = 64, kernel_size = c(3, 3), activation =
121   "relu", padding = "same")
122 # Upsampling + skip connection from conv1
123 upconv1 <- layer_conv_2d_transpose(conv6, filters = 32, kernel_size = c(3, 3),
124   strides = c(2, 2), padding = "same")
125 concat1 <- layer_concatenate(list(upconv1, conv1))
126 conv7 <- layer_conv_2d(concat1, filters = 32, kernel_size = c(3, 3), activation =
127   "relu", padding = "same")
128 conv7 <- layer_conv_2d(conv7, filters = 32, kernel_size = c(3, 3), activation =
129   "relu", padding = "same")
130 # Feature extraction using VGG16
131 features <- base_model(input_image)
132 features <- layer_global_average_pooling_2d(features)
133 # Flatten and concatenate segmentation output with VGG16 features
134 segmentation_features <- layer_flatten(conv5)
135 combined_features <- layer_concatenate(list(features, segmentation_features))
136 # Dense layers with L2 regularization, batch normalization, and increased dropout
137 x <- layer_dense(combined_features, units = 512, activation = "relu",
138   kernel_regularizer = regularizer_l2(0.01))
139 x <- layer_batch_normalization(x)
140 x <- layer_dropout(x, rate = 0.7) # Increased dropout rate to 0.7
141 x <- layer_dense(x, units = 256, activation = "relu", kernel_regularizer =
142   regularizer_l2(0.01))
143 x <- layer_batch_normalization(x)
144 x <- layer_dropout(x, rate = 0.7) # Increased dropout rate to 0.7
145 x <- layer_dense(x, units = 128, activation = "relu", kernel_regularizer =
146   regularizer_l2(0.01))
147 x <- layer_batch_normalization(x)
148 x <- layer_dropout(x, rate = 0.7) # Increased dropout rate to 0.7
149 output <- layer_dense(x, units = length(unique(full_df$class)), activation =
150   "softmax")
151 # Create the model
152 model <- keras_model(inputs = input_image, outputs = output)
153 # Compile the model with a lower learning rate to prevent overfitting
154 model %>% compile(
155   loss = "categorical_crossentropy",
156   optimizer = optimizer_adam(learning_rate = 1e-5),
157   metrics = c("accuracy")
158 )
159
160 print(model)
161
162 # Define the class-to-index mapping
163 classes <- c("Benign", "Early", "Pre", "Pro")
164 class_to_index <- setNames(0:(length(classes) - 1), classes)
165
166 # Add numeric columns to both full and training DataFrames
167 full_df$class_numeric <- as.numeric(factor(full_df$class, levels = classes)) - 1

```

```

168 train_df$class_numeric <- as.numeric(factor(train_df$class, levels = classes)) - 1
169
170 # Verify the class-to-index mapping
171 cat("Class-to-Index Mapping:\n")
172 print(class_to_index)
173
174 # Create improved mapping displays with counts
175 cat("\nMappings and counts in full dataset:\n")
176 unique_mapping_full <- data.frame(
177   original_class = classes,
178   numeric_class = 0:(length(classes)-1),
179   count = as.vector(table(full_df$class))
180 )
181 print(unique_mapping_full)
182
183 cat("\nMappings and counts in training dataset:\n")
184 unique_mapping_train <- data.frame(
185   original_class = classes,
186   numeric_class = 0:(length(classes)-1),
187   count = as.vector(table(train_df$class))
188 )
189 print(unique_mapping_train)
190
191 # Calculate class weights using the train dataset
192 n_samples <- nrow(train_df)
193 n_classes <- length(classes)
194
195 # Get class frequencies from train dataset
196 class_counts <- table(train_df$class_numeric)
197
198 # Calculate balanced weights
199 class_weights <- n_samples / (n_classes * class_counts)
200
201 # Convert to a named list with numeric indices as names
202 class_weights <- as.list(class_weights)
203 names(class_weights) <- as.character(0:(n_classes-1))
204
205 # Print diagnostics
206 cat("\nClass distribution in full dataset:\n")
207 print(table(full_df$class))
208 cat("\nClass distribution in training dataset:\n")
209 print(table(train_df$class))
210
211 cat("\nFinal Class Weights (based on train dataset):\n")
212 print(class_weights)
213
214 # Validation steps
215 cat("\nValidation:\n")
216 cat("Total samples in full dataset:", n_samples, "\n")
217 cat("Sum of class counts:", sum(class_counts), "\n")
218 cat("Number of classes:", n_classes, "\n")
219 cat("Maximum weight ratio:", max(unlist(class_weights)) /
220     min(unlist(class_weights)), "\n\n")
221
222 # Additional validation checks
223 cat("Verification of weight calculation:\n")
224 for(i in 0:(n_classes-1)) {
225   class_name <- classes[i+1]
226   weight <- class_weights[[as.character(i)]]

```

```

226 count <- class_counts[as.character(i)]
227 cat(sprintf("%s: count=%d, weight=%.4f, count*weight=%.4f\n",
228       class_name, count, weight, count*weight))
229 }
230
231 # Create a data generator with augmentation (no augmentation on validation set)
232 datagen_train <- image_data_generator(
233   rescale = 1/255,
234   rotation_range = 40, # Increased rotation range for more variability
235   width_shift_range = 0.4, # Increased shift range for more variability
236   height_shift_range = 0.4, # Increased shift range for more variability
237   shear_range = 0.4, # Increased shear range for more variability
238   zoom_range = 0.4, # Increased zoom range for more variability
239   horizontal_flip = TRUE
240 )
241
242 datagen_val <- image_data_generator(
243   rescale = 1/255
244 )
245
246 datagen_test <- image_data_generator(
247   rescale = 1/255
248 )
249
250 # Custom generator for training and validation (no augmentation on validation set)
251 train_generator <- flow_images_from_dataframe(
252   train_df,
253   directory = data_dir,
254   x_col = "image_filename",
255   y_col = "class",
256   target_size = c(img_height, img_width),
257   batch_size = batch_size,
258   class_mode = "categorical",
259   generator = datagen_train
260 )
261
262 val_generator <- flow_images_from_dataframe(
263   val_df,
264   directory = data_dir,
265   x_col = "image_filename",
266   y_col = "class",
267   target_size = c(img_height, img_width),
268   batch_size = batch_size,
269   class_mode = "categorical",
270   generator = datagen_val,
271   shuffle = FALSE,
272 )
273
274 test_generator <- flow_images_from_dataframe(
275   test_df,
276   directory = data_dir,
277   x_col = "image_filename",
278   y_col = "class",
279   target_size = c(img_height, img_width),
280   batch_size = batch_size,
281   class_mode = "categorical",
282   generator = datagen_test,
283   shuffle = FALSE,
284 )

```

```

285
286 # Initialize an empty vector to store the learning rate values
287 lr_values <- numeric(0)
288
289 # Custom callback to track the learning rate during training
290 callback_lr_tracker <- callback_lambda(
291   on_epoch_begin = function(epoch, logs) {
292     lr_values <- c(lr_values, as.numeric(model$optimizer$lr)) # Store learning
293     # rate at the start of each epoch
294   }
295 )
296
297 # Train the model with class weights and validation
298 history <- model %>% fit(
299   train_generator,
300   epochs = epochs,
301   steps_per_epoch = nrow(train_df) / batch_size,
302   validation_data = val_generator,
303   validation_steps = nrow(val_df) / batch_size,
304   class_weight = class_weights,
305   callbacks = list(
306     callback_early_stopping(monitor = "val_loss", patience = 8,
307     restore_best_weights = TRUE), # Increased patience
308     callback_reduce_lr_on_plateau(monitor = "val_loss", factor = 0.1, patience =
309     6),
310     callback_model_checkpoint("best_model.h5", monitor = "val_loss",
311     save_best_only = TRUE), # Save the best model
312     callback_lr_tracker
313   )
314 )
315
316 # Retrieve the class labels as defined in the test generator
317 class_indices <- test_generator$class_indices # This is a named list
318 # Reverse the class_indices to map indices back to labels
319 index_to_label <- names(class_indices)[order(unlist(class_indices))]
320
321 # Evaluate the model on the test data
322 evaluation_metrics <- model %>% evaluate(
323   test_generator,
324   steps = nrow(test_df) / batch_size
325 )
326
327 # Extract the evaluation metrics (test loss and accuracy)
328 test_loss <- evaluation_metrics[1]
329 test_accuracy <- evaluation_metrics[2]
330
331 cat("Test Loss:", test_loss, "\n")
332 cat("Test Accuracy:", test_accuracy, "\n")
333
334 # Calculate the number of steps for test predictions
335 steps <- ceiling(nrow(test_df) / batch_size)
336
337 # Make predictions on the test set
338 predictions <- model %>% predict(test_generator, steps = steps)
339
340 # Map the predicted indices to class labels using index_to_label
341 predicted_classes <- factor(apply(predictions, 1, function(x)
342   index_to_label[which.max(x)]), levels = index_to_label)

```

```

339
340 # Get true labels from the test data frame (already in factor form)
341 true_classes <- factor(test_df$class, levels = index_to_label)
342
343 # Calculate the confusion matrix
344 conf_matrix <- confusionMatrix(predicted_classes, true_classes)
345
346 # Print confusion matrix details
347 print(conf_matrix)
348
349 cat("\nClass-Specific Statistics:\n")
350 print(conf_matrix$byClass)
351
352 cat("\nConfusion Matrix Mode:\n")
353 print(conf_matrix$mode)
354
355 cat("\nAdditional Arguments Passed to confusionMatrix():\n")
356 print(conf_matrix$dots)
357
358 # Extract the confusion matrix table as a data frame
359 cm_table <- as.data.frame(as.table(conf_matrix))
360
361 # Rename the columns for easier reference
362 colnames(cm_table) <- c("Predicted", "Actual", "Count")
363
364 # Plot the confusion matrix as a heatmap using ggplot2
365 ggplot(cm_table, aes(x = Predicted, y = Actual, fill = Count)) +
366   geom_tile() +
367   scale_fill_gradient(low = "white", high = "blue") +
368   theme_minimal() +
369   labs(title = "Confusion Matrix Heatmap",
370       x = "Predicted Class",
371       y = "Actual Class") +
372   theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
373   geom_text(aes(label = Count), color = "black", size = 5) # Add text labels
374   with the count in each tile
374
375 plot(history)
376
377 library(ggplot2)
378
379 # Create a data frame for plotting
380 history_df <- data.frame(
381   epoch = 1:length(history$metrics$loss),
382   loss = history$metrics$loss,
383   val_loss = history$metrics$val_loss,
384   accuracy = history$metrics$accuracy,
385   val_accuracy = history$metrics$val_accuracy,
386   lr = lr_values # Add the learning rate history
387 )
388
389 # Create history_df to include loss, accuracy, and learning rate values
390 history_df <- data.frame(
391   epoch = 1:length(history$metrics$loss),
392   loss = history$metrics$loss,
393   val_loss = history$metrics$val_loss,
394   accuracy = history$metrics$accuracy,
395   val_accuracy = history$metrics$val_accuracy,
396   lr = lr_values # Add the learning rate history

```

```

397 )
398
399 library(ggplot2)
400 library(patchwork)
401 library(scales) # for trans_format and log_breaks
402
403 # Linear scale plots
404 plot1 <- ggplot(history_df, aes(x = epoch)) +
405   geom_line(aes(y = loss, color = 'Training Loss')) +
406   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
407   geom_line(aes(y = accuracy * 100, color = 'Accuracy (%)')) +
408   scale_y_continuous(sec.axis = sec_axis(~./100, name = "Accuracy (%)" )) +
409   labs(title = "All Metrics (Linear Scale)", x = "Epoch", y = "Value") +
410   theme_minimal() +
411   scale_color_manual(values = c("blue", "red", "green"))
412
413 # Logarithmic scale plots with custom log base labels
414 plot2 <- ggplot(history_df, aes(x = epoch)) +
415   geom_line(aes(y = loss, color = 'Training Loss')) +
416   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
417   geom_line(aes(y = accuracy * 100, color = 'Accuracy (%)')) +
418   scale_y_log10(labels = trans_format("log10", math_format(10^.x))) +
419   labs(title = "All Metrics (Logarithmic Scale)", x = "Epoch", y = "Value (log
420     scale)") +
421   theme_minimal() +
422   scale_color_manual(values = c("blue", "red", "green"))
423
424 # Loss linear scale
425 plot3 <- ggplot(history_df, aes(x = epoch)) +
426   geom_line(aes(y = loss, color = 'Training Loss')) +
427   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
428   labs(title = "Losses (Linear Scale)", x = "Epoch", y = "Loss") +
429   theme_minimal() +
430   scale_color_manual(values = c("blue", "red"))
431
432 # Loss logarithmic scale with custom log base labels
433 plot4 <- ggplot(history_df, aes(x = epoch)) +
434   geom_line(aes(y = loss, color = 'Training Loss')) +
435   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
436   scale_y_log10(labels = trans_format("log10", math_format(10^.x))) +
437   labs(title = "Losses (Logarithmic Scale)", x = "Epoch", y = "Loss (log scale)")
438   +
439   theme_minimal() +
440   scale_color_manual(values = c("blue", "red"))
441
442 # Combine all plots into a single grid
443 combined_plot <- (plot1 | plot2) / (plot3 | plot4) +
444   plot_annotation(title = "Training Progress")
445
446 # Print combined plot
447 print(combined_plot)
448
449 # Get final training and validation accuracy from the history object
450 final_train_accuracy <- history$metrics$accuracy[length(history$metrics$accuracy)]
451 final_val_accuracy <-
452   history$metrics$val_accuracy[length(history$metrics$val_accuracy)]
453
454 # Extract the final training loss from the history object
455 final_train_loss <- history$metrics$loss[length(history$metrics$loss)]

```

```

453 final_val_loss <- history$metrics$loss[length(history$metrics$val_loss)]
454
455 # Print the final training loss
456 cat("Final Training Loss:", final_train_loss, "\n")
457 cat("Final Validation Loss:", final_val_loss, "\n")
458 cat("Final Train Accuracy:", final_train_accuracy, "\n")
459 cat("Final Validation Accuracy:", final_val_accuracy, "\n")
460
461 # Evaluating classification cost
462 # Define the cost matrix (replace with your actual costs)
463 cost_matrix <- matrix(
464   c(0, 10, 20, 30,    # Benign -> Benign, Early, Pre, Pro
465     10, 0, 15, 25,    # Early -> Benign, Early, Pre, Pro
466     20, 15, 0, 10,    # Pre -> Benign, Early, Pre, Pro
467     30, 25, 10, 0),  # Pro -> Benign, Early, Pre, Pro
468   nrow = 4, byrow = TRUE
469 )
470 rownames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
471 colnames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
472
473 # Extract confusion matrix data
474 cm_table <- conf_matrix$table
475 classes <- rownames(cm_table)
476
477 visualize_cost_confusion_matrix <- function(conf_matrix, cost_matrix) {
478   # Extract the confusion matrix from caret's result
479   cm_table <- conf_matrix$table
480
481   # Create cost-weighted matrix
482   cost_weighted_cm <- matrix(0,
483                             nrow = nrow(cost_matrix),
484                             ncol = ncol(cost_matrix))
485
486   for (i in 1:nrow(cm_table)) {
487     for (j in 1:ncol(cm_table)) {
488       if (i != j) {
489         cost_weighted_cm[i, j] <- cm_table[i, j] * cost_matrix[i, j]
490       }
491     }
492   }
493
494   # Set row and column names
495   rownames(cost_weighted_cm) <- rownames(cost_matrix)
496   colnames(cost_weighted_cm) <- colnames(cost_matrix)
497
498   return(cost_weighted_cm)
499 }
500
501 # Usage
502 cost_weighted_matrix <- visualize_cost_confusion_matrix(conf_matrix, cost_matrix)
503 cat("\nCost-Weighted Misclassification Matrix:\n")
504 print(cost_weighted_matrix)
505
506 # Initialize total misclassification cost
507 total_cost <- 0
508
509 # Calculate total misclassification cost by iterating over confusion matrix
      entries
510 for (i in 1:nrow(cm_table)) {

```

```

511  for (j in 1:ncol(cm_table)) {
512    if (i != j) { # Misclassifications (i != j)
513      total_cost <- total_cost + cm_table[i, j] * cost_matrix[i, j]
514    }
515  }
516}
517
518 # Output the total misclassification cost
519 cat("Total Misclassification Cost:", total_cost, "\n")
520
521 # Calculate average cost per sample
522 total_samples <- sum(cm_table)
523 avg_cost_per_sample <- total_cost / total_samples
524 cat("Average Cost Per Sample:", avg_cost_per_sample, "\n")
525
526 # Calculate cost-sensitive accuracy
527 # Compute the maximum possible cost (if all instances were misclassified to the
# most costly class)
528 max_possible_cost <- sum(cost_matrix) * (total_samples / nrow(cost_matrix))
529 cost_sensitive_accuracy <- 1 - (total_cost / max_possible_cost)
530 cat("Cost-Sensitive Accuracy:", cost_sensitive_accuracy, "\n")
531
532 # Using cost matrix for cost-weighted F1 calculation
533 f1_scores <- conf_matrix$byClass[, "F1"]
534 weights <- 1 / (diag(cost_matrix) + 1) # Add 1 to avoid divide-by-zero
535 weights <- weights / sum(weights)
536 weighted_f1 <- sum(f1_scores * weights, na.rm = TRUE)
537 cat("Weighted F1-score", weighted_f1, "\n")
538
539 # Save the trained model as an HDF5 file
540 save_model_hdf5(model, "vgg16_model.h5")

```

### 9.1.3 DenseNet-201

```

1 # Load necessary libraries
2 library(keras)
3 library(tensorflow)
4 library(dplyr)
5 library(tidyverse)
6 library(caret)
7
8 # Set the path to your dataset
9 data_dir <- "/kaggle/input/leukemia-images/Original"
10
11 # Define image size and parameters for DenseNet-201
12 img_height <- 224
13 img_width <- 224
14 batch_size <- 32
15 epochs <- 220
16
17 # Gather file paths and corresponding labels
18 file_paths <- c()
19 class_labels <- c()
20
21 for (class in c("Benign", "Early", "Pre", "Pro")) {
22   class_path <- file.path(data_dir, class)
23   files <- list.files(class_path, full.names = TRUE)
24   unique_files <- unique(files)

```

```

25  file_paths <- c(file_paths, unique_files)
26  class_labels <- c(class_labels, rep(class, length(unique_files)))
27 }
28
29 # Create DataFrame with the necessary structure
30 full_df <- data.frame(
31   image_filename = file_paths,
32   class = class_labels,
33   stringsAsFactors = FALSE
34 )
35
36 # Stratified split into training (64%), validation (20%), and test (16%) sets
37 set.seed(123)
38 trainIndex <- createDataPartition(full_df$class, p = 0.64, list = FALSE)
39 train_df <- full_df[trainIndex, ]
40 remaining_df <- full_df[-trainIndex, ]
41
42 # Shuffle only train_df
43 train_df <- train_df[sample(nrow(train_df)), ]
44
45 # Split the remaining 36% into validation (20%) and test (16%)
46 validationIndex <- createDataPartition(remaining_df$class, p = 20 / (20 + 16),
47   list = FALSE)
47 val_df <- remaining_df[validationIndex, ]
48 test_df <- remaining_df[-validationIndex, ]
49
50 # Load DenseNet-201 model from TensorFlow
51 base_model <- tf$keras$applications$DenseNet201(
52   weights = 'imagenet',
53   include_top = FALSE,
54   input_shape = as.integer(c(img_height, img_width, 3))
55 )
56
57 # Freeze most of the layers of the base model, unfreezing the top few for
58 # fine-tuning
59 for (layer in base_model$layers[1:500]) {
60   layer$trainable <- FALSE
61 }
62 for (layer in base_model$layers[501:length(base_model$layers)]) {
63   layer$trainable <- TRUE
64 }
65
66 # Define the input and segmentation layer, as before
67 input_image <- layer_input(shape = c(img_height, img_width, 3))
68
69 # Define the color threshold for purple in HSV space
70 lower_purple <- c(0.7, 0.5, 0.4)
71 upper_purple <- c(0.8, 1.0, 1.0)
72
73 # Convert the input image to HSV and create a binary mask for purple regions
74 input_image_hsv <- tf$image$rgb_to_hsv(input_image)
75 mask_purple <- tf$math$logical_and(
76   input_image_hsv[, , 1] >= lower_purple[1], input_image_hsv[, , 1] <=
77     upper_purple[1]
78 )
79 mask_purple <- tf$math$logical_and(mask_purple, input_image_hsv[, , 2] >=
80   lower_purple[2])
81 mask_purple <- tf$math$logical_and(mask_purple, input_image_hsv[, , 3] >=
82   lower_purple[3])

```

```

79 mask_purple <- tf$cast(mask_purple, dtype = tf$float32)
80 mask_purple <- tf$expand_dims(mask_purple, axis = as.integer(-1))
81 mask_purple <- tf$tile(mask_purple, multiples = as.integer(c(1, 1, 1, 3)))
82 segmented_image <- tf$math$multiply(input_image, mask_purple)
83
84 # Contracting Path (Encoder) and Expansive Path (Decoder) as before
85 conv1 <- layer_conv_2d(segmented_image, filters = 32, kernel_size = c(3, 3),
86   activation = "relu", padding = "same")
86 conv1 <- layer_conv_2d(conv1, filters = 32, kernel_size = c(3, 3), activation =
87   "relu", padding = "same")
87 pool1 <- layer_max_pooling_2d(conv1, pool_size = c(2, 2))
88
89 conv2 <- layer_conv_2d(pool1, filters = 64, kernel_size = c(3, 3), activation =
90   "relu", padding = "same")
90 conv2 <- layer_conv_2d(conv2, filters = 64, kernel_size = c(3, 3), activation =
91   "relu", padding = "same")
91 pool2 <- layer_max_pooling_2d(conv2, pool_size = c(2, 2))
92
93 conv3 <- layer_conv_2d(pool2, filters = 128, kernel_size = c(3, 3), activation =
94   "relu", padding = "same")
94 conv3 <- layer_conv_2d(conv3, filters = 128, kernel_size = c(3, 3), activation =
95 pool3 <- layer_max_pooling_2d(conv3, pool_size = c(2, 2))
96
97 conv4 <- layer_conv_2d(pool3, filters = 256, kernel_size = c(3, 3), activation =
98   "relu", padding = "same")
98 conv4 <- layer_conv_2d(conv4, filters = 256, kernel_size = c(3, 3), activation =
99   "relu", padding = "same")
99
100 # Decoder with skip connections
101 upconv3 <- layer_conv_2d_transpose(conv4, filters = 128, kernel_size = c(3, 3),
102   strides = c(2, 2), padding = "same")
102 concat3 <- layer_concatenate(list(upconv3, conv3))
103 conv5 <- layer_conv_2d(concat3, filters = 128, kernel_size = c(3, 3), activation =
104   "relu", padding = "same")
104 conv5 <- layer_conv_2d(conv5, filters = 128, kernel_size = c(3, 3), activation =
105   "relu", padding = "same")
105
106 upconv2 <- layer_conv_2d_transpose(conv5, filters = 64, kernel_size = c(3, 3),
107   strides = c(2, 2), padding = "same")
107 concat2 <- layer_concatenate(list(upconv2, conv2))
108 conv6 <- layer_conv_2d(concat2, filters = 64, kernel_size = c(3, 3), activation =
109   "relu", padding = "same")
109 conv6 <- layer_conv_2d(conv6, filters = 64, kernel_size = c(3, 3), activation =
110   "relu", padding = "same")
110
111 upconv1 <- layer_conv_2d_transpose(conv6, filters = 32, kernel_size = c(3, 3),
112   strides = c(2, 2), padding = "same")
112 concat1 <- layer_concatenate(list(upconv1, conv1))
113 conv7 <- layer_conv_2d(concat1, filters = 32, kernel_size = c(3, 3), activation =
114   "relu", padding = "same")
114 conv7 <- layer_conv_2d(conv7, filters = 32, kernel_size = c(3, 3), activation =
115   "relu", padding = "same")
115
116 # Feature extraction using DenseNet-201
117 features <- base_model(input_image)
118 features <- layer_global_average_pooling_2d(features)
119
120 # Combine segmentation and DenseNet-201 features

```

```

121 segmentation_features <- layer_flatten(conv5)
122 combined_features <- layer_concatenate(list(features, segmentation_features))
123
124 # Dense layers with regularization, batch normalization, and dropout
125 x <- layer_dense(combined_features, units = 512, activation = "relu",
126   kernel_regularizer = regularizer_l2(0.01))
127 x <- layer_batch_normalization(x)
128 x <- layer_dropout(x, rate = 0.7)
129
130 x <- layer_dense(x, units = 256, activation = "relu", kernel_regularizer =
131   regularizer_l2(0.01))
132 x <- layer_batch_normalization(x)
133 x <- layer_dropout(x, rate = 0.7)
134
135 x <- layer_dense(x, units = 128, activation = "relu", kernel_regularizer =
136   regularizer_l2(0.01))
137 x <- layer_batch_normalization(x)
138 x <- layer_dropout(x, rate = 0.7)
139
140 output <- layer_dense(x, units = length(unique(full_df$class)), activation =
141   "softmax")
142
143 # Create the final model
144 model <- keras_model(inputs = input_image, outputs = output)
145
146 # Compile the model with adjusted learning rate
147 model %>% compile(
148   loss = loss_categorical_crossentropy(),
149   optimizer = optimizer_adam(learning_rate = 1e-4),
150   metrics = c("accuracy"))
151
152 print(model)
153
154 # Define the class-to-index mapping
155 classes <- c("Benign", "Early", "Pre", "Pro")
156 class_to_index <- setNames(0:(length(classes) - 1), classes)
157
158 # Add numeric columns to both full and training DataFrames
159 full_df$class_numeric <- as.numeric(factor(full_df$class, levels = classes)) - 1
160 train_df$class_numeric <- as.numeric(factor(train_df$class, levels = classes)) - 1
161
162 # Verify the class-to-index mapping
163 cat("Class-to-Index Mapping:\n")
164 print(class_to_index)
165
166 # Create improved mapping displays with counts
167 cat("\nMappings and counts in full dataset:\n")
168 unique_mapping_full <- data.frame(
169   original_class = classes,
170   numeric_class = 0:(length(classes)-1),
171   count = as.vector(table(full_df$class)))
172
173 print(unique_mapping_full)
174
175 cat("\nMappings and counts in training dataset:\n")
176 unique_mapping_train <- data.frame(
177   original_class = classes,
178   numeric_class = 0:(length(classes)-1),

```

```

176 count = as.vector(table(train_df$class))
177 )
178 print(unique_mapping_train)
179
180 # Calculate class weights using the train dataset
181 n_samples <- nrow(train_df)
182 n_classes <- length(classes)
183
184 # Get class frequencies from train dataset
185 class_counts <- table(train_df$class_numeric)
186
187 # Calculate balanced weights
188 class_weights <- n_samples / (n_classes * class_counts)
189
190 # Convert to a named list with numeric indices as names
191 class_weights <- as.list(class_weights)
192 names(class_weights) <- as.character(0:(n_classes-1))
193
194 # Print diagnostics
195 cat("\nClass distribution in full dataset:\n")
196 print(table(full_df$class))
197 cat("\nClass distribution in training dataset:\n")
198 print(table(train_df$class))
199
200 cat("\nFinal Class Weights (based on train dataset):\n")
201 print(class_weights)
202
203 # Validation steps
204 cat("\nValidation:\n")
205 cat("Total samples in full dataset:", n_samples, "\n")
206 cat("Sum of class counts:", sum(class_counts), "\n")
207 cat("Number of classes:", n_classes, "\n")
208 cat("Maximum weight ratio:", max(unlist(class_weights)) /
      min(unlist(class_weights)), "\n\n")
209
210 # Additional validation checks
211 cat("Verification of weight calculation:\n")
212 for(i in 0:(n_classes-1)) {
213   class_name <- classes[i+1]
214   weight <- class_weights[[as.character(i)]]
215   count <- class_counts[as.character(i)]
216   cat(sprintf("%s: count=%d, weight=%.4f, count*weight=%.4f\n",
217           class_name, count, weight, count*weight))
218 }
219
220 # Create a data generator with augmentation (no augmentation on validation set)
221 datagen_train <- image_data_generator(
222   rescale = 1/255,
223   rotation_range = 40, # Increased rotation range for more variability
224   width_shift_range = 0.4, # Increased shift range for more variability
225   height_shift_range = 0.4, # Increased shift range for more variability
226   shear_range = 0.4, # Increased shear range for more variability
227   zoom_range = 0.4, # Increased zoom range for more variability
228   horizontal_flip = TRUE
229 )
230
231 datagen_val <- image_data_generator(
232   rescale = 1/255
233 )

```

```

234
235 datagen_test <- image_data_generator(
236   rescale = 1/255
237 )
238
239 # Custom generator for training and validation (no augmentation on validation set)
240 train_generator <- flow_images_from_dataframe(
241   train_df,
242   directory = data_dir,
243   x_col = "image_filename",
244   y_col = "class",
245   target_size = c(img_height, img_width),
246   batch_size = batch_size,
247   class_mode = "categorical",
248   generator = datagen_train
249 )
250
251 val_generator <- flow_images_from_dataframe(
252   val_df,
253   directory = data_dir,
254   x_col = "image_filename",
255   y_col = "class",
256   target_size = c(img_height, img_width),
257   batch_size = batch_size,
258   class_mode = "categorical",
259   generator = datagen_val,
260   shuffle = FALSE,
261 )
262
263 test_generator <- flow_images_from_dataframe(
264   test_df,
265   directory = data_dir,
266   x_col = "image_filename",
267   y_col = "class",
268   target_size = c(img_height, img_width),
269   batch_size = batch_size,
270   class_mode = "categorical",
271   generator = datagen_test,
272   shuffle = FALSE,
273 )
274
275 # Initialize an empty vector to store the learning rate values
276 lr_values <- numeric(0)
277
278 # Custom callback to track the learning rate during training
279 callback_lr_tracker <- callback_lambda(
280   on_epoch_begin = function(epoch, logs) {
281     lr_values <- c(lr_values, as.numeric(model$optimizer$lr)) # Store learning
282     rate at the start of each epoch
283   }
284 )
285
286 # Train the model with class weights and validation
287 history <- model %>% fit(
288   train_generator,
289   epochs = epochs,
290   steps_per_epoch = nrow(train_df) / batch_size,
291   validation_data = val_generator,

```

```

292 validation_steps = nrow(val_df) / batch_size,
293 class_weight = class_weights,
294 callbacks = list(
295   callback_early_stopping(monitor = "val_loss", patience = 5,
296   restore_best_weights = TRUE), # Increased patience
297   callback_reduce_lr_on_plateau(monitor = "val_loss", factor = 0.1, patience =
298   3),
299   callback_model_checkpoint("best_model.h5", monitor = "val_loss",
300   save_best_only = TRUE), # Save the best model
301   callback_lr_tracker
302 )
303
304 # Retrieve the class labels as defined in the test generator
305 class_indices <- test_generator$class_indices # This is a named list
306 # Reverse the class_indices to map indices back to labels
307 index_to_label <- names(class_indices)[order(unlist(class_indices))]
308
309 # Evaluate the model on the test data
310 evaluation_metrics <- model %>% evaluate(
311   test_generator,
312   steps = nrow(test_df) / batch_size
313 )
314
315 # Extract the evaluation metrics (test loss and accuracy)
316 test_loss <- evaluation_metrics[1]
317 test_accuracy <- evaluation_metrics[2]
318
319 cat("Test Loss:", test_loss, "\n")
320 cat("Test Accuracy:", test_accuracy, "\n")
321
322 # Calculate the number of steps for test predictions
323 steps <- ceiling(nrow(test_df) / batch_size)
324
325 # Make predictions on the test set
326 predictions <- model %>% predict(test_generator, steps = steps)
327
328 # Map the predicted indices to class labels using index_to_label
329 predicted_classes <- factor(apply(predictions, 1, function(x)
330   index_to_label[which.max(x)]), levels = index_to_label)
331
332 # Get true labels from the test data frame (already in factor form)
333 true_classes <- factor(test_df$class, levels = index_to_label)
334
335 # Calculate the confusion matrix
336 conf_matrix <- confusionMatrix(predicted_classes, true_classes)
337
338 # Print confusion matrix details
339 print(conf_matrix)
340 cat("\nClass-Specific Statistics:\n")
341 print(conf_matrix$byClass)
342
343 cat("\nConfusion Matrix Mode:\n")
344 print(conf_matrix$mode)
345
346 cat("\nAdditional Arguments Passed to confusionMatrix():\n")

```

```

347 print(conf_matrix$dots)
348
349 # Extract the confusion matrix table as a data frame
350 cm_table <- as.data.frame(as.table(conf_matrix))
351
352 # Rename the columns for easier reference
353 colnames(cm_table) <- c("Predicted", "Actual", "Count")
354
355 # Plot the confusion matrix as a heatmap using ggplot2
356 ggplot(cm_table, aes(x = Predicted, y = Actual, fill = Count)) +
357   geom_tile() +
358   scale_fill_gradient(low = "white", high = "blue") +
359   theme_minimal() +
360   labs(title = "Confusion Matrix Heatmap",
361       x = "Predicted Class",
362       y = "Actual Class") +
363   theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
364   geom_text(aes(label = Count), color = "black", size = 5) # Add text labels
365   with the count in each tile
366
367 plot(history)
368
369 library(ggplot2)
370
371 # Create a data frame for plotting
372 history_df <- data.frame(
373   epoch = 1:length(history$metrics$loss),
374   loss = history$metrics$loss,
375   val_loss = history$metrics$val_loss,
376   accuracy = history$metrics$accuracy,
377   val_accuracy = history$metrics$val_accuracy,
378   lr = lr_values # Add the learning rate history
379 )
380
381 # Create history_df to include loss, accuracy, and learning rate values
382 history_df <- data.frame(
383   epoch = 1:length(history$metrics$loss),
384   loss = history$metrics$loss,
385   val_loss = history$metrics$val_loss,
386   accuracy = history$metrics$accuracy,
387   val_accuracy = history$metrics$val_accuracy,
388   lr = lr_values # Add the learning rate history
389 )
390
391 library(ggplot2)
392 library(patchwork)
393 library(scales) # for trans_format and log_breaks
394
395 # Linear scale plots
396 plot1 <- ggplot(history_df, aes(x = epoch)) +
397   geom_line(aes(y = loss, color = 'Training Loss')) +
398   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
399   geom_line(aes(y = accuracy * 100, color = 'Accuracy (%)')) +
400   scale_y_continuous(sec.axis = sec_axis(~./100, name = "Accuracy (%))) +
401   labs(title = "All Metrics (Linear Scale)", x = "Epoch", y = "Value") +
402   theme_minimal() +
403   scale_color_manual(values = c("blue", "red", "green"))
404
405 # Logarithmic scale plots with custom log base labels

```

```

405 plot2 <- ggplot(history_df, aes(x = epoch)) +
406   geom_line(aes(y = loss, color = 'Training Loss')) +
407   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
408   geom_line(aes(y = accuracy * 100, color = 'Accuracy (%)')) +
409   scale_y_log10(labels = trans_format("log10", math_format(10^.x))) +
410   labs(title = "All Metrics (Logarithmic Scale)", x = "Epoch", y = "Value (log
411     scale)") +
412   theme_minimal() +
413   scale_color_manual(values = c("blue", "red", "green"))
414 
415 # Loss linear scale
416 plot3 <- ggplot(history_df, aes(x = epoch)) +
417   geom_line(aes(y = loss, color = 'Training Loss')) +
418   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
419   labs(title = "Losses (Linear Scale)", x = "Epoch", y = "Loss") +
420   theme_minimal() +
421   scale_color_manual(values = c("blue", "red"))
422 
423 # Loss logarithmic scale with custom log base labels
424 plot4 <- ggplot(history_df, aes(x = epoch)) +
425   geom_line(aes(y = loss, color = 'Training Loss')) +
426   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
427   scale_y_log10(labels = trans_format("log10", math_format(10^.x))) +
428   labs(title = "Losses (Logarithmic Scale)", x = "Epoch", y = "Loss (log scale)")
429   +
430   theme_minimal() +
431   scale_color_manual(values = c("blue", "red"))
432 
433 # Combine all plots into a single grid
434 combined_plot <- (plot1 | plot2) / (plot3 | plot4) +
435   plot_annotation(title = "Training Progress")
436 
437 # Print combined plot
438 print(combined_plot)
439 
440 # Get final training and validation accuracy from the history object
441 final_train_accuracy <- history$metrics$accuracy[length(history$metrics$accuracy)]
442 final_val_accuracy <-
443   history$metrics$val_accuracy[length(history$metrics$val_accuracy)]
444 
445 # Extract the final training loss from the history object
446 final_train_loss <- history$metrics$loss[length(history$metrics$loss)]
447 final_val_loss <- history$metrics$loss[length(history$metrics$val_loss)]
448 
449 # Print the final training loss
450 cat("Final Training Loss:", final_train_loss, "\n")
451 cat("Final Validation Loss:", final_val_loss, "\n")
452 cat("Final Train Accuracy:", final_train_accuracy, "\n")
453 cat("Final Validation Accuracy:", final_val_accuracy, "\n")
454 
455 # Evaluating classification cost
456 # Define the cost matrix (replace with your actual costs)
457 cost_matrix <- matrix(
458   c(0, 10, 20, 30,    # Benign -> Benign, Early, Pre, Pro
459     10, 0, 15, 25,    # Early -> Benign, Early, Pre, Pro
460     20, 15, 0, 10,    # Pre -> Benign, Early, Pre, Pro
461     30, 25, 10, 0),   # Pro -> Benign, Early, Pre, Pro
462   nrow = 4, byrow = TRUE
463 )

```

```

461 rownames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
462 colnames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
463
464 # Extract confusion matrix data
465 cm_table <- conf_matrix$table
466 classes <- rownames(cm_table)
467
468 visualize_cost_confusion_matrix <- function(conf_matrix, cost_matrix) {
469   # Extract the confusion matrix from caret's result
470   cm_table <- conf_matrix$table
471
472   # Create cost-weighted matrix
473   cost_weighted_cm <- matrix(0,
474                               nrow = nrow(cost_matrix),
475                               ncol = ncol(cost_matrix))
476
477   for (i in 1:nrow(cm_table)) {
478     for (j in 1:ncol(cm_table)) {
479       if (i != j) {
480         cost_weighted_cm[i, j] <- cm_table[i, j] * cost_matrix[i, j]
481       }
482     }
483   }
484
485   # Set row and column names
486   rownames(cost_weighted_cm) <- rownames(cost_matrix)
487   colnames(cost_weighted_cm) <- colnames(cost_matrix)
488
489   return(cost_weighted_cm)
490 }
491
492 # Usage
493 cost_weighted_matrix <- visualize_cost_confusion_matrix(conf_matrix, cost_matrix)
494 cat("\nCost-Weighted Misclassification Matrix:\n")
495 print(cost_weighted_matrix)
496
497 # Initialize total misclassification cost
498 total_cost <- 0
499
500 # Calculate total misclassification cost by iterating over confusion matrix
      entries
501 for (i in 1:nrow(cm_table)) {
502   for (j in 1:ncol(cm_table)) {
503     if (i != j) { # Misclassifications (i != j)
504       total_cost <- total_cost + cm_table[i, j] * cost_matrix[i, j]
505     }
506   }
507 }
508
509 # Output the total misclassification cost
510 cat("Total Misclassification Cost:", total_cost, "\n")
511
512 # Calculate average cost per sample
513 total_samples <- sum(cm_table)
514 avg_cost_per_sample <- total_cost / total_samples
515 cat("Average Cost Per Sample:", avg_cost_per_sample, "\n")
516
517 # Calculate cost-sensitive accuracy

```

```

518 # Compute the maximum possible cost (if all instances were misclassified to the
      most costly class)
519 max_possible_cost <- sum(cost_matrix) * (total_samples / nrow(cost_matrix))
520 cost_sensitive_accuracy <- 1 - (total_cost / max_possible_cost)
521 cat("Cost-Sensitive Accuracy:", cost_sensitive_accuracy, "\n")
522
523 # Using cost matrix for cost-weighted F1 calculation
524 f1_scores <- conf_matrix$byClass[, "F1"]
525 weights <- 1 / (diag(cost_matrix) + 1) # Add 1 to avoid divide-by-zero
526 weights <- weights / sum(weights)
527 weighted_f1 <- sum(f1_scores * weights, na.rm = TRUE)
528 cat("Weighted F1-score", weighted_f1, "\n")
529
530 # Save the trained model as an HDF5 file
531 save_model_hdf5(model, "densenet201_model.h5")

```

#### 9.1.4 ConvNeXt

```

1 # Load necessary libraries
2 library(keras)
3 library(tensorflow)
4 library(dplyr) # For data manipulation
5 library(tidyverse) # For tidyverse functions
6 library(caret) # For confusionMatrix and stratified sampling
7
8 # Set the path to your dataset
9 data_dir <- "/kaggle/input/leukemia-images/Original"
10
11 # Define image size and parameters for ConvNeXt
12 img_height <- 224
13 img_width <- 224
14 batch_size <- 32
15 epochs <- 220
16
17 # Gather file paths and corresponding labels
18 file_paths <- c()
19 class_labels <- c()
20
21 for (class in c("Benign", "Early", "Pre", "Pro")) {
22   class_path <- file.path(data_dir, class)
23   files <- list.files(class_path, full.names = TRUE)
24   unique_files <- unique(files)
25   file_paths <- c(file_paths, unique_files)
26   class_labels <- c(class_labels, rep(class, length(unique_files)))
27 }
28
29 # Create DataFrame with the necessary structure
30 full_df <- data.frame(
31   image_filename = file_paths,
32   class = class_labels,
33   stringsAsFactors = FALSE
34 )
35
36 # Stratified split into training (64%), validation (20%), and test (16%) sets
37 set.seed(123)
38 trainIndex <- createDataPartition(full_df$class, p = 0.64, list = FALSE)
39 train_df <- full_df[trainIndex, ]
40 remaining_df <- full_df[-trainIndex, ]

```

```

41 # Shuffle only train_df
42 train_df <- train_df[sample(nrow(train_df)), ]
43
44 # Split the remaining 36% into validation (20%) and test (16%)
45 validationIndex <- createDataPartition(remaining_df$class, p = 20 / (20 + 16),
46   list = FALSE)
47 val_df <- remaining_df[validationIndex, ]
48 test_df <- remaining_df[-validationIndex, ]
49
50 # Load ConvNeXT model from TensorFlow using tf$keras$applications$ConvNeXtBase
51 base_model <- tf$keras$applications$ConvNeXtBase(
52   weights = 'imagenet',
53   include_top = FALSE,
54   input_shape = as.integer(c(img_height, img_width, 3))
55 )
56
57 # Freeze most of the layers of the base model, but unfreeze the top few layers
58 # for fine-tuning
59 for (layer in base_model$layers[1:100]) {
60   layer$trainable <- FALSE
61 }
62 for (layer in base_model$layers[101:length(base_model$layers)]) {
63   layer$trainable <- TRUE
64 }
65
66 # Set the color threshold for purple in HSV space
67 lower_purple <- c(0.7, 0.5, 0.4) # Adjust as needed
68 upper_purple <- c(0.8, 1.0, 1.0) # Adjust as needed
69
70 # Define the model input layer
71 input_image <- layer_input(shape = c(img_height, img_width, 3))
72
73 # Convert the input image to HSV
74 input_image_hsv <- tf$image$rgb_to_hsv(input_image)
75
76 # Create a binary mask for purple regions (lymphoblasts) based on the HSV
77 # threshold
78 mask_purple <- tf$math$logical_and(
79   input_image_hsv[,,1] >= lower_purple[1], input_image_hsv[,,1] <=
80     upper_purple[1]
81 )
82 mask_purple <- tf$math$logical_and(mask_purple, input_image_hsv[,,2] >=
83   lower_purple[2])
84 mask_purple <- tf$math$logical_and(mask_purple, input_image_hsv[,,3] >=
85   lower_purple[3])
86
87 # Convert logical mask to float and expand dims to match RGB channels
88 mask_purple <- tf$cast(mask_purple, dtype = tf$float32) # Ensure the mask is a
89 # float tensor
90 mask_purple <- tf$expand_dims(mask_purple, axis = as.integer(-1)) # Add a channel
# dimension
91 mask_purple <- tf$tile(mask_purple, multiples = as.integer(c(1, 1, 1, 3))) #
# Repeat for each color channel
92
93 # Apply the mask to the original RGB image
94 segmented_image <- tf$math$multiply(input_image, mask_purple)
95
96 # Pass the segmented image into segmentation block

```

```

91 # Modified U-Net inspired segmentation block with color-based segmentation
92 # Contracting Path (Encoder) with fewer filters
93 conv1 <- layer_conv_2d(segmented_image, filters = 32, kernel_size = c(3, 3),
94   activation = "relu", padding = "same")
95 conv1 <- layer_conv_2d(conv1, filters = 32, kernel_size = c(3, 3), activation =
96   "relu", padding = "same")
97 pool1 <- layer_max_pooling_2d(conv1, pool_size = c(2, 2))
98
99 conv2 <- layer_conv_2d(pool1, filters = 64, kernel_size = c(3, 3), activation =
100  "relu", padding = "same")
101 conv2 <- layer_conv_2d(conv2, filters = 64, kernel_size = c(3, 3), activation =
102  "relu", padding = "same")
103 pool2 <- layer_max_pooling_2d(conv2, pool_size = c(2, 2))
104
105 conv3 <- layer_conv_2d(pool2, filters = 128, kernel_size = c(3, 3), activation =
106  "relu", padding = "same")
107 conv3 <- layer_conv_2d(conv3, filters = 128, kernel_size = c(3, 3), activation =
108  "relu", padding = "same")
109 pool3 <- layer_max_pooling_2d(conv3, pool_size = c(2, 2))
110
111 # Bottleneck layer with fewer filters
112 conv4 <- layer_conv_2d(pool3, filters = 256, kernel_size = c(3, 3), activation =
113  "relu", padding = "same")
114 conv4 <- layer_conv_2d(conv4, filters = 256, kernel_size = c(3, 3), activation =
115  "relu", padding = "same")
116
117 # Expansive Path (Decoder) with skip connections and fewer filters
118 # Upsampling + skip connection from conv3
119 upconv3 <- layer_conv_2d_transpose(conv4, filters = 128, kernel_size = c(3, 3),
120   strides = c(2, 2), padding = "same")
121 concat3 <- layer_concatenate(list(upconv3, conv3))
122 conv5 <- layer_conv_2d(concat3, filters = 128, kernel_size = c(3, 3), activation =
123  "relu", padding = "same")
124 conv5 <- layer_conv_2d(conv5, filters = 128, kernel_size = c(3, 3), activation =
125  "relu", padding = "same")
126
127 # Upsampling + skip connection from conv2
128 upconv2 <- layer_conv_2d_transpose(conv5, filters = 64, kernel_size = c(3, 3),
129   strides = c(2, 2), padding = "same")
130 concat2 <- layer_concatenate(list(upconv2, conv2))
131 conv6 <- layer_conv_2d(concat2, filters = 64, kernel_size = c(3, 3), activation =
132  "relu", padding = "same")
133 conv6 <- layer_conv_2d(conv6, filters = 64, kernel_size = c(3, 3), activation =
134  "relu", padding = "same")
135
136 # Upsampling + skip connection from conv1
137 upconv1 <- layer_conv_2d_transpose(conv6, filters = 32, kernel_size = c(3, 3),
138   strides = c(2, 2), padding = "same")
139 concat1 <- layer_concatenate(list(upconv1, conv1))
140 conv7 <- layer_conv_2d(concat1, filters = 32, kernel_size = c(3, 3), activation =
141  "relu", padding = "same")
142 conv7 <- layer_conv_2d(conv7, filters = 32, kernel_size = c(3, 3), activation =
143  "relu", padding = "same" )
144
145 # Feature extraction using ConvNeXT
146 features <- base_model(input_image)
147 features <- layer_global_average_pooling_2d(features)
148
```

```

133 # Flatten and concatenate segmentation output with ConvNeXT features
134 segmentation_features <- layer_flatten(conv5)
135 combined_features <- layer_concatenate(list(features, segmentation_features))
136
137 # Dense layers with L2 regularization, batch normalization, and increased dropout
138 x <- layer_dense(combined_features, units = 512, activation = "relu",
139   kernel_regularizer = regularizer_l2(0.01))
140 x <- layer_batch_normalization(x)
141 x <- layer_dropout(x, rate = 0.7) # Increased dropout rate to 0.7
142
143 x <- layer_dense(x, units = 256, activation = "relu", kernel_regularizer =
144   regularizer_l2(0.01))
145 x <- layer_batch_normalization(x)
146 x <- layer_dropout(x, rate = 0.7) # Increased dropout rate to 0.7
147
148 x <- layer_dense(x, units = 128, activation = "relu", kernel_regularizer =
149   regularizer_l2(0.01))
150 x <- layer_batch_normalization(x)
151 x <- layer_dropout(x, rate = 0.7) # Increased dropout rate to 0.7
152
153 output <- layer_dense(x, units = length(unique(full_df$class)), activation =
154   "softmax")
155
156 # Create the model
157 model <- keras_model(inputs = input_image, outputs = output)
158
159 # Compile the model with a lower learning rate to prevent overfitting
160 model %>% compile(
161   loss = loss_categorical_crossentropy(),
162   optimizer = optimizer_adam(learning_rate = 1e-4), # Lower learning rate to
163   prevent overfitting
164   metrics = c("accuracy"))
165
166 print(model)
167
168 # Define the class-to-index mapping
169 classes <- c("Benign", "Early", "Pre", "Pro")
170 class_to_index <- setNames(0:(length(classes) - 1), classes)
171
172 # Add numeric columns to both full and training DataFrames
173 full_df$class_numeric <- as.numeric(factor(full_df$class, levels = classes)) - 1
174 train_df$class_numeric <- as.numeric(factor(train_df$class, levels = classes)) - 1
175
176 # Verify the class-to-index mapping
177 cat("Class-to-Index Mapping:\n")
178 print(class_to_index)
179
180 # Create improved mapping displays with counts
181 cat("\nMappings and counts in full dataset:\n")
182 unique_mapping_full <- data.frame(
183   original_class = classes,
184   numeric_class = 0:(length(classes)-1),
185   count = as.vector(table(full_df$class)))
186
187 print(unique_mapping_full)
188
189 cat("\nMappings and counts in training dataset:\n")
190 unique_mapping_train <- data.frame(

```

```

187 original_class = classes,
188 numeric_class = 0:(length(classes)-1),
189 count = as.vector(table(train_df$class))
190 )
191 print(unique_mapping_train)
192
193 # Calculate class weights using the train dataset
194 n_samples <- nrow(train_df)
195 n_classes <- length(classes)
196
197 # Get class frequencies from train dataset
198 class_counts <- table(train_df$class_numeric)
199
200 # Calculate balanced weights
201 class_weights <- n_samples / (n_classes * class_counts)
202
203 # Convert to a named list with numeric indices as names
204 class_weights <- as.list(class_weights)
205 names(class_weights) <- as.character(0:(n_classes-1))
206
207 # Print diagnostics
208 cat("\nClass distribution in full dataset:\n")
209 print(table(full_df$class))
210 cat("\nClass distribution in training dataset:\n")
211 print(table(train_df$class))
212
213 cat("\nFinal Class Weights (based on train dataset):\n")
214 print(class_weights)
215
216 # Validation steps
217 cat("\nValidation:\n")
218 cat("Total samples in full dataset:", n_samples, "\n")
219 cat("Sum of class counts:", sum(class_counts), "\n")
220 cat("Number of classes:", n_classes, "\n")
221 cat("Maximum weight ratio:", max(unlist(class_weights)) /
222     min(unlist(class_weights)), "\n\n")
223
224 # Additional validation checks
225 cat("Verification of weight calculation:\n")
226 for(i in 0:(n_classes-1)) {
227   class_name <- classes[i+1]
228   weight <- class_weights[[as.character(i)]]
229   count <- class_counts[as.character(i)]
230   cat(sprintf("%s: count=%d, weight=%.4f, count*weight=%.4f\n",
231           class_name, count, weight, count*weight))
232 }
233
234 # Create a data generator with augmentation (no augmentation on validation set)
235 datagen_train <- image_data_generator(
236   rescale = 1/255,
237   rotation_range = 40, # Increased rotation range for more variability
238   width_shift_range = 0.4, # Increased shift range for more variability
239   height_shift_range = 0.4, # Increased shift range for more variability
240   shear_range = 0.4, # Increased shear range for more variability
241   zoom_range = 0.4, # Increased zoom range for more variability
242   horizontal_flip = TRUE
243 )
244 datagen_val <- image_data_generator(

```

```

245     rescale = 1/255
246 )
247
248 datagen_test <- image_data_generator(
249   rescale = 1/255
250 )
251
252 # Custom generator for training and validation (no augmentation on validation set)
253 train_generator <- flow_images_from_dataframe(
254   train_df,
255   directory = data_dir,
256   x_col = "image_filename",
257   y_col = "class",
258   target_size = c(img_height, img_width),
259   batch_size = batch_size,
260   class_mode = "categorical",
261   generator = datagen_train
262 )
263
264 val_generator <- flow_images_from_dataframe(
265   val_df,
266   directory = data_dir,
267   x_col = "image_filename",
268   y_col = "class",
269   target_size = c(img_height, img_width),
270   batch_size = batch_size,
271   class_mode = "categorical",
272   generator = datagen_val,
273   shuffle = FALSE,
274 )
275
276 test_generator <- flow_images_from_dataframe(
277   test_df,
278   directory = data_dir,
279   x_col = "image_filename",
280   y_col = "class",
281   target_size = c(img_height, img_width),
282   batch_size = batch_size,
283   class_mode = "categorical",
284   generator = datagen_test,
285   shuffle = FALSE,
286 )
287
288 # Initialize an empty vector to store the learning rate values
289 lr_values <- numeric(0)
290
291 # Custom callback to track the learning rate during training
292 callback_lr_tracker <- callback_lambda(
293   on_epoch_begin = function(epoch, logs) {
294     lr_values <- c(lr_values, as.numeric(model$optimizer$lr)) # Store learning
295     rate at the start of each epoch
296   }
297
298
299 # Train the model with class weights and validation
300 history <- model %>% fit(
301   train_generator,
302   epochs = epochs,

```

```

303 steps_per_epoch = nrow(train_df) / batch_size,
304 validation_data = val_generator,
305 validation_steps = nrow(val_df) / batch_size,
306 class_weight = class_weights,
307 callbacks = list(
308     callback_early_stopping(monitor = "val_loss", patience = 8,
309     restore_best_weights = TRUE), # Increased patience
310     callback_reduce_lr_on_plateau(monitor = "val_loss", factor = 0.1, patience =
311     5),
312     callback_model_checkpoint("best_model.h5", monitor = "val_loss",
313     save_best_only = TRUE), # Save the best model
314     callback_lr_tracker
315 )
316
317 # Retrieve the class labels as defined in the test generator
318 class_indices <- test_generator$class_indices # This is a named list
319 # Reverse the class_indices to map indices back to labels
320 index_to_label <- names(class_indices)[order(unlist(class_indices))]
321
322 # Evaluate the model on the test data
323 evaluation_metrics <- model %>% evaluate(
324     test_generator,
325     steps = nrow(test_df) / batch_size
326 )
327
328 # Extract the evaluation metrics (test loss and accuracy)
329 test_loss <- evaluation_metrics[1]
330 test_accuracy <- evaluation_metrics[2]
331
332 cat("Test Loss:", test_loss, "\n")
333 cat("Test Accuracy:", test_accuracy, "\n")
334
335 # Calculate the number of steps for test predictions
336 steps <- ceiling(nrow(test_df) / batch_size)
337
338 # Make predictions on the test set
339 predictions <- model %>% predict(test_generator, steps = steps)
340
341 # Map the predicted indices to class labels using index_to_label
342 predicted_classes <- factor(apply(predictions, 1, function(x)
343     index_to_label[which.max(x)]), levels = index_to_label)
344
345 # Get true labels from the test data frame (already in factor form)
346 true_classes <- factor(test_df$class, levels = index_to_label)
347
348 # Calculate the confusion matrix
349 conf_matrix <- confusionMatrix(predicted_classes, true_classes)
350
351 # Print confusion matrix details
352 print(conf_matrix)
353 cat("\nClass-Specific Statistics:\n")
354 print(conf_matrix$byClass)
355
356 cat("\nConfusion Matrix Mode:\n")
357 print(conf_matrix$mode)

```

```

358
359 cat("\nAdditional Arguments Passed to confusionMatrix():\n")
360 print(conf_matrix$dots)
361
362 # Extract the confusion matrix table as a data frame
363 cm_table <- as.data.frame(as.table(conf_matrix))
364
365 # Rename the columns for easier reference
366 colnames(cm_table) <- c("Predicted", "Actual", "Count")
367
368 # Plot the confusion matrix as a heatmap using ggplot2
369 ggplot(cm_table, aes(x = Predicted, y = Actual, fill = Count)) +
370   geom_tile() +
371   scale_fill_gradient(low = "white", high = "blue") +
372   theme_minimal() +
373   labs(title = "Confusion Matrix Heatmap",
374        x = "Predicted Class",
375        y = "Actual Class") +
376   theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
377   geom_text(aes(label = Count), color = "black", size = 5) # Add text labels
378   with the count in each tile
379
380 plot(history)
381
382 library(ggplot2)
383
384 # Create a data frame for plotting
385 history_df <- data.frame(
386   epoch = 1:length(history$metrics$loss),
387   loss = history$metrics$loss,
388   val_loss = history$metrics$val_loss,
389   accuracy = history$metrics$accuracy,
390   val_accuracy = history$metrics$val_accuracy,
391   lr = lr_values # Add the learning rate history
392 )
393
394 # Create history_df to include loss, accuracy, and learning rate values
395 history_df <- data.frame(
396   epoch = 1:length(history$metrics$loss),
397   loss = history$metrics$loss,
398   val_loss = history$metrics$val_loss,
399   accuracy = history$metrics$accuracy,
400   val_accuracy = history$metrics$val_accuracy,
401   lr = lr_values # Add the learning rate history
402 )
403
404 library(ggplot2)
405 library(patchwork)
406
407 # Linear scale plots
408 plot1 <- ggplot(history_df, aes(x = epoch)) +
409   geom_line(aes(y = loss, color = 'Training Loss')) +
410   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
411   geom_line(aes(y = accuracy * 100, color = 'Accuracy (%)')) +
412   scale_y_continuous(sec.axis = sec_axis(~./100, name = "Accuracy (%))) +
413   labs(title = "All Metrics (Linear Scale)", x = "Epoch", y = "Value") +
414   theme_minimal() +

```

```

415 scale_color_manual(values = c("blue", "red", "green"))
416
417 # Logarithmic scale plots
418 plot2 <- ggplot(history_df, aes(x = epoch)) +
419   geom_line(aes(y = loss, color = 'Training Loss')) +
420   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
421   geom_line(aes(y = accuracy * 100, color = 'Accuracy (%)')) +
422   scale_y_log10() +
423   labs(title = "All Metrics (Logarithmic Scale)", x = "Epoch", y = "Value (log
424     scale)") +
425   theme_minimal() +
426   scale_color_manual(values = c("blue", "red", "green"))
427
428 # Loss linear scale
429 plot3 <- ggplot(history_df, aes(x = epoch)) +
430   geom_line(aes(y = loss, color = 'Training Loss')) +
431   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
432   labs(title = "Losses (Linear Scale)", x = "Epoch", y = "Loss") +
433   theme_minimal() +
434   scale_color_manual(values = c("blue", "red"))
435
436 # Loss logarithmic scale
437 plot4 <- ggplot(history_df, aes(x = epoch)) +
438   geom_line(aes(y = loss, color = 'Training Loss')) +
439   geom_line(aes(y = val_loss, color = 'Validation Loss')) +
440   scale_y_log10() +
441   labs(title = "Losses (Logarithmic Scale)", x = "Epoch", y = "Loss (log scale)") +
442   theme_minimal() +
443   scale_color_manual(values = c("blue", "red"))
444
445 # Combine all plots into a single grid
446 combined_plot <- (plot1 | plot2) / (plot3 | plot4) +
447   plot_annotation(title = "Training Progress")
448
449 # Print combined plot
450 print(combined_plot)
451
452 # Get final training and validation accuracy from the history object
453 final_train_accuracy <- history$metrics$accuracy[length(history$metrics$accuracy)]
454 final_val_accuracy <-
455   history$metrics$val_accuracy[length(history$metrics$val_accuracy)]
456
457 # Extract the final training loss from the history object
458 final_train_loss <- history$metrics$loss[length(history$metrics$loss)]
459 final_val_loss <- history$metrics$loss[length(history$metrics$val_loss)]
460
461 # Print the final training loss
462 cat("Final Training Loss:", final_train_loss, "\n")
463 cat("Final Validation Loss:", final_val_loss, "\n")
464 cat("Final Train Accuracy:", final_train_accuracy, "\n")
465 cat("Final Validation Accuracy:", final_val_accuracy, "\n")
466
467 # Evaluating classification cost
468 # Define the cost matrix (replace with your actual costs)
469 cost_matrix <- matrix(
470   c(0, 10, 20, 30,    # Benign -> Benign, Early, Pre, Pro
471     10, 0, 15, 25,    # Early -> Benign, Early, Pre, Pro
472     20, 15, 0, 10,    # Pre -> Benign, Early, Pre, Pro

```

```

471     30, 25, 10, 0),    # Pro -> Benign, Early, Pre, Pro
472   nrow = 4, byrow = TRUE
473 )
474 rownames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
475 colnames(cost_matrix) <- c("Benign", "Early", "Pre", "Pro")
476
477 # Extract confusion matrix data
478 cm_table <- conf_matrix$table
479 classes <- rownames(cm_table)
480
481 visualize_cost_confusion_matrix <- function(conf_matrix, cost_matrix) {
482   # Extract the confusion matrix from caret's result
483   cm_table <- conf_matrix$table
484
485   # Create cost-weighted matrix
486   cost_weighted_cm <- matrix(0,
487                               nrow = nrow(cost_matrix),
488                               ncol = ncol(cost_matrix))
489
490   for (i in 1:nrow(cm_table)) {
491     for (j in 1:ncol(cm_table)) {
492       if (i != j) {
493         cost_weighted_cm[i, j] <- cm_table[i, j] * cost_matrix[i, j]
494       }
495     }
496   }
497
498   # Set row and column names
499   rownames(cost_weighted_cm) <- rownames(cost_matrix)
500   colnames(cost_weighted_cm) <- colnames(cost_matrix)
501
502   return(cost_weighted_cm)
503 }
504
505 # Usage
506 cost_weighted_matrix <- visualize_cost_confusion_matrix(conf_matrix, cost_matrix)
507 cat("\nCost-Weighted Misclassification Matrix:\n")
508 print(cost_weighted_matrix)
509
510 # Initialize total misclassification cost
511 total_cost <- 0
512
513 # Calculate total misclassification cost by iterating over confusion matrix
      entries
514 for (i in 1:nrow(cm_table)) {
515   for (j in 1:ncol(cm_table)) {
516     if (i != j) { # Misclassifications (i != j)
517       total_cost <- total_cost + cm_table[i, j] * cost_matrix[i, j]
518     }
519   }
520 }
521
522 # Output the total misclassification cost
523 cat("Total Misclassification Cost:", total_cost, "\n")
524
525 # Calculate average cost per sample
526 total_samples <- sum(cm_table)
527 avg_cost_per_sample <- total_cost / total_samples
528 cat("Average Cost Per Sample:", avg_cost_per_sample, "\n")

```

```

529
530 # Calculate cost-sensitive accuracy
531 # Compute the maximum possible cost (if all instances were misclassified to the
531   # most costly class)
532 max_possible_cost <- sum(cost_matrix) * (total_samples / nrow(cost_matrix))
533 cost_sensitive_accuracy <- 1 - (total_cost / max_possible_cost)
534 cat("Cost-Sensitive Accuracy:", cost_sensitive_accuracy, "\n")
535
536 # Using cost matrix for cost-weighted F1 calculation
537 f1_scores <- conf_matrix$byClass[, "F1"]
538 weights <- 1 / (diag(cost_matrix) + 1) # Add 1 to avoid divide-by-zero
539 weights <- weights / sum(weights)
540 weighted_f1 <- sum(f1_scores * weights, na.rm = TRUE)
541 cat("Weighted F1-score", weighted_f1, "\n")
542
543 # Save the trained model as an HDF5 file
544 save_model_hdf5(model, "convnext_model.h5")

```